

Chapter 3


Transport Layer

Computer Networking: A Top-Down Approach

8th Edition

Jim Kurose, Keith Ross

Pearson, 2021



Modified from the following

James Kurose and Keith Ross, Computer Networking: A Top-Down Approach, 8th Edition, Pearson Education Limited, 2021.

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

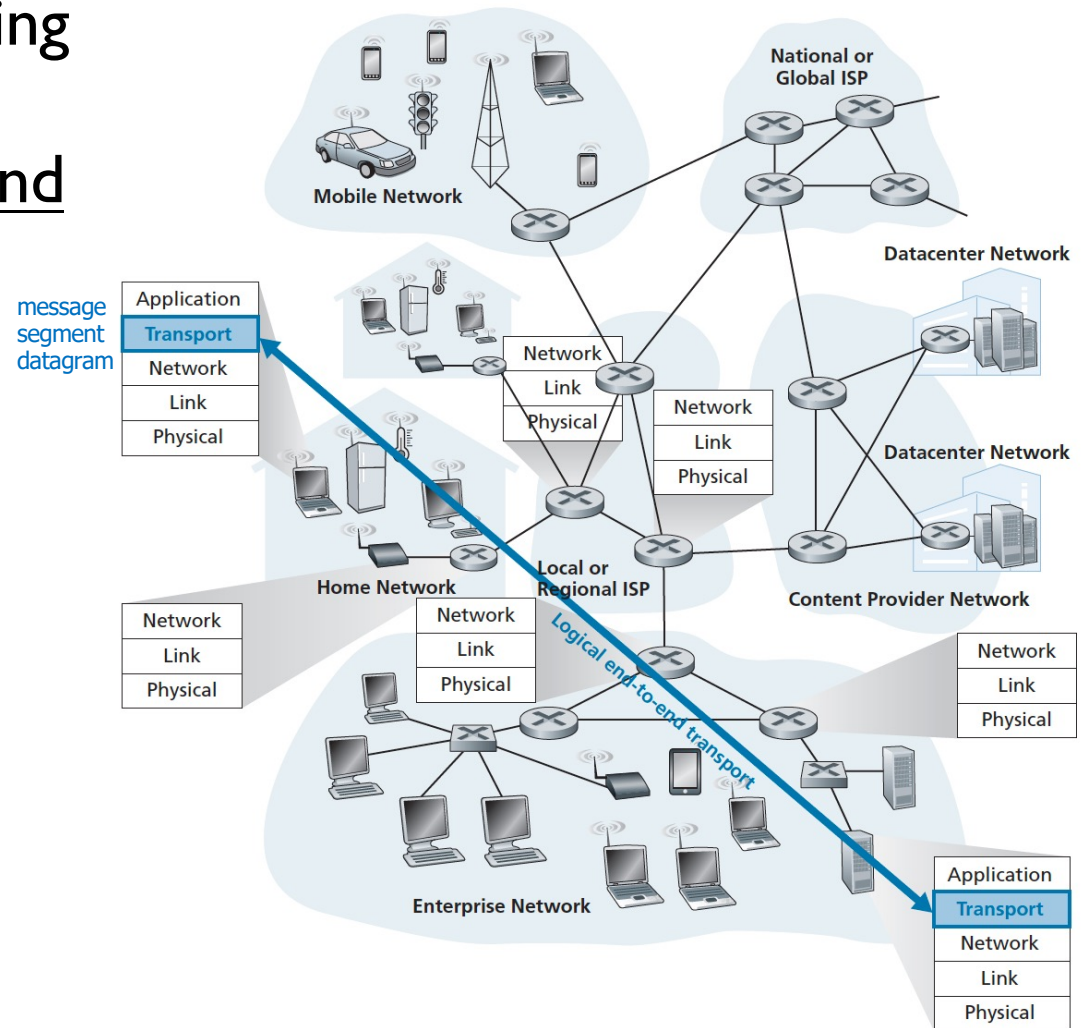
- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

Transport services and protocols

- ❖ Provide *logical communication* between app *processes* running on different hosts
- ❖ Transport protocols run in end systems
 - *Send side*: breaks app messages into *segments*, passes to network layer
 - *Rcv side*: reassembles segments into *messages*, passes to app layer
- ❖ More than one transport protocol available to apps (Internet)
 - TCP (*byte stream*)
 - UDP (*datagram*)



Transport vs. network layer

- ❖ *Network layer*
logical communication
between **hosts**
 - IP
- ❖ *Transport layer*
logical communication
between **processes**
 - Relies on and enhances
network layer services
 - IP + port number

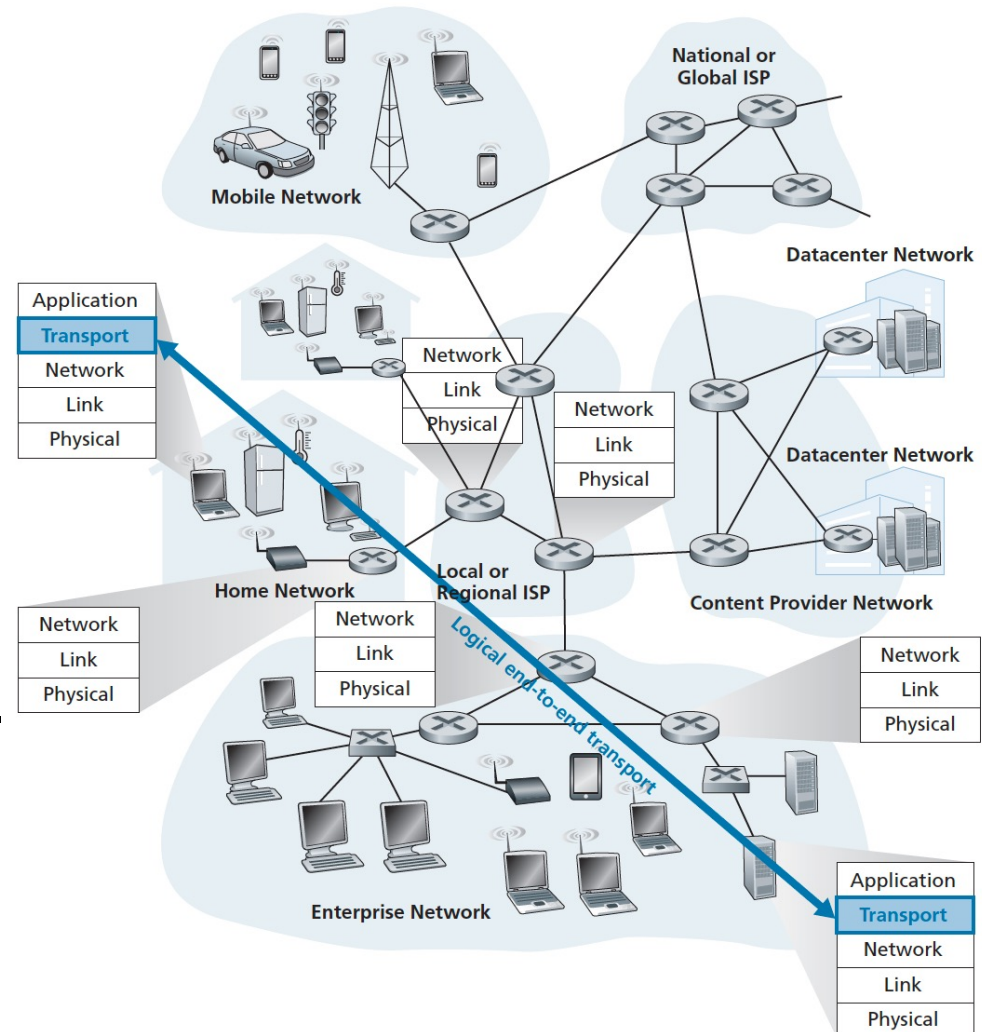
Household analogy

12 kids in Ann's house
sending letters to 12 kids in
Bill's house:

- ❖ **Hosts** = houses
- ❖ **Processes** = kids
- ❖ **App messages**
= letters in envelopes
- ❖ **Network-layer protocol**
= postal service
- ❖ **Transport protocol**
= Ann and Bill who demux
to in-house siblings

Internet transport-layer protocols

- ❖ Reliable, in-order delivery (TCP)
 - Connection setup
 - Flow control (**sender / receiver**)
 - Congestion control (**network**)
- ❖ Unreliable, unordered delivery (UDP)
 - No-frills extension of “best-effort” IP service
- ❖ IP services not available
 - Delay guarantees
 - Bandwidth guarantees



Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

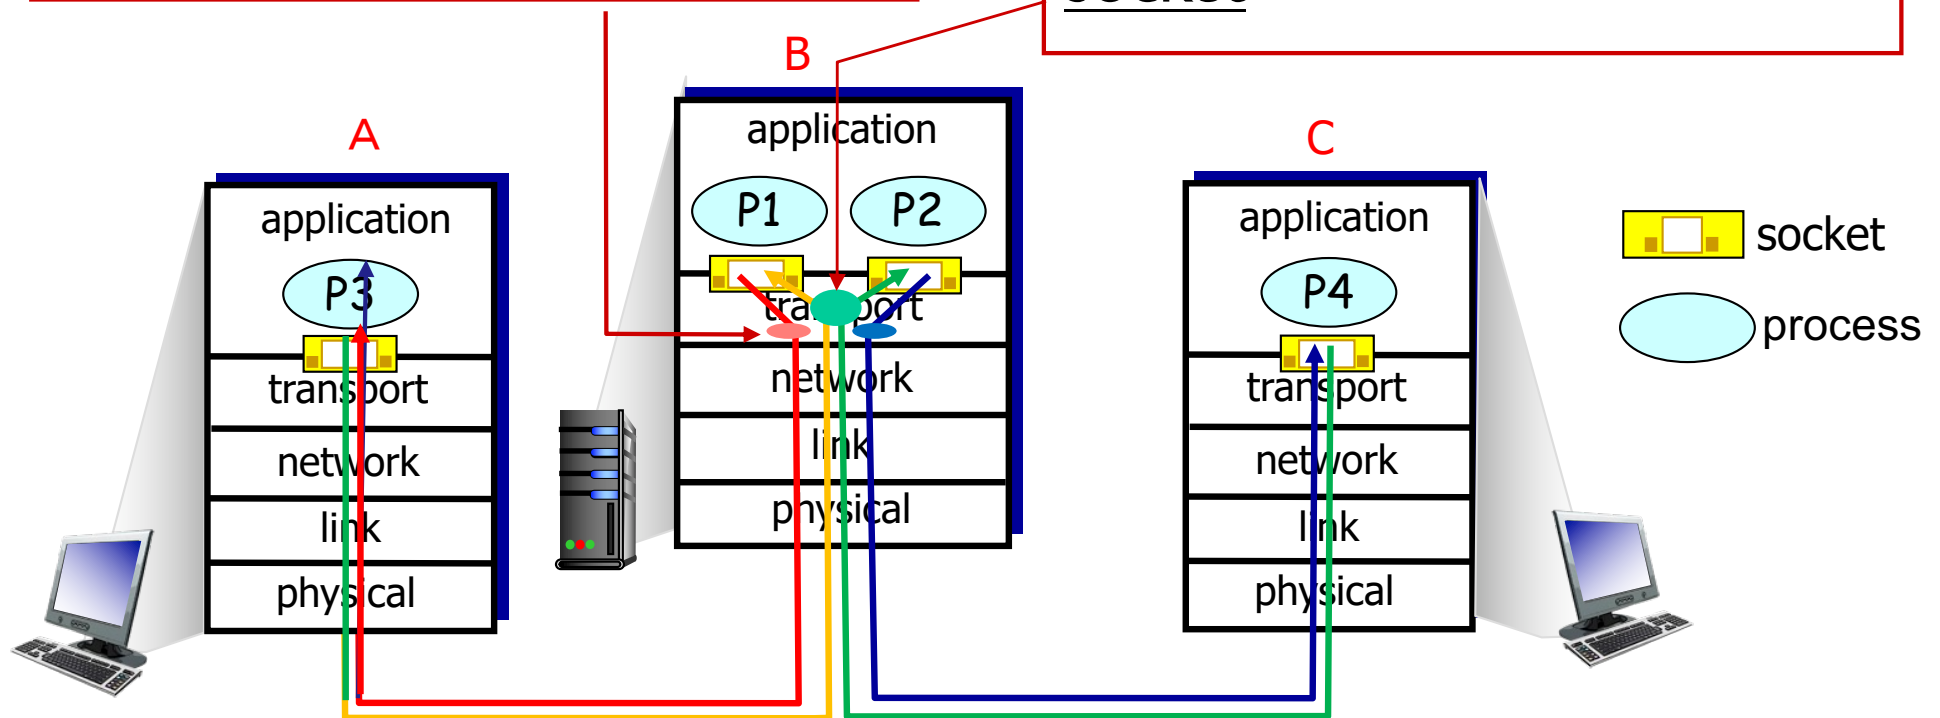
Multiplexing/demultiplexing

Multiplexing at sender

Handle data from multiple sockets, add transport header (later used for demultiplexing)

Demultiplexing at receiver

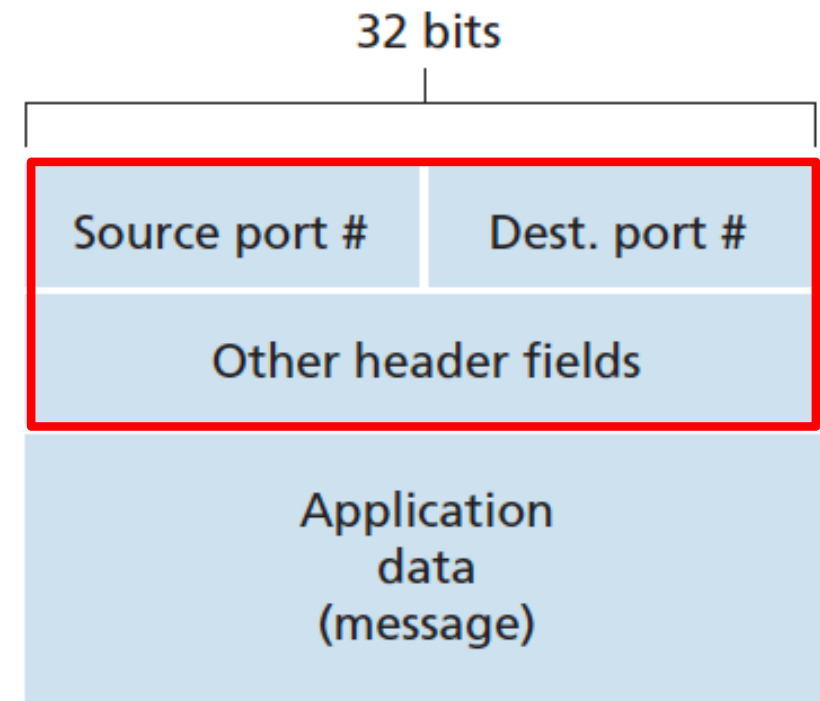
Use header info to deliver received segments to correct socket



How demultiplexing works

- ❖ Host receives IP datagrams
 - Each datagram has network-layer source IP address, destination **IP address**
 - Each datagram carries one transport-layer segment
 - Each segment has source, destination **port number**
- ❖ Host uses **IP addresses** & **port numbers** to direct segment to appropriate socket

Segment: transport layer (TCP, UDP) data unit
Datagram: network layer (IP) data unit
Datagram = IP header + Segment



Source and destination port-number fields in a transport-layer **segment**

Connectionless demultiplexing (UDP)

- ❖ Created socket has host-local port #:

```
DatagramSocket mySocket1  
= new  
DatagramSocket (12534) ;  
                host-local port#
```

- ❖ When creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

- ❖ When host receives UDP segment

- Checks destination port # in segment
- Directs UDP segment to socket with that port #



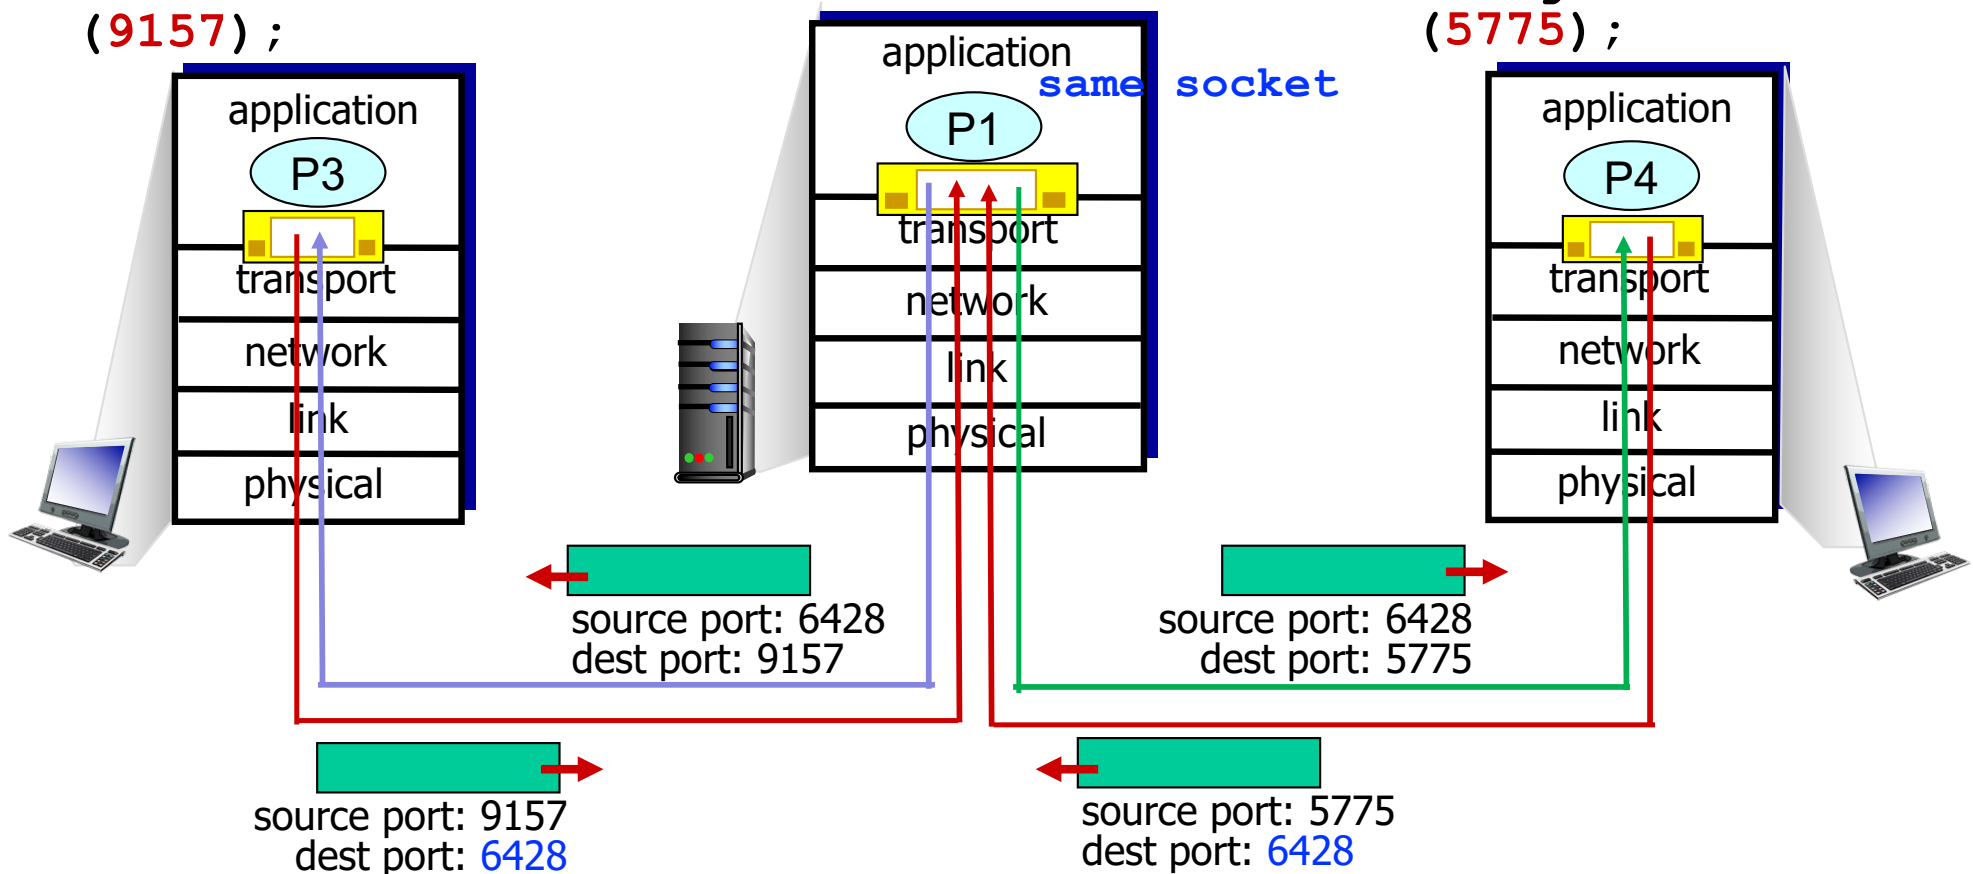
IP datagrams with same dest. port #, but different source IP addresses and/or source port numbers will be directed to the same socket at dest

Connectionless demux: example

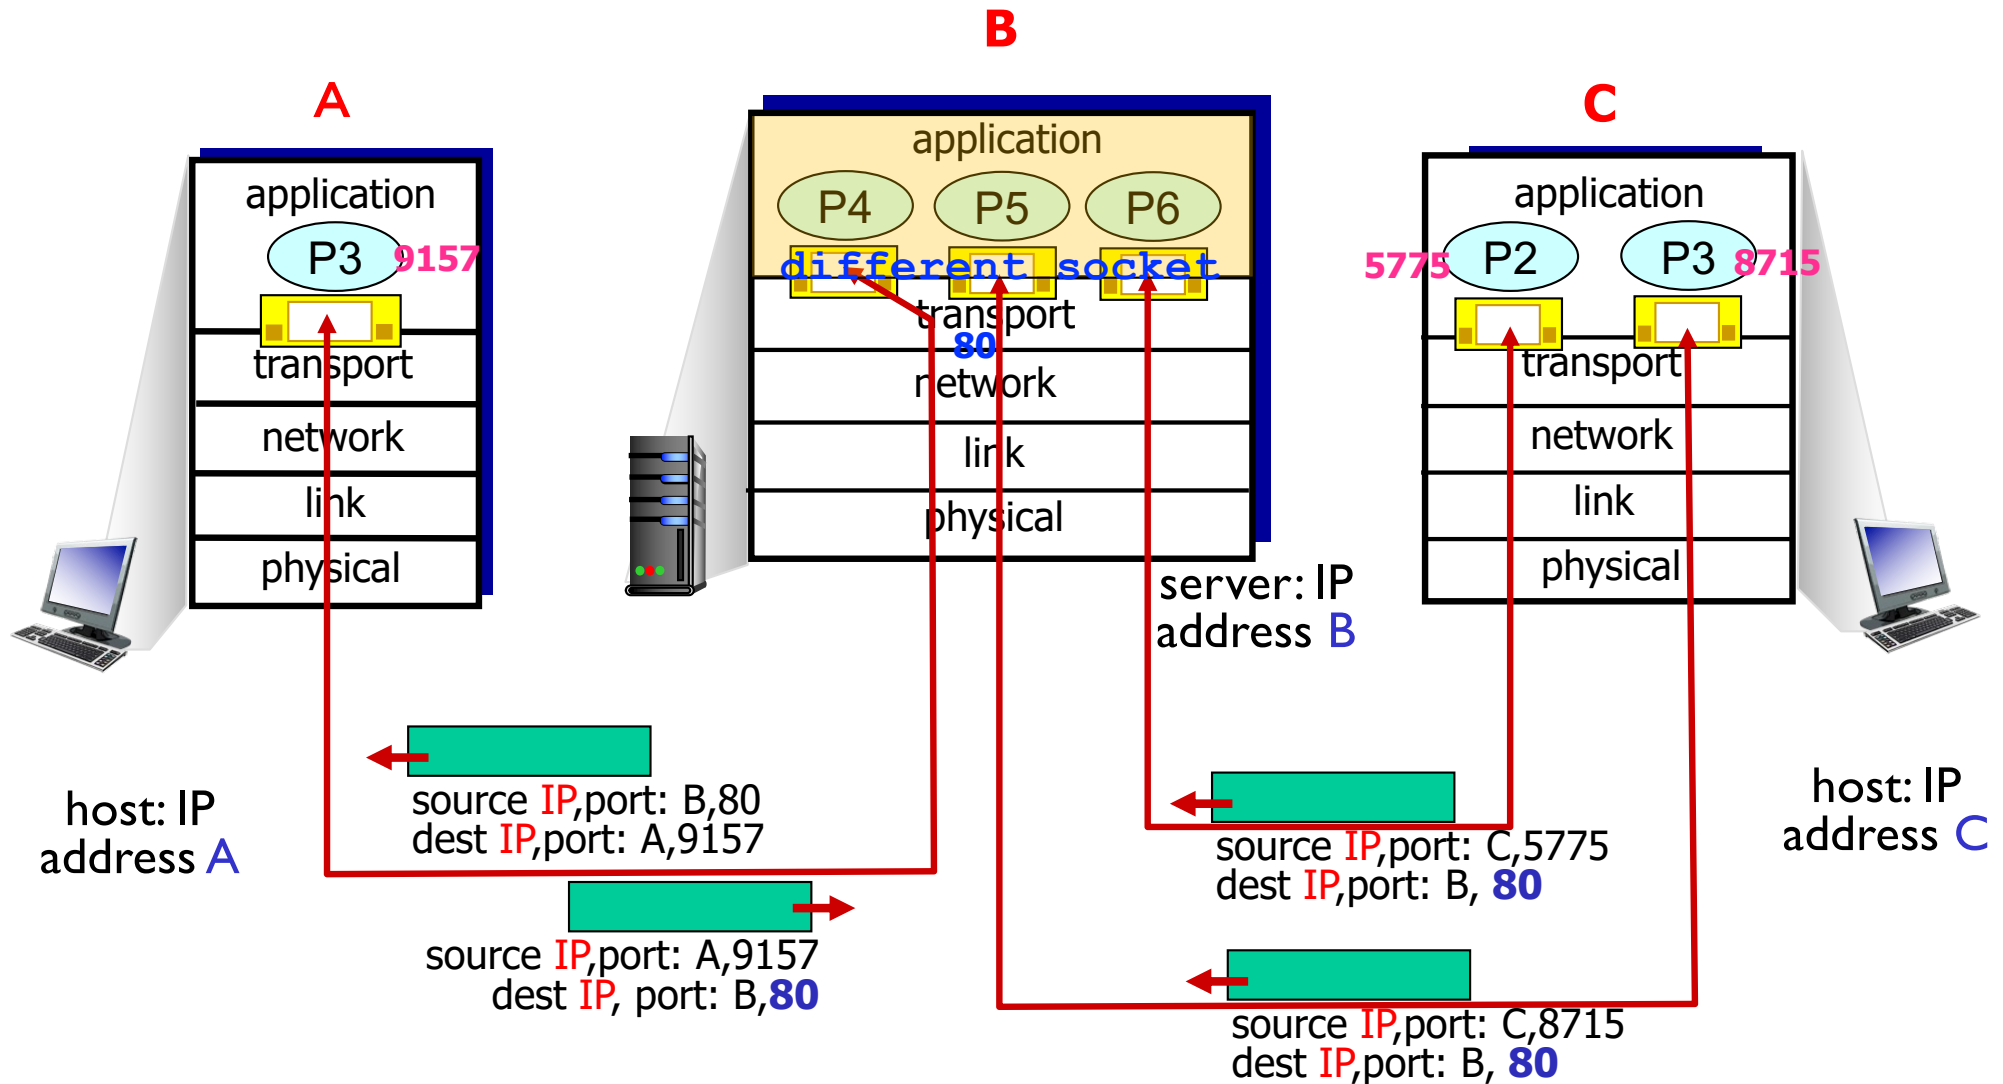
```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428) ;
```

```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775) ;
```



Connection-oriented demux: example

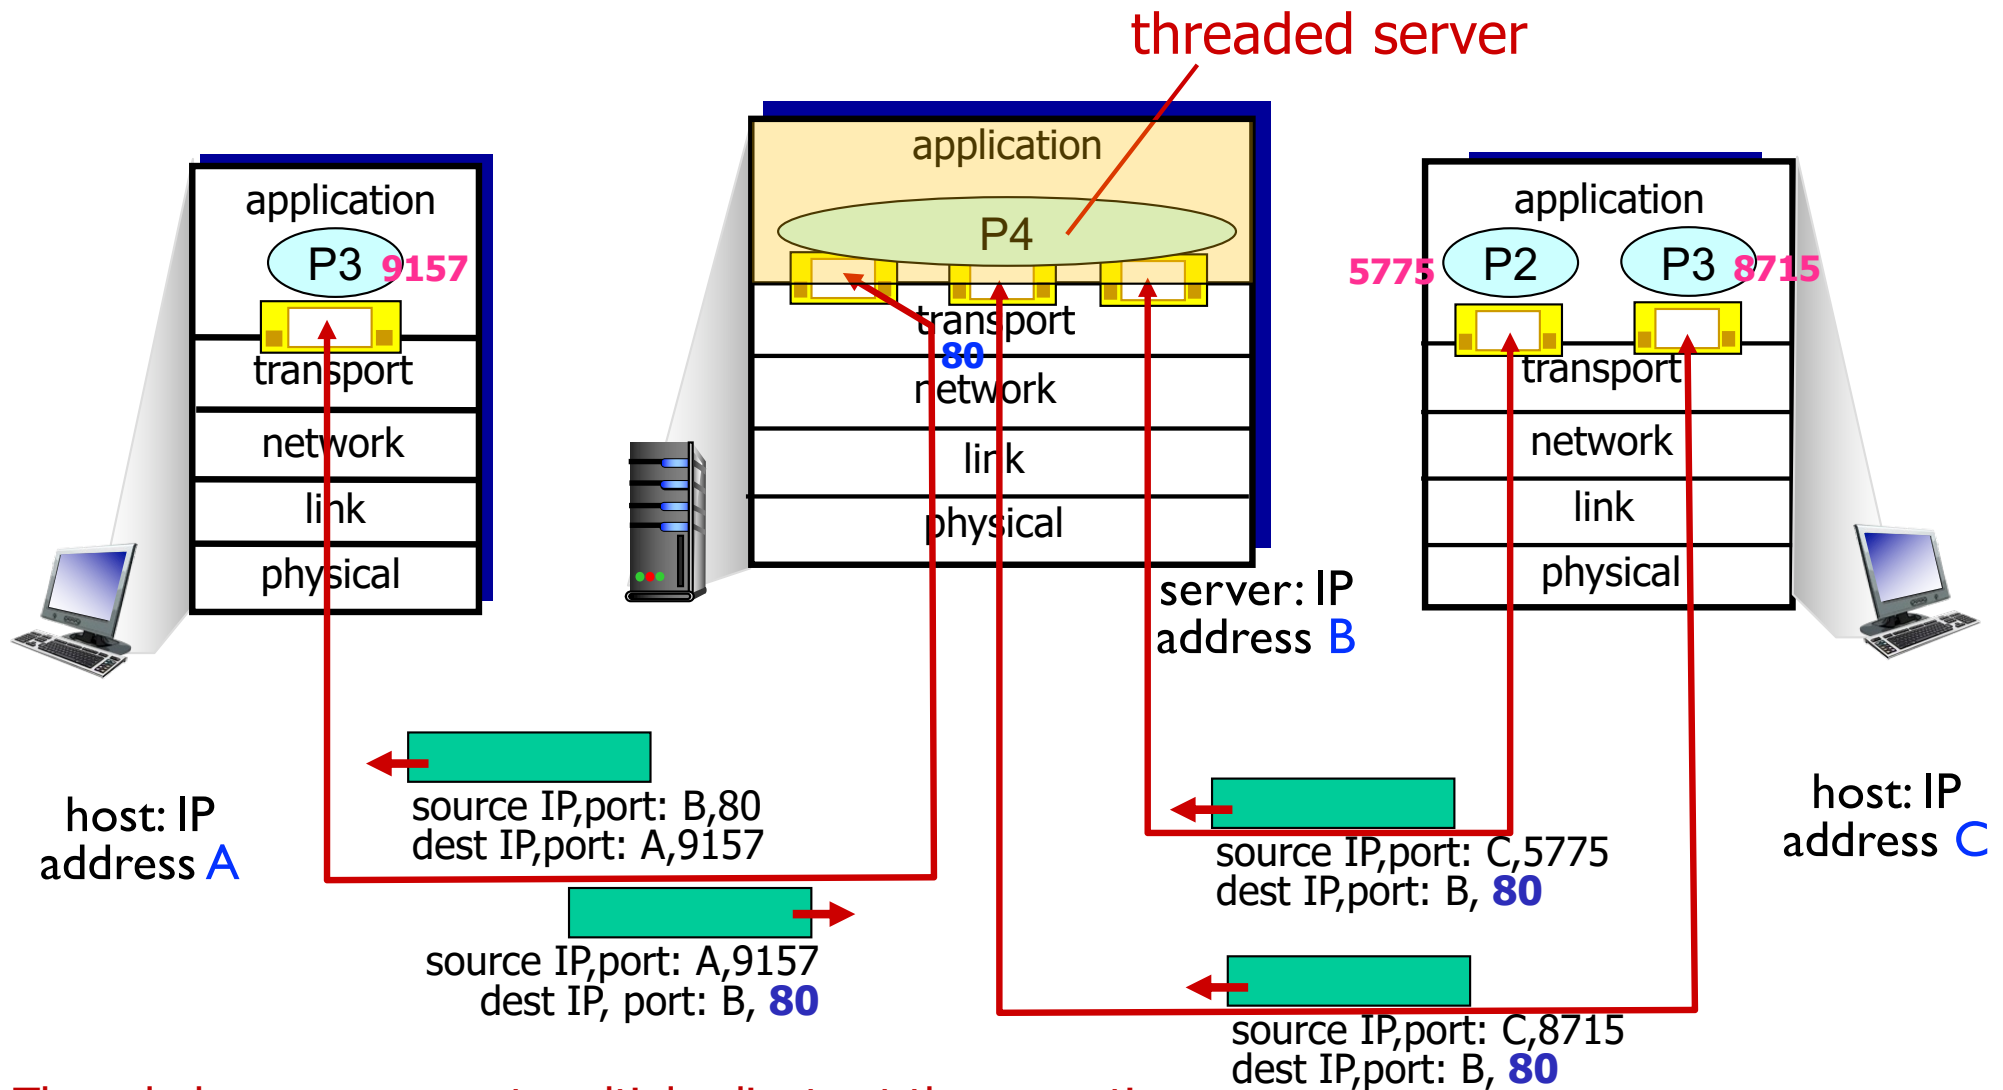


- * Three segments, all destined to IP address: B
- * Dest port: 80 are demultiplexed to *different* sockets

Connection-oriented demux (TCP)

- ❖ TCP socket identified by 4-tuple
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- ❖ Demux: receiver uses all four values to direct segment to appropriate socket
- ❖ Server host may support many **simultaneous TCP sockets**
 - Each socket identified by its own 4-tuple
- ❖ Web servers have **different sockets** for each connecting **client**
 - Non-persistent HTTP will have different socket for each request

Connection-oriented demux: example



Threaded server: support multiple clients at the same time
(whenever a client request comes, a separate thread can
be assigned for handling each request)

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

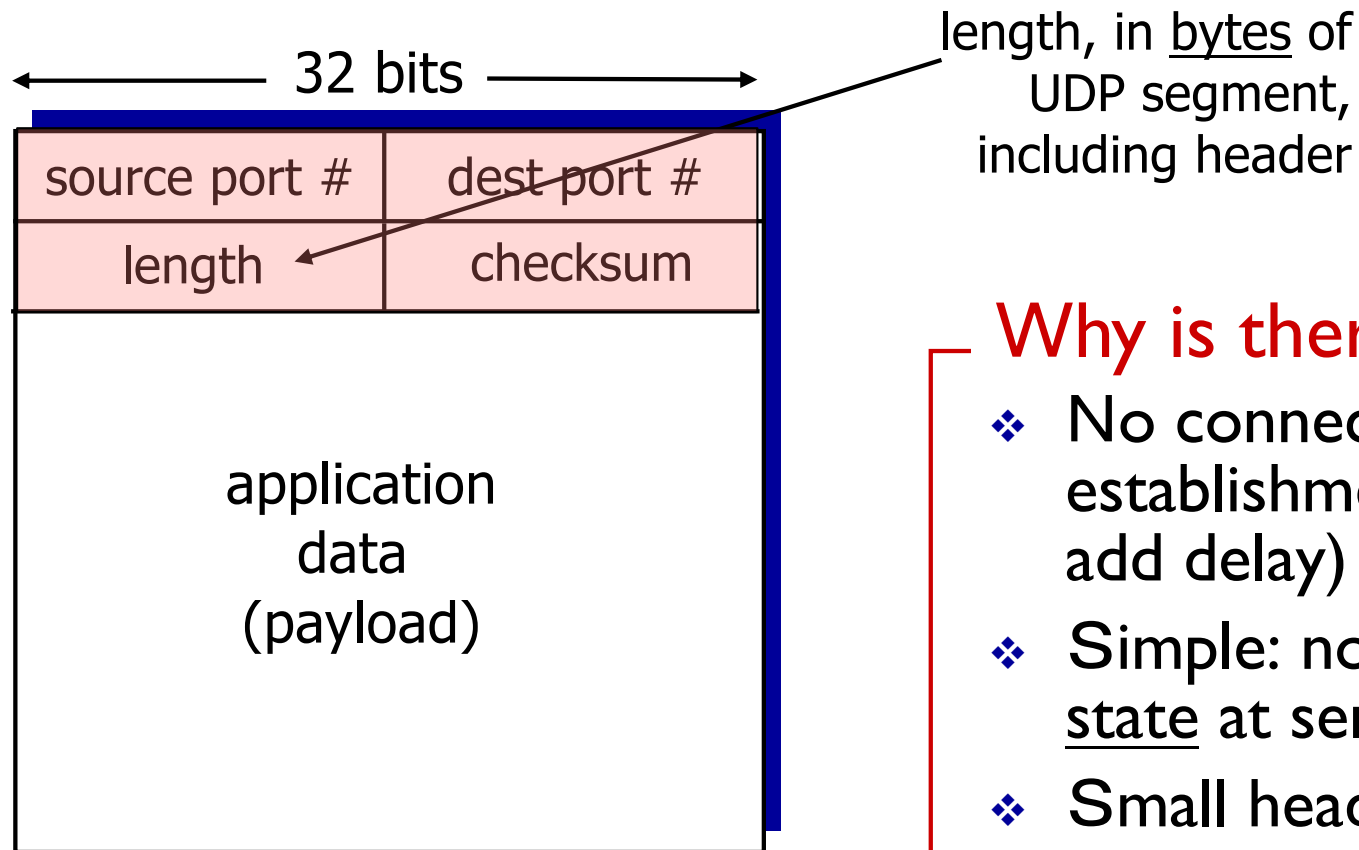
3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

- ❖ “No frills,” “bare bones” Internet transport protocol
- ❖ “Best effort” service, UDP segments may be
 - Lost
 - Delivered out-of-order to app
- ❖ *Connectionless*
 - No handshaking between UDP sender, receiver
 - Each UDP segment handled independently of others
- ❖ UDP apps
 - Streaming multimedia apps (loss tolerant, data rate sensitive)
 - DNS
 - SNMP (Simple Network Management Protocol)
 - Eg., query no. of ports, network interface of port, turn off port, etc.
- ❖ Reliable transfer over UDP
 - Add reliability (simple) at application layer
 - Application-specific error recovery!

No IP address fields
(network layer header)

UDP: segment header



UDP **segment** format

Why is there a UDP?

- ❖ No connection establishment (which can add delay)
- ❖ Simple: no connection state at sender, receiver
- ❖ Small header size
- ❖ No congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender

- ❖ Treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ **Checksum**: addition (one's complement sum) of segment contents
- ❖ Sender puts checksum value into UDP checksum field

Receiver

- ❖ Compute checksum of received segment
- ❖ Check if computed checksum equals checksum field value
 - **NO** - error detected
 - **YES** - no error detected
But maybe errors nonetheless? More later
....

Internet checksum: example

Example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum																1
checksum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP


- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

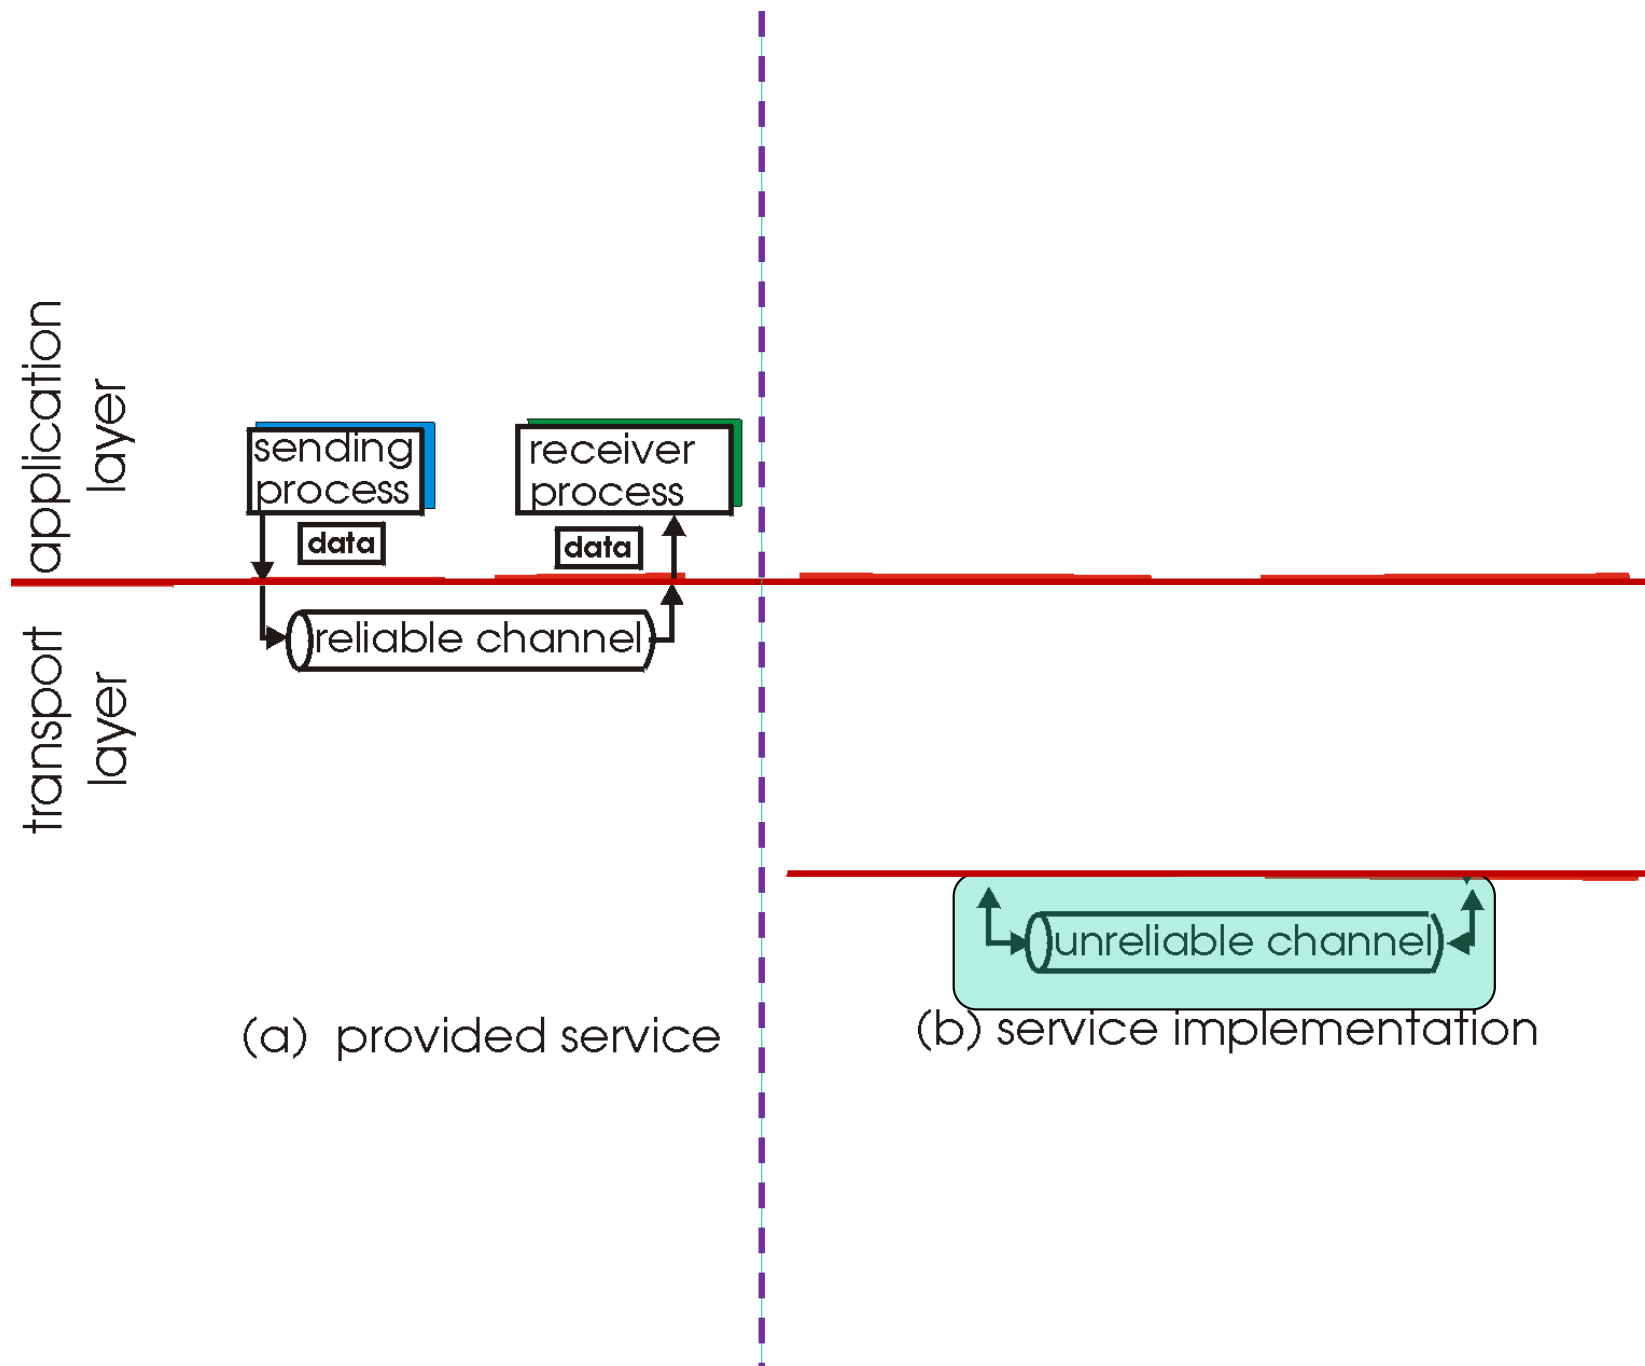
Principles of reliable data transfer

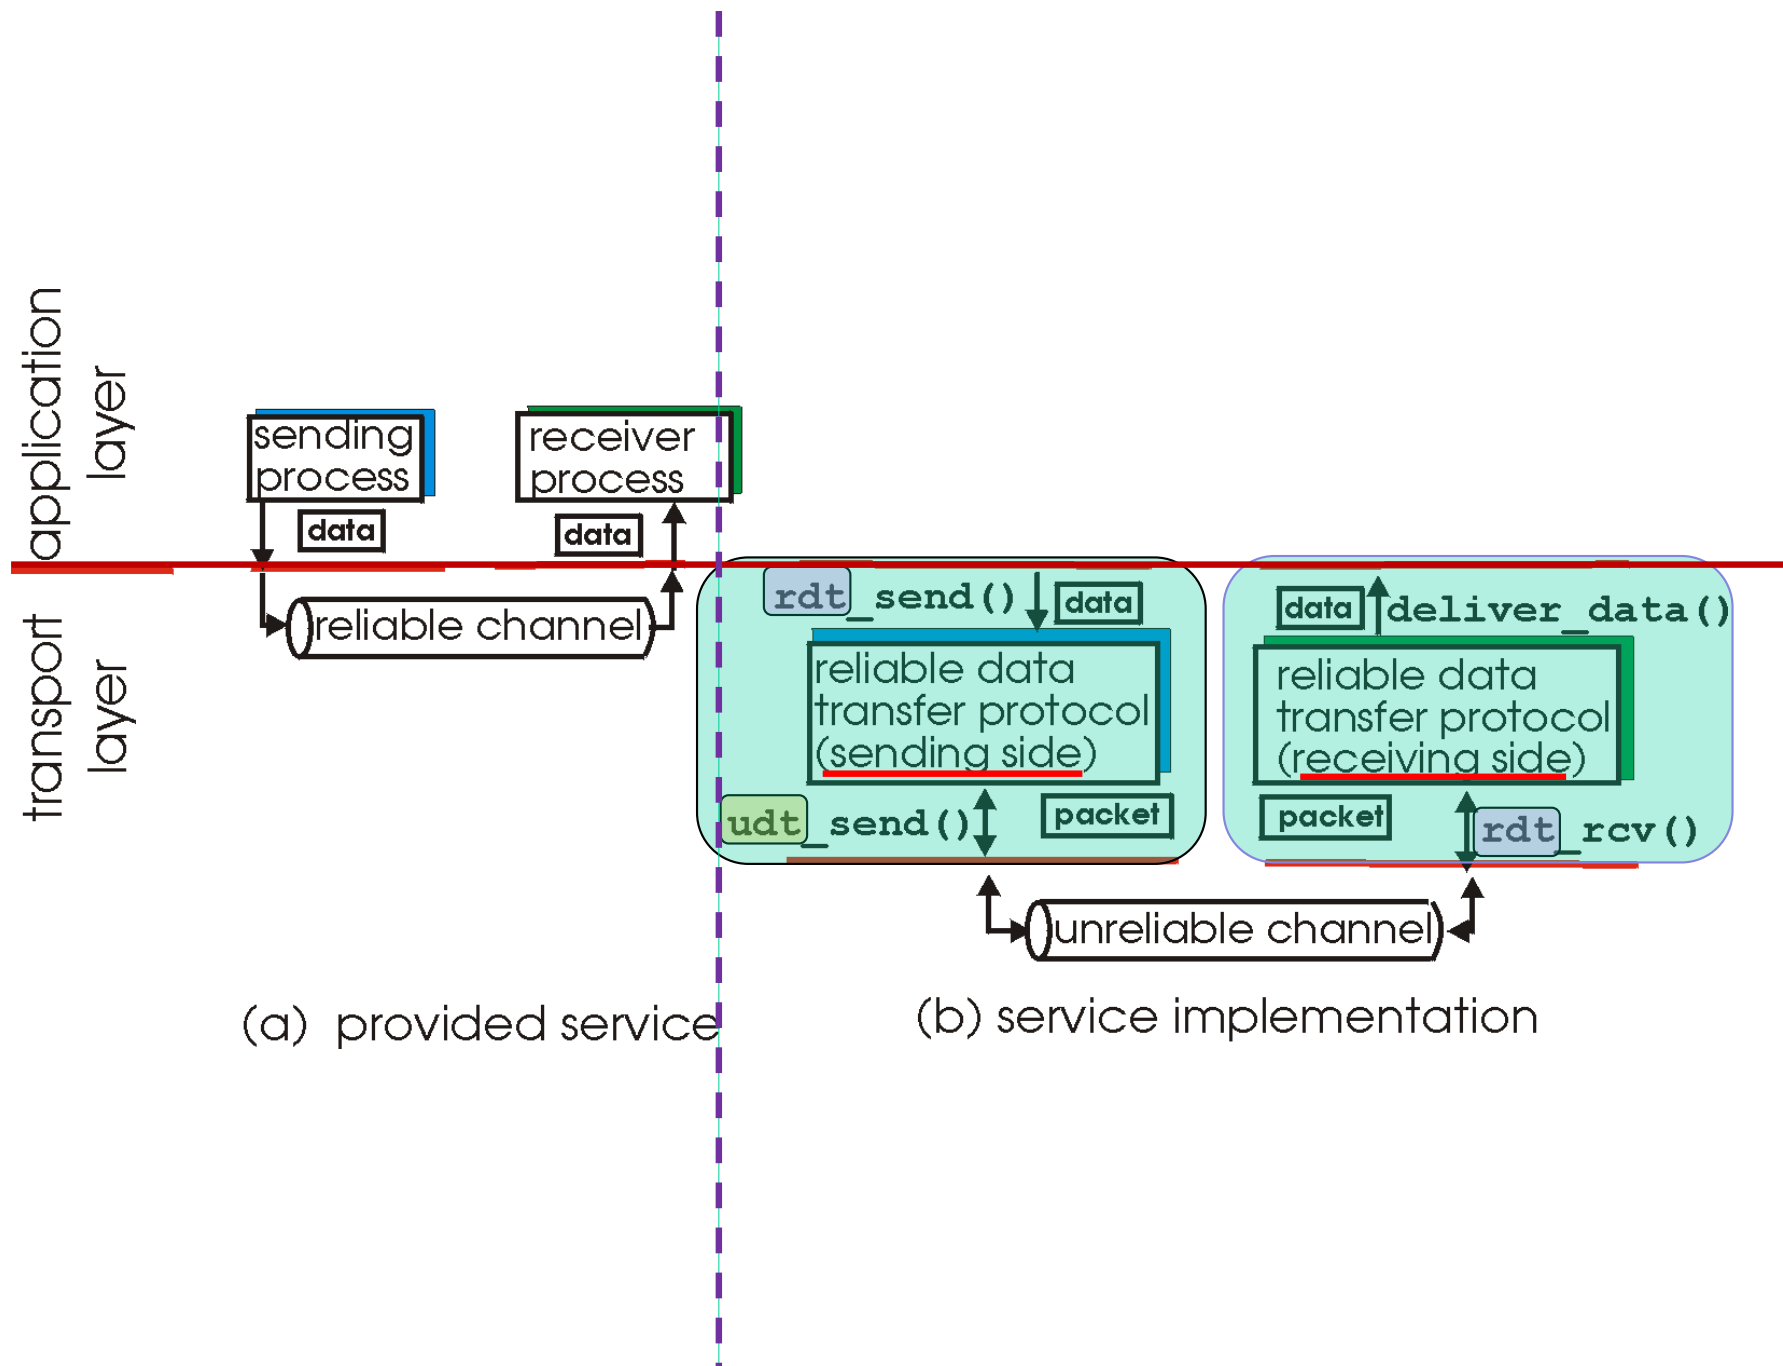
- ❖ Important in application, transport, link layers
 - Top-10 list of important networking topics!



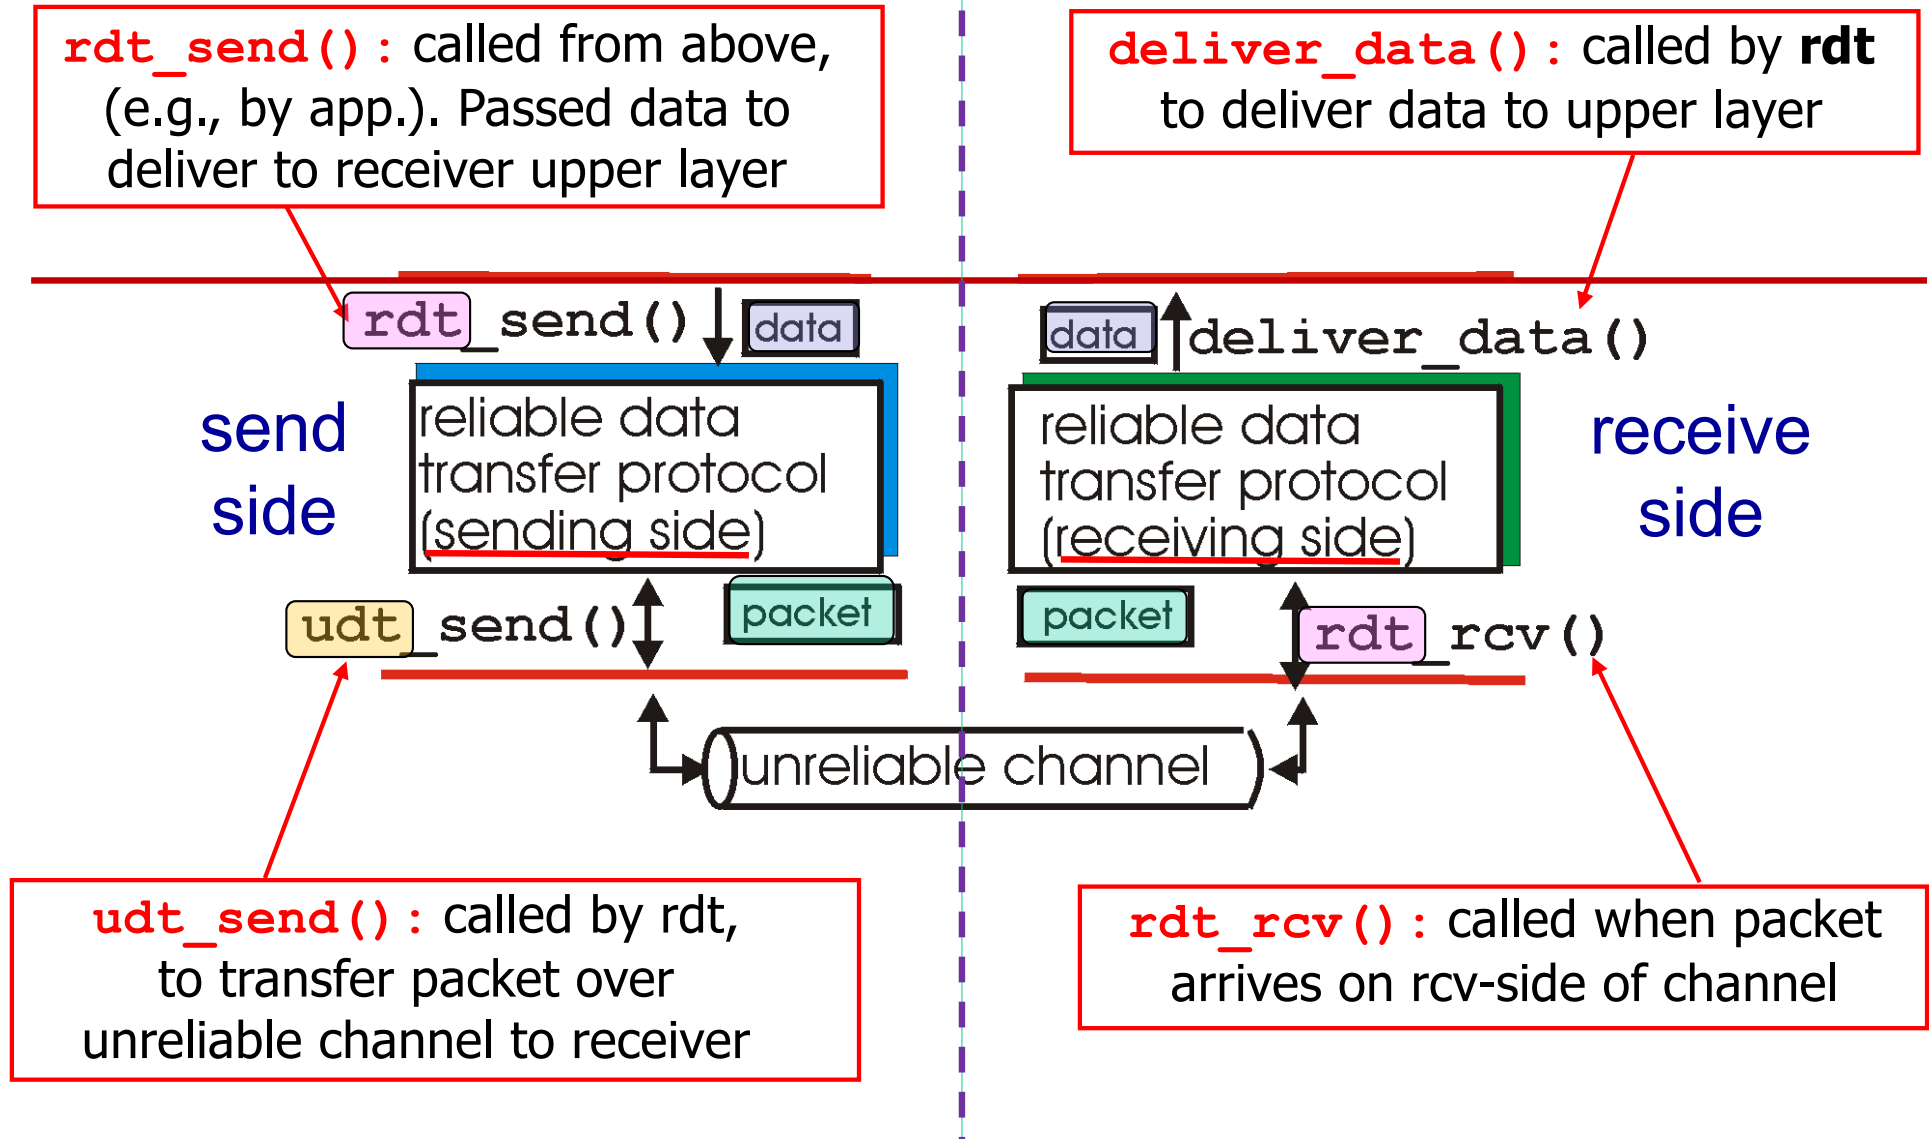
(a) provided service

- ❖ Characteristics of unreliable channel will determine complexity of reliable data transfer protocol (**rdt**)

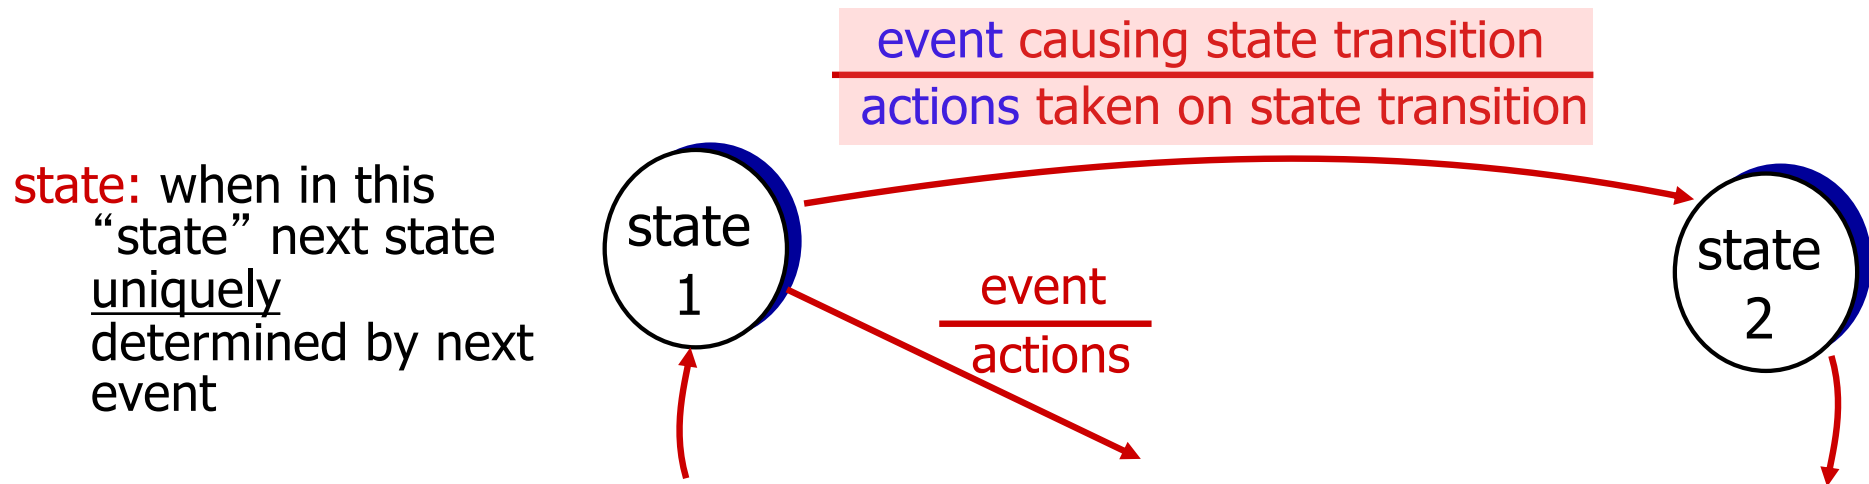




Reliable data transfer: getting started

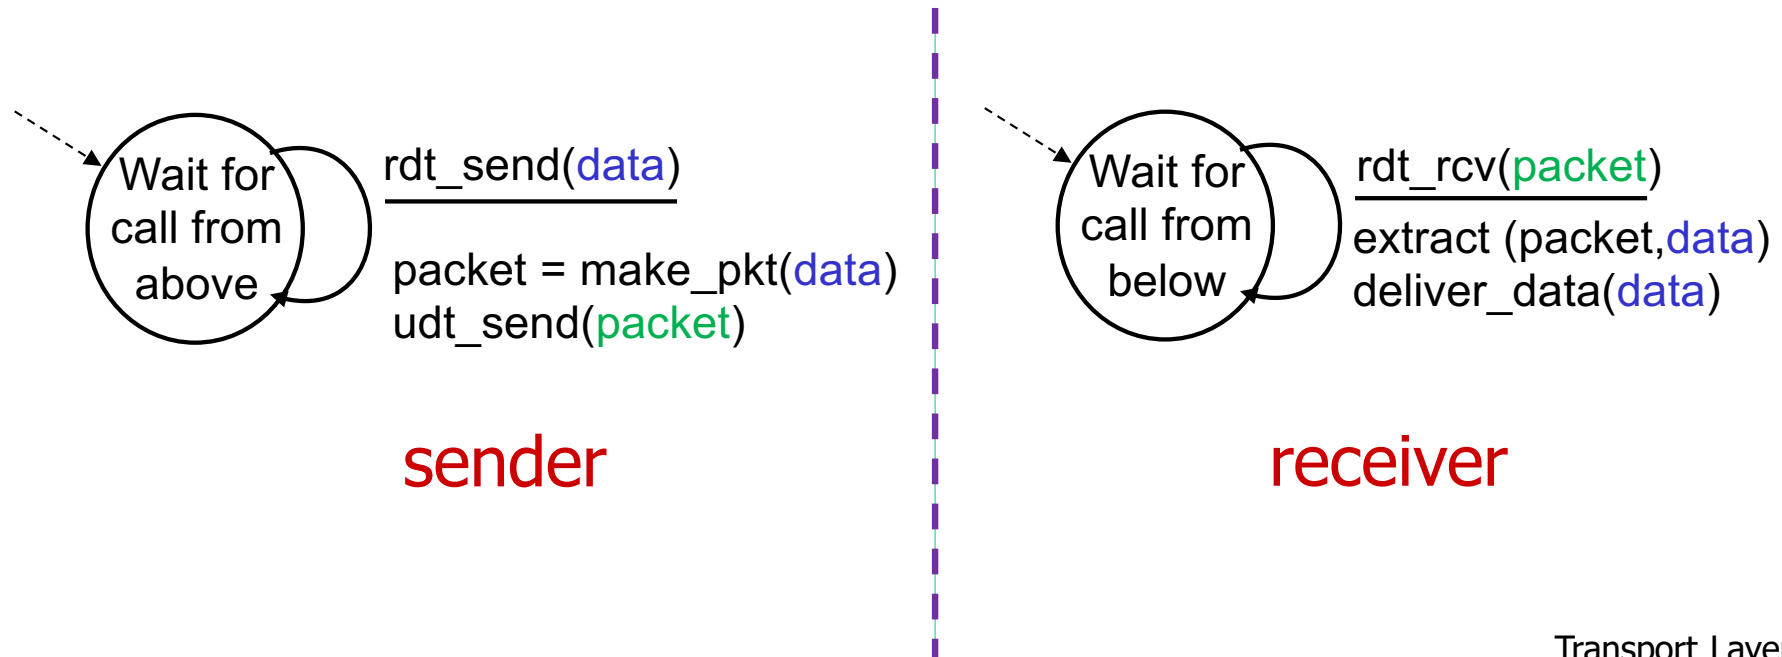


- ❖ Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ Consider only unidirectional data transfer
 - But control info (for connection) will flow on both directions
- ❖ Use finite state machines (FSM) to specify sender, receiver



rdt1.0: reliable transfer over a reliable channel

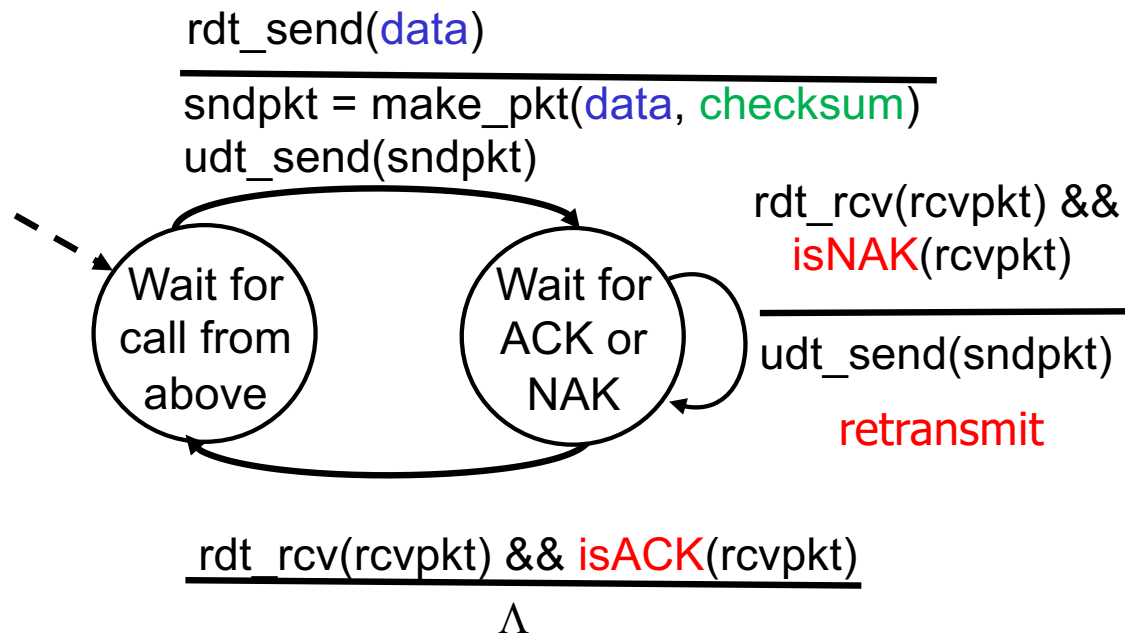
- ❖ Underlying channel perfectly reliable
 - No bit errors
 - No loss of packets
- ❖ Separate FSMs for sender, receiver
 - Sender sends data into underlying channel
 - Receiver reads data from underlying channel



rdt2.0: channel with bit errors

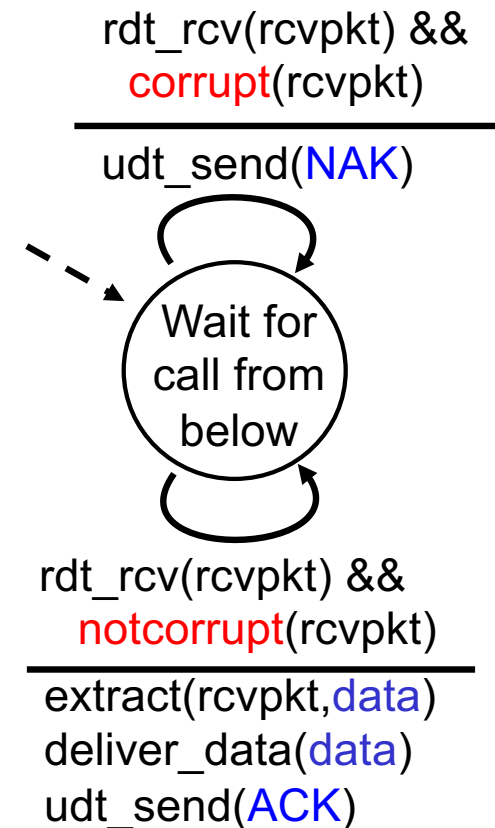
- ❖ Underlying channel may flip bits in packet
 - Checksum to detect bit errors
- ❖ Question: how to recover from errors
 - *Acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *Negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - Sender retransmits pkt on receipt of NAK
- ❖ New mechanisms in rdt2.0 (beyond rdt1.0)
 - Error detection
 - Receiver feedback: control msgs (ACK, NAK) from receiver to sender

rdt2.0: FSM specification

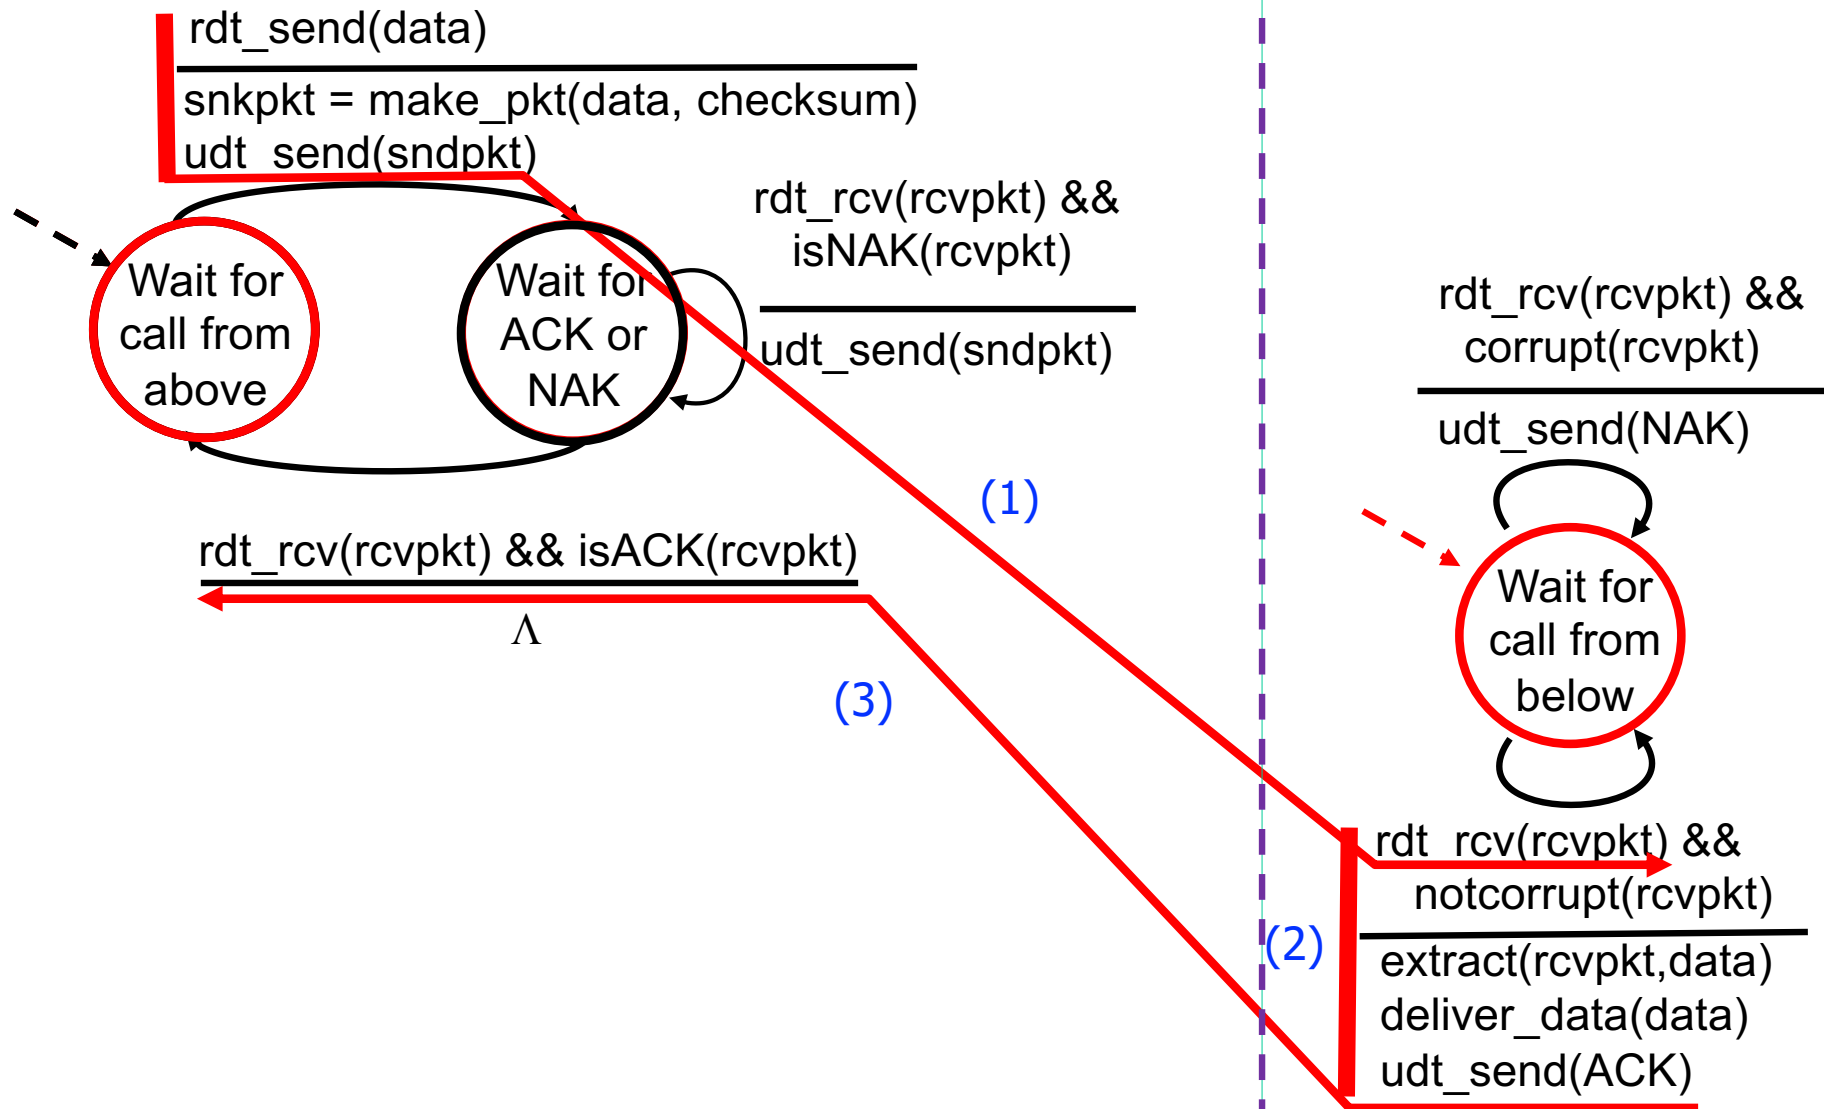


sender

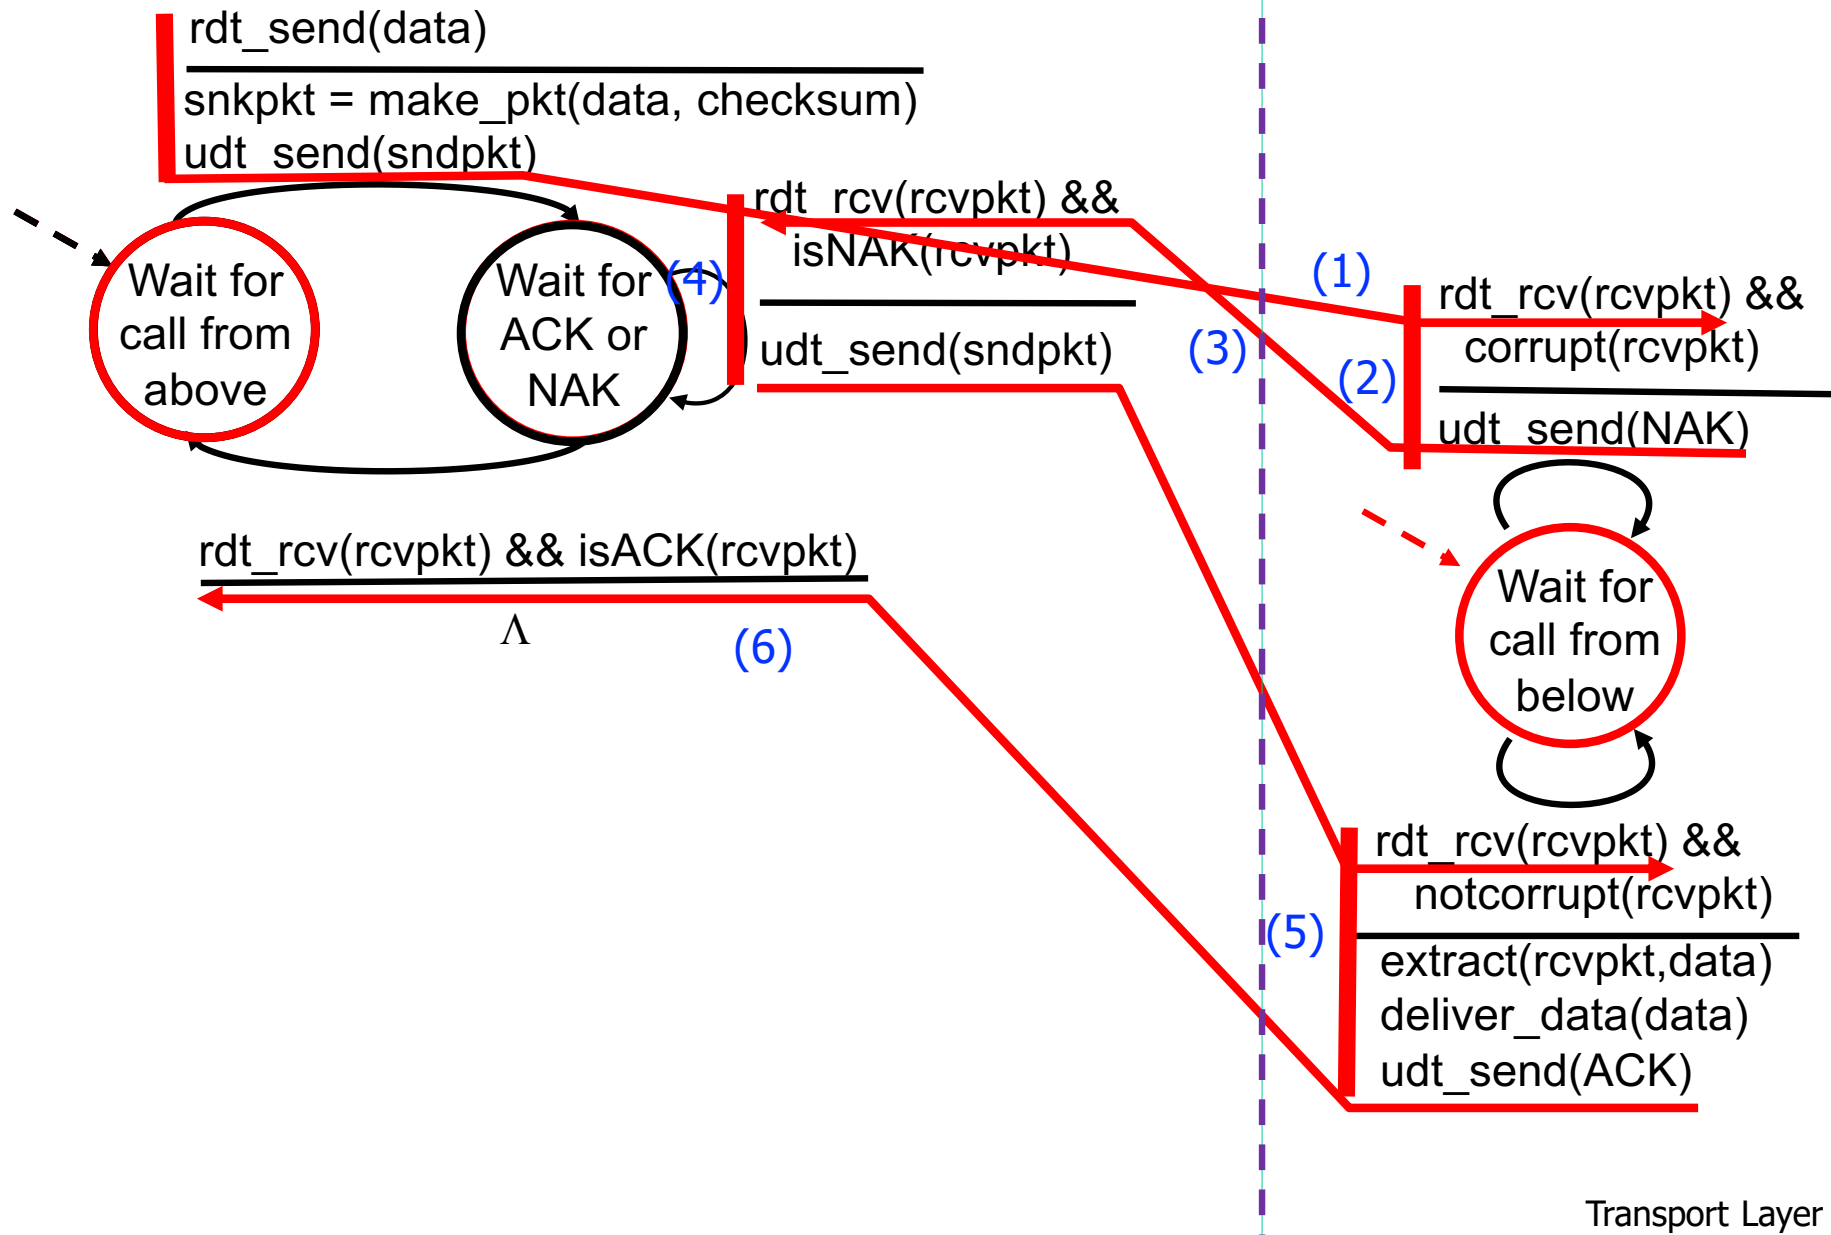
receiver



rdt2.0: operation with no errors



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

What happens if ACK/NAK corrupted?

- ❖ Sender doesn't know what happened at receiver!
- ❖ Can't just retransmit: possible duplicate

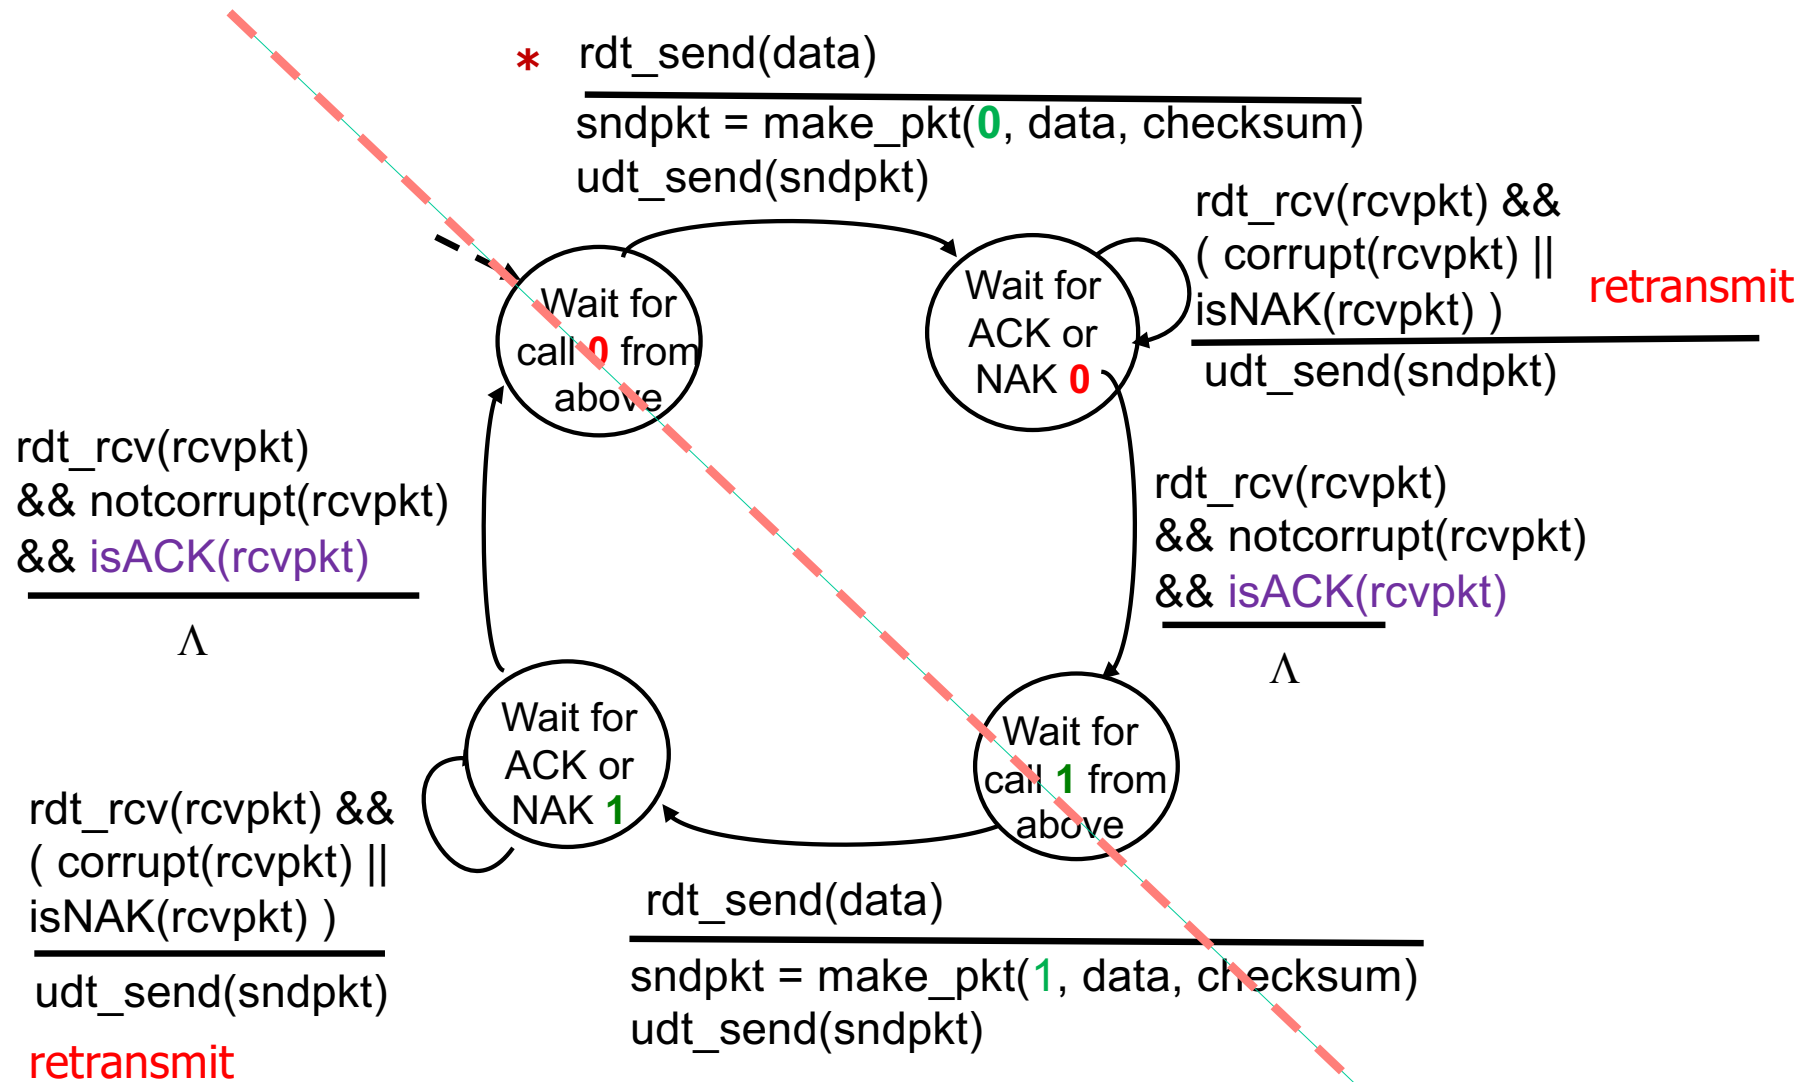
Handling duplicates

- ❖ Sender retransmits current pkt if ACK/NAK corrupted
- ❖ Sender adds *sequence number* (0 or 1) to each pkt to prevent duplicate pkt due to receiving garbled ACK/NAK
- ❖ Receiver discards (doesn't deliver up) duplicate pkt

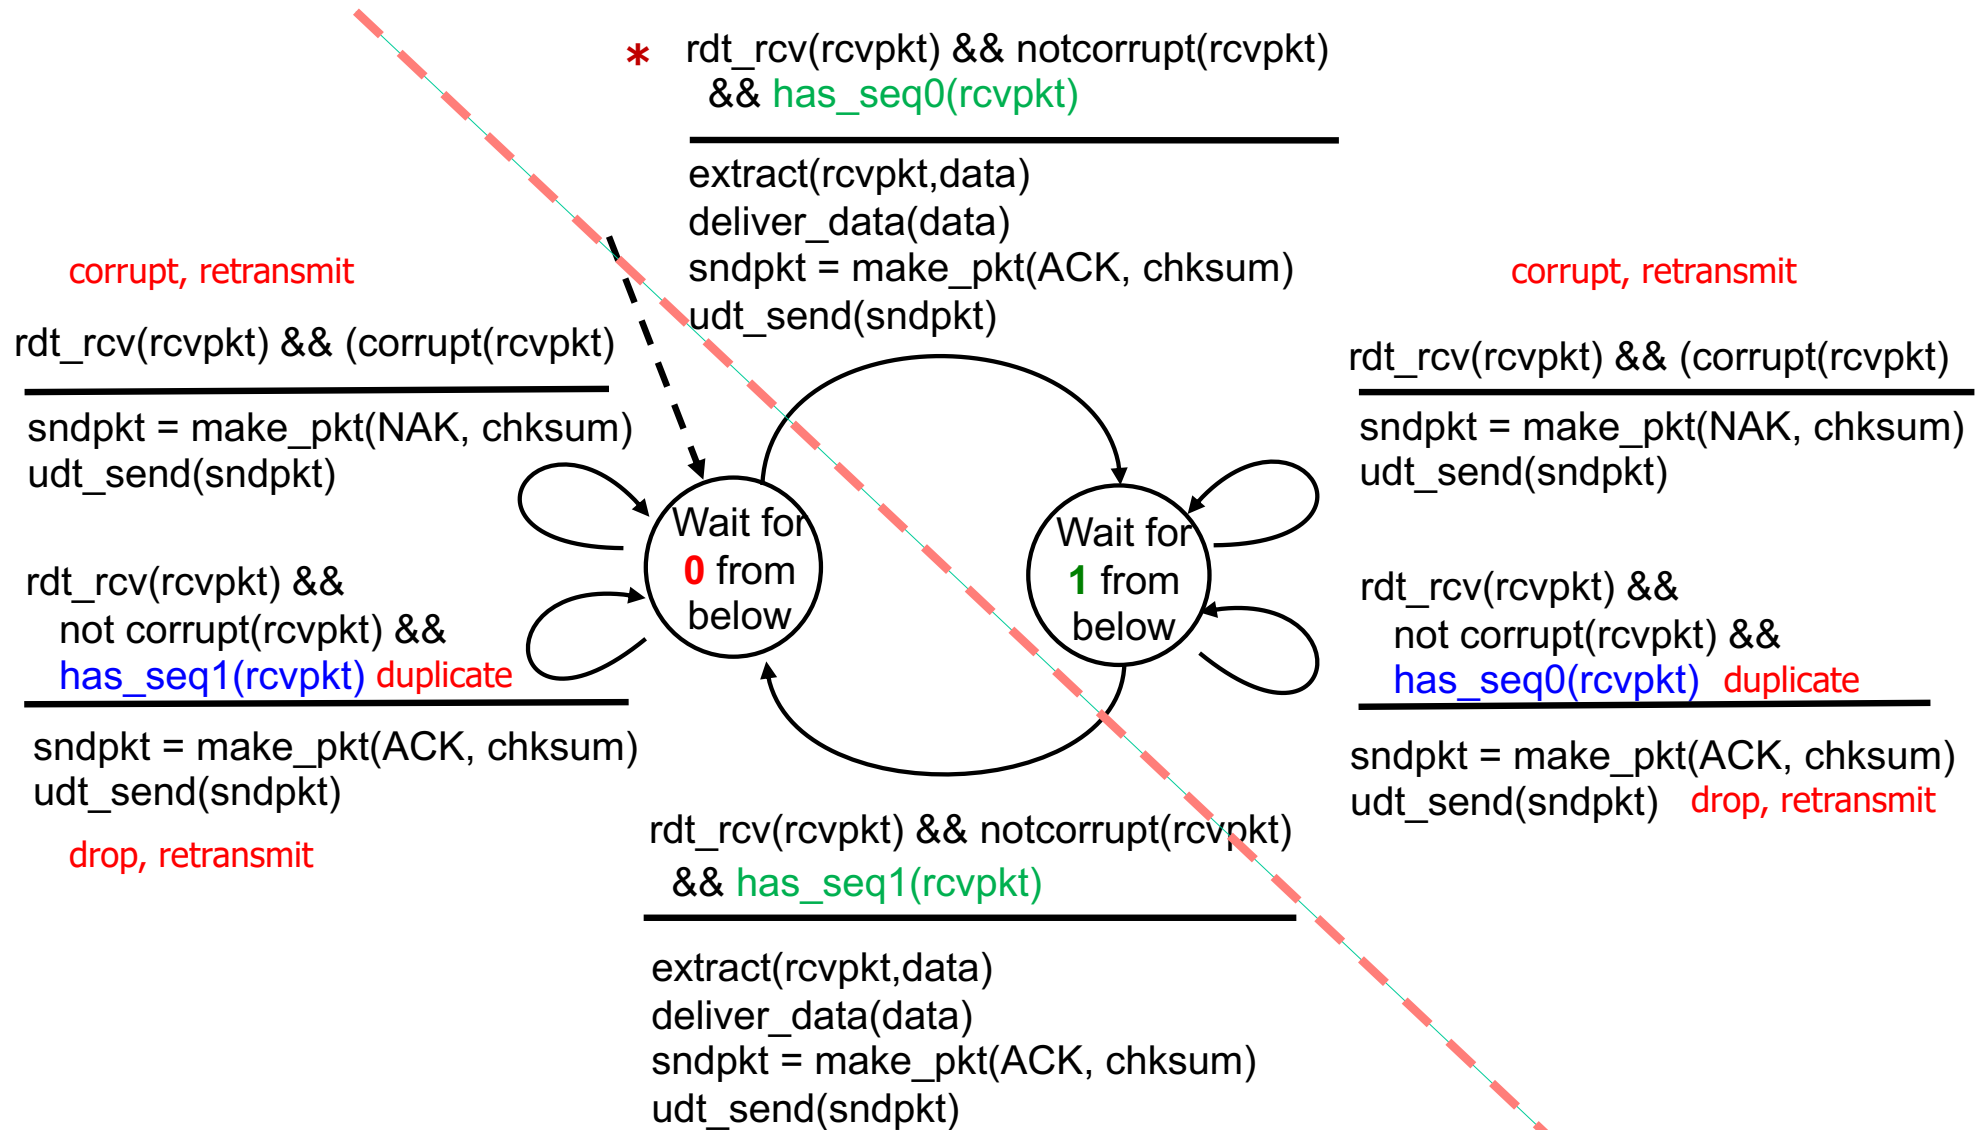
Stop and wait

Sender sends one packet,
then waits for receiver
response

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles garbled ACK/NAKs



rdt2.1: discussion

Sender

- ❖ seq # added to pkt
- ❖ Two seq. #'s (0,1) is suffice
- ❖ Must check if received ACK/NAK corrupted
- ❖ Twice as many states
 - State must “remember” whether “expected” pkt should have seq # of 0 or 1

Receiver

- ❖ Must check if received packet is duplicate
 - State indicates whether 0 or 1 is expected pkt seq #
- ❖ Note: receiver *can not know* if its last ACK/NAK received OK at sender

rdt2.2: a NAK-free protocol

- ❖ Same functionality as rdt2.1, using ACKs only
- ❖ Instead of NAK, receiver sends ACK for last pkt successfully received OK
 - Receiver must *explicitly include seq #* of pkt being ACKed
- ❖ Duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: sender, receiver fragments

sender FSM fragment

* rdt_send(data)

sndpkt = make_pkt(0, data, checksum)

udt_send(sndpkt)

Wait for
call 0 from
above

Wait for
ACK
0

rdt_rcv(rcvpkt) &&
(corrupt(rcvpkt) ||
isACK(rcvpkt,1))

udt_send(sndpkt)
drop, retransmit

rdt_rcv(rcvpkt)
&& notcorrupt(rcvpkt)
&& **isACK(rcvpkt,0)**

Λ

rdt_rcv(rcvpkt) &&
(corrupt(rcvpkt) ||
has_seq1(rcvpkt))

**sndpkt = make_pkt(ACK,1,
chksum)**
udt_send(sndpkt)

drop, retransmit

Wait for
0 from
below

Wait for
1 from
below

rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
&& **has_seq1(rcvpkt)**

extract(rcvpkt,data)

deliver_data(data)

sndpkt = make_pkt(ACK,1, chksum)

udt_send(sndpkt)

receiver FSM fragment

rdt3.0: channels with errors and loss

New assumption

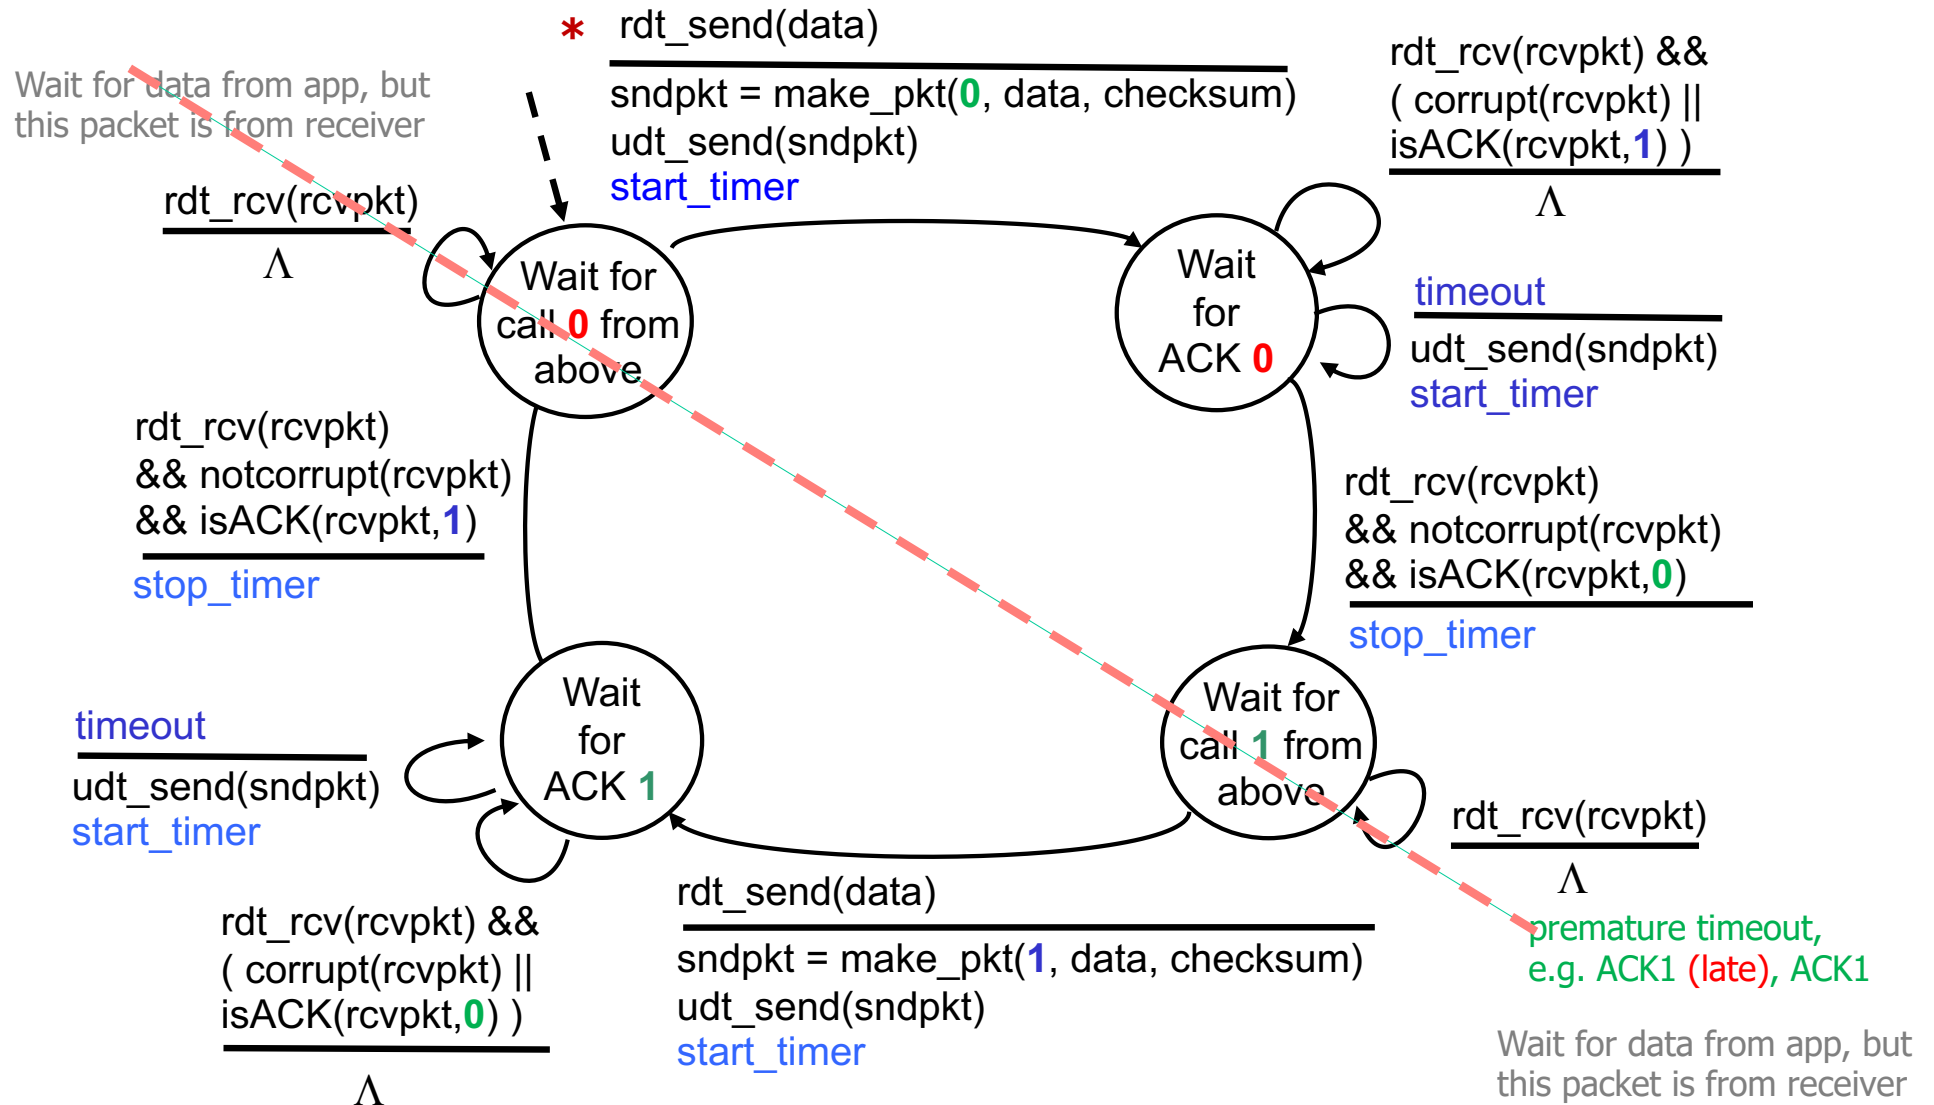
underlying channel can also lose packets (data, ACKs)

- Checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

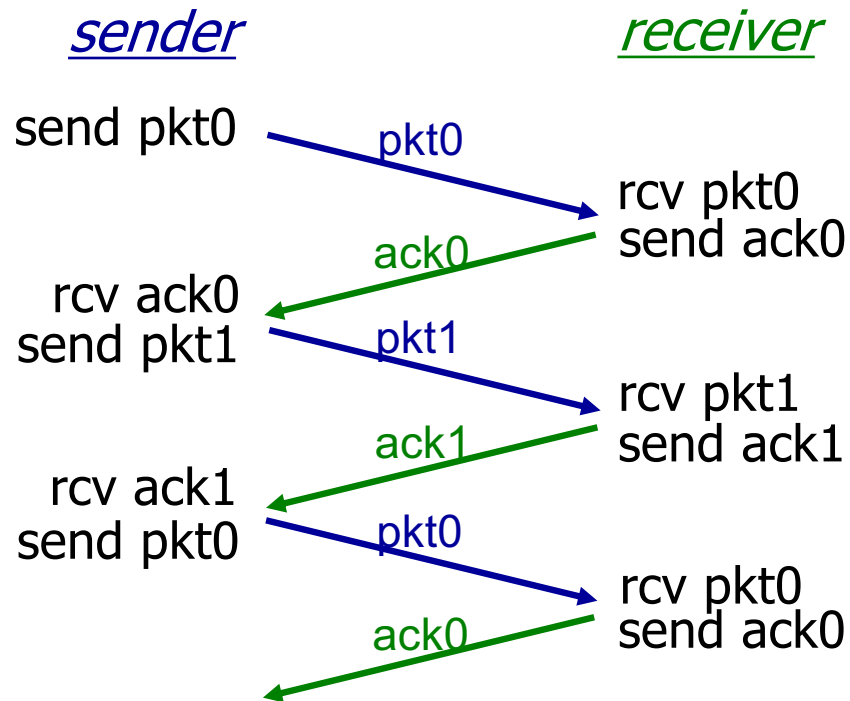
Approach: sender waits “reasonable” amount of **time** for ACK (round trip time + process time)

- ❖ Retransmits if no ACK received in this time
- ❖ If pkt (or ACK) just delayed (not lost)
 - Retransmission will be duplicate, but seq. #'s already handles this
 - Receiver must specify seq # of pkt being ACKed
- ❖ Requires countdown timer (RTT+ tolerate time)

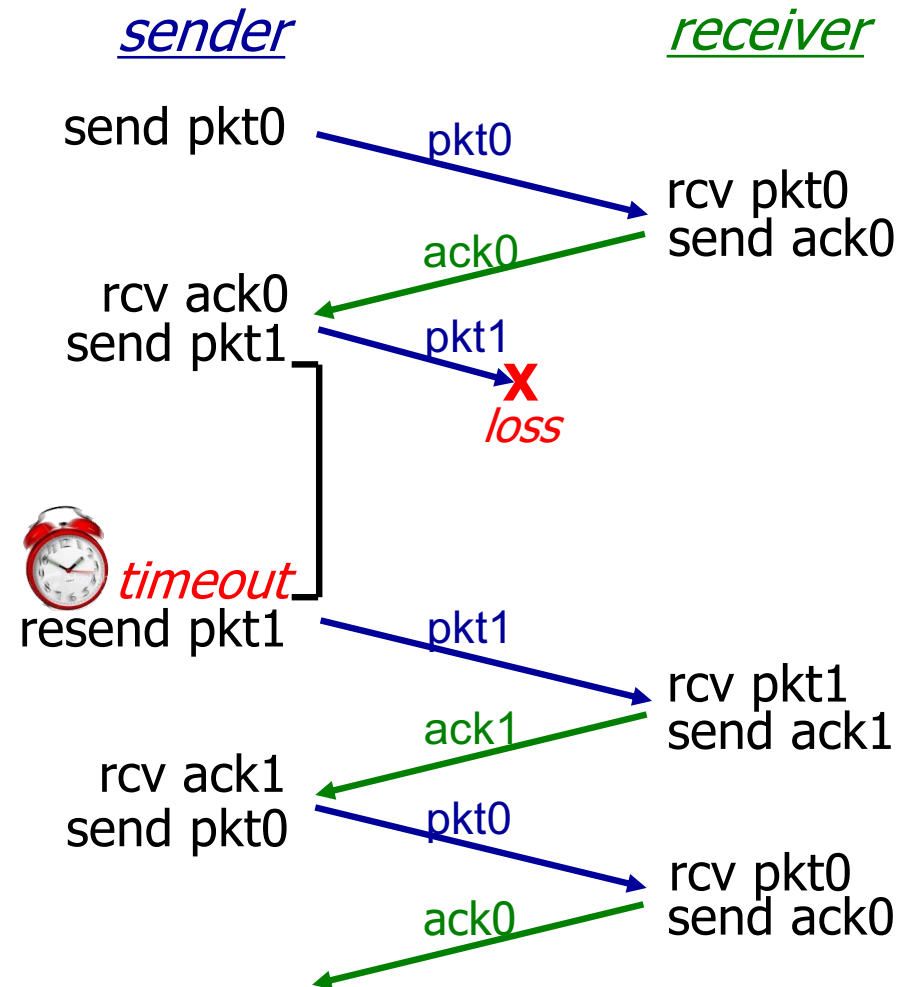
rdt3.0 sender



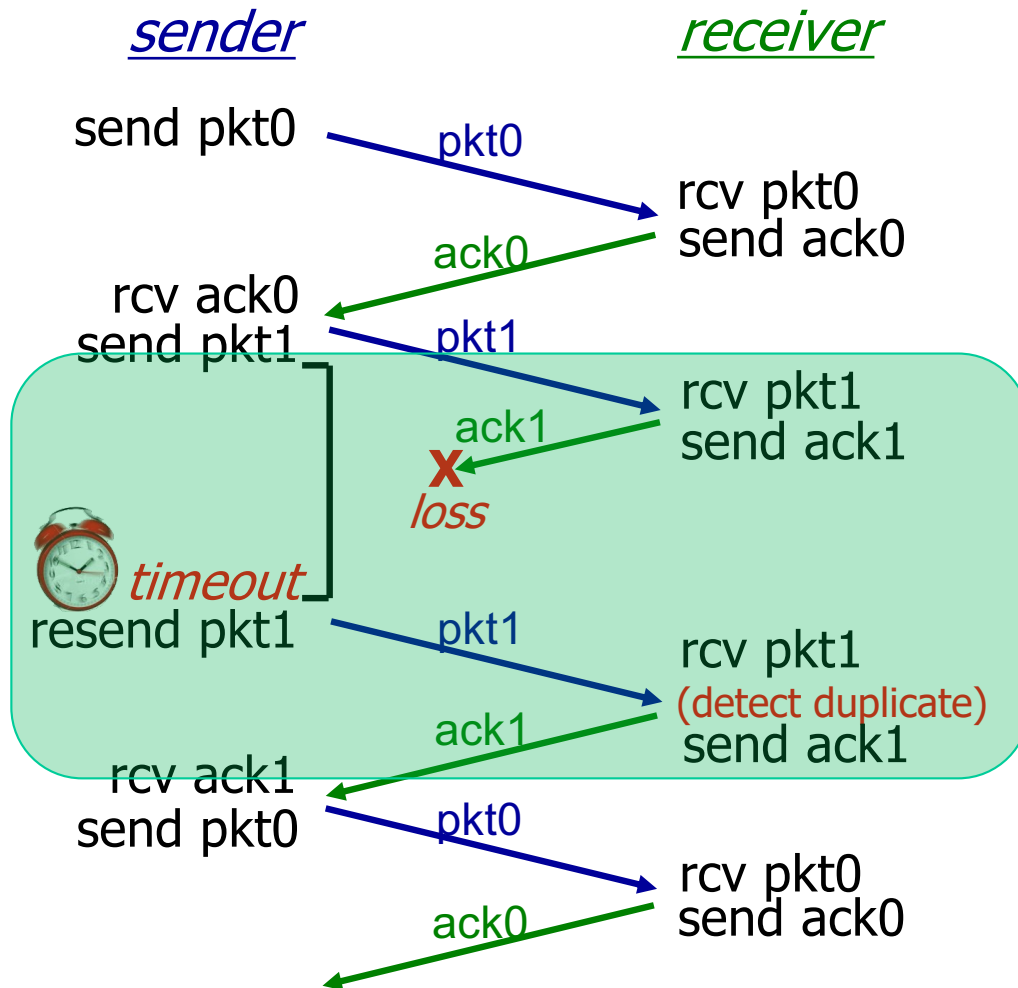
rdt3.0 in action



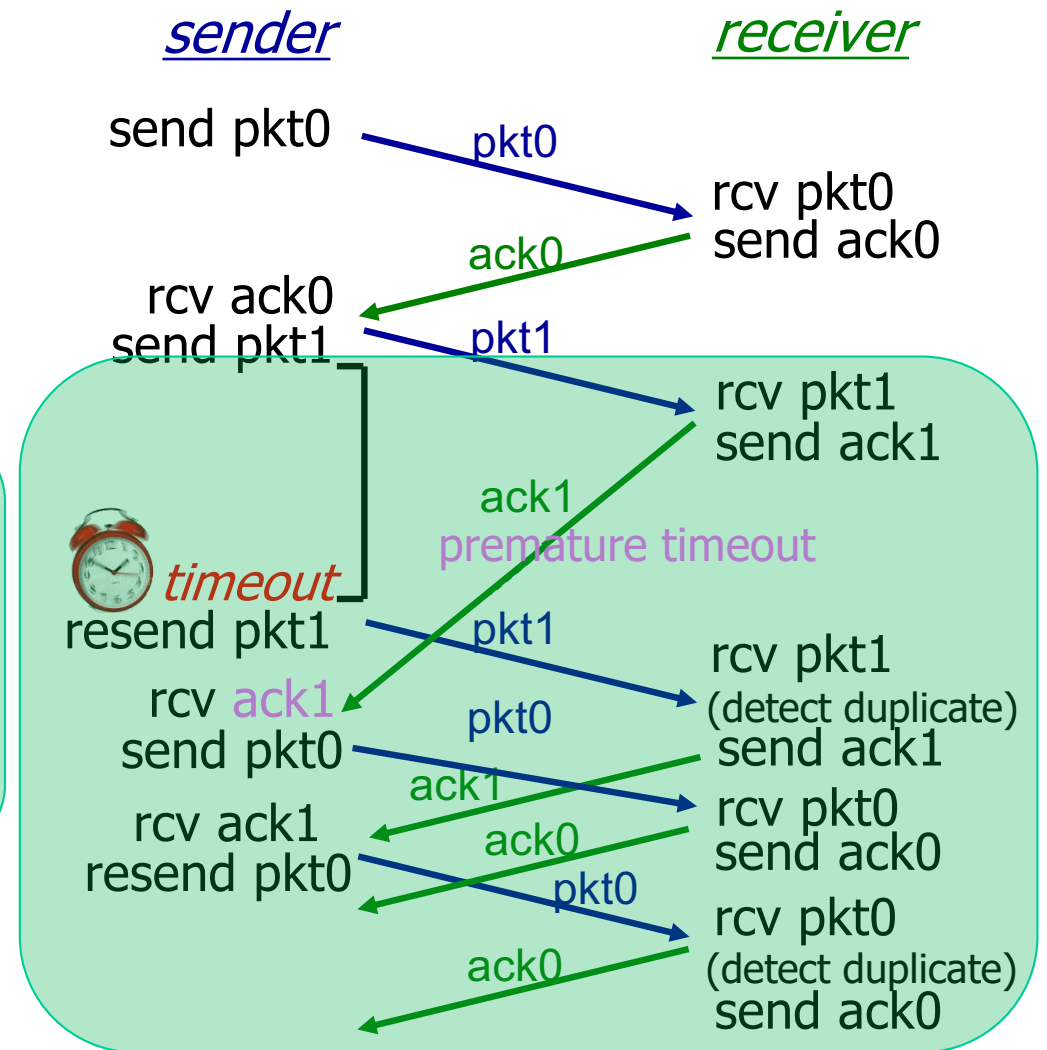
(a) no loss



(b) packet loss



(c) ACK loss



(d) premature timeout/ delayed ACK

Performance of rdt3.0

- ❖ rdt3.0 is correct, but performance stinks
- ❖ e.g.: 1 Gbps link (R), 15 ms prop. delay, 8000 bit packet (L), $RTT=30$ msec :


$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs} \quad (0.008 \text{ msec})$$

- U_{sender} : **utilization** – fraction of time sender busy sending

$$U_{sender} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.000267$$

- If $RTT=30$ msec, 8000 bits every 30.008 msec equals to **267 kbps thruput over 1 Gbps link**
- ❖ Network protocol limits use of physical resources!

rdt3.0: stop-and-wait operation




$$U_{sender} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.000267$$

Pipelined protocols

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- Range of sequence numbers must be increased [more than 2]
- **Buffering** at sender and/or receiver



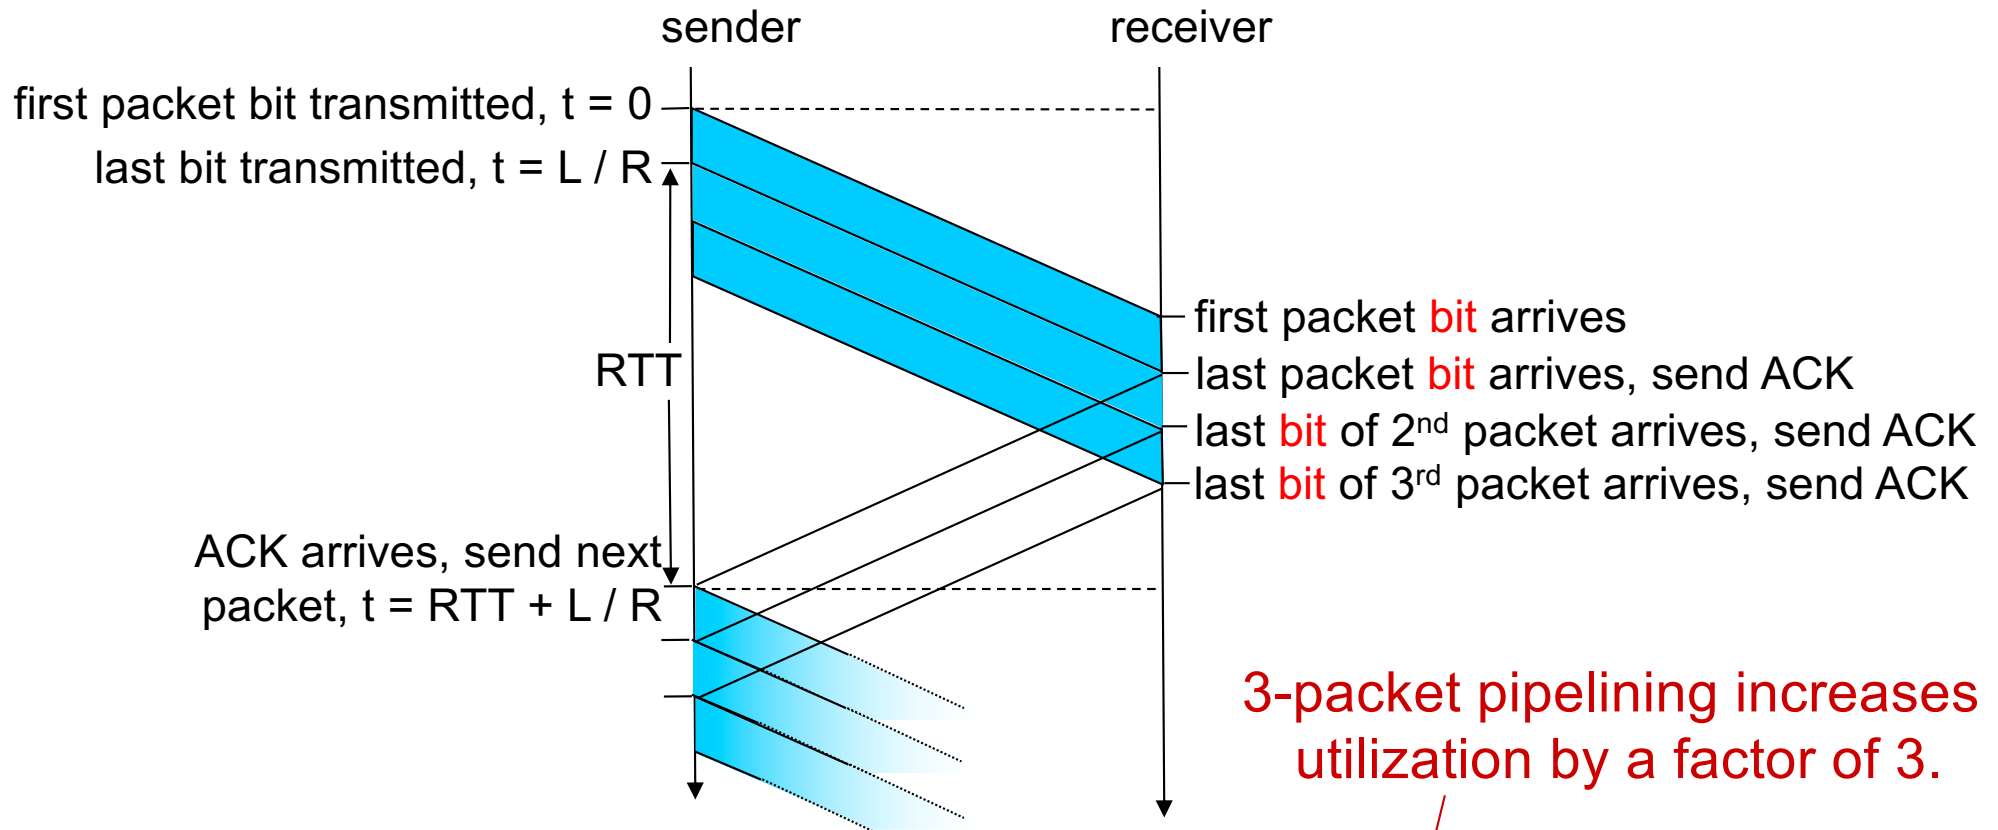
(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

❖ Two generic forms of pipelined protocols

- ❖ **Go-Back-N** (*one timer, timeout then resend all unacked pkts*)
- ❖ **Selective repeat** (*each timer per packet, timeout only resend the timeout pkt*)

Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Pipelined protocols: overview

Go-back-N


- ❖ Sender can have up to N unacked packets in pipeline
- ❖ Receiver only sends *cumulative ack*
 - Doesn't ack packet if there's a gap (e.g. 0, 1, 2, 4, 5)
- ❖ Sender has only one timer for the oldest unacked packet
 - When timer expires, retransmit all unacked packets

Selective Repeat

- ❖ Sender can have up to N unacked packets in pipeline
- ❖ Receiver sends *individual ack* for each packet
- ❖ Sender maintains timer for each unacked packet
 - When timer expires, retransmit only that unacked packet


Go-Back-N: sender

- ❖ K -bit seq # in pkt header (each pkt is assigned a seq# (n), length $k = \log_2 n$)
- ❖ “Window size” of up to N , consecutive unack’ed pkts allowed



- ❖ $ACK(n)$: ACKs **all pkts** up to, including seq # n - “**cumulative ACK**” (eg. $n=7$, $ACK(7)$ means pkts with seq#0, 1, 2, ... 7 are received)
 - May receive duplicate ACKs (see receiver) (eg. $ACK(5), ACK(5)$, ...pkt might be lost)
- ❖ Timer for the **oldest** in-flight pkt (only one timer)
- ❖ $Timeout(n)$: retransmit packet n and **all** higher seq # pkts (**yellow part**) in window

GBN: sender extended FSM



* Λ
 send_base=1
 nextseqnum=1


rdt_rcv(rcvpkt)
 && corrupt(rcvpkt)

drop

rdt_send(data)

```

if (nextseqnum < send_base+N) {
    sndpkt[nextseqnum] = make_pkt(nextseqnum,data,chksum)
    udt_send(sndpkt[nextseqnum])
    if (send_base == nextseqnum)
        start_timer (the oldest in-flight pkt)
    nextseqnum++
}
else
    refuse_data(data)
    
```



timeout

```

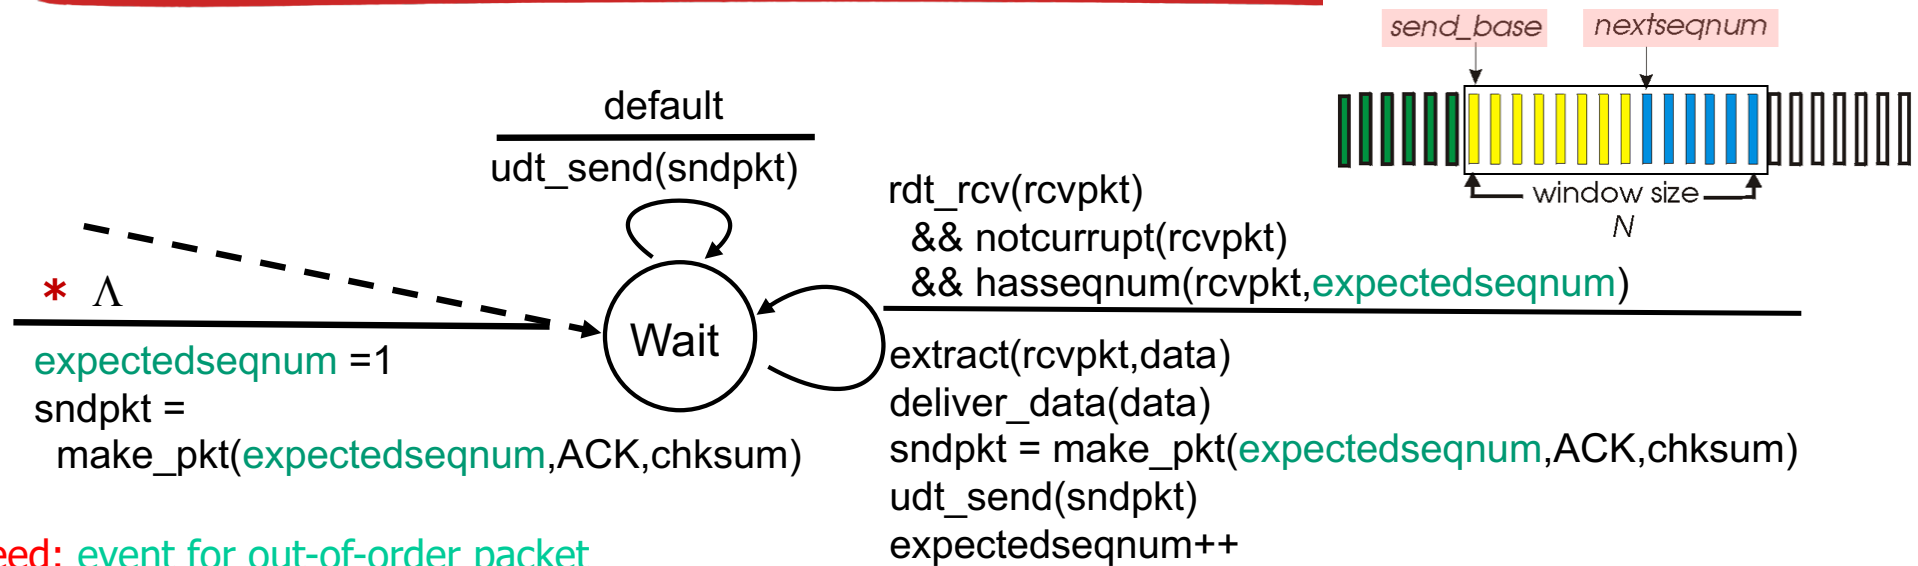
start_timer
udt_send(sndpkt[send_base])
udt_send(sndpkt[send_base+1])
...
udt_send(sndpkt[nextseqnum-1])
    
```

```

rdt_rcv(rcvpkt) &&
notcorrupt(rcvpkt)

send_base = getacknum(rcvpkt)+1
If (send_base == nextseqnum)
    stop_timer (all in-flight pkts arrived)
else
    start_timer (the oldest in-flight pkt)
    
```

GBN: receiver extended FSM



Need: event for out-of-order packet
(standard not specified)

ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- May generate duplicate ACKs (eg. ACK(5), ACK(5), ...pkt might be lost)
- Receiver needs only remember **expectedseqnum**

❖ Out-of-order pkt

- Discard (don't buffer): *no receiver buffering!*
- Re-ACK pkt with **highest** in-order seq #

GBN in action

sender window ($N=4$)

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

sender

send pkt0
 send pkt1
 send pkt2
 send pkt3
 (wait)

rcv ack0, send pkt4
 rcv ack1, send pkt5

ignore duplicate ACK



pkt 2 timeout

send pkt2
 send pkt3
 send pkt4
 send pkt5

receiver

rcv pkt0, deliver, send ack0
 rcv pkt1, deliver, send ack1

rcv pkt3, discard,
 (re)send ack1

rcv pkt4, discard,
 (re)send ack1


rcv pkt5, discard,
 (re)send ack1

rcv pkt2, deliver, send ack2
 rcv pkt3, deliver, send ack3
 rcv pkt4, deliver, send ack4
 rcv pkt5, deliver, send ack5


Selective repeat

- ❖ Receiver *individually* acknowledges all correctly received pkts
 - Buffer pkts, as needed, for eventual in-order delivery to upper layer
- ❖ Sender only resends pkts for which ACK not received
 - Sender timer for each unACKed pkt
- ❖ Sender window
 - N consecutive seq #'s
 - Limits seq #'s of sent, unACKed pkts

Selective repeat: sender, receiver windows



Selective repeat



Sender

Data from above

- ❖ If next available seq # in window, send pkt

Timeout(n)

- ❖ Resend pkt n , restart timer

ACK(n) in $[send_base, send_base+N]$

- ❖ Mark pkt n as received
- ❖ [Sliding sender window] if n is the **smallest unACKed** pkt, advance window base to **next unACKed** seq #

Receiver

Rceive pkt n in $[rcv_base, rcv_base+N-1]$

- ❖ Send ACK(n)
- ❖ Out-of-order: buffer
- ❖ [Sliding receiver window] in-order: deliver (also deliver buffered, in-order pkts), advance window base to **next not-yet-received** pkt

Pkt n in $[rcv_base-N, rcv_base-1]$

- ❖ ACK(n) [delay packet]

Otherwise

- ❖ Ignore

Selective repeat in action

sender window ($N=4$)

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

sender

send pkt0
 send pkt1
 send pkt2
 send pkt3
 (wait)

rcv ack0, send pkt4
 rcv ack1, send pkt5

record ack3 arrived



pkt 2 timeout

send pkt2

record ack4 arrived

record ack5 arrived

receiver

rcv pkt0, deliver, send ack0
 rcv pkt1, deliver, send ack1

rcv pkt3, **buffer**, send ack3

rcv pkt4, **buffer**, send ack4

rcv pkt5, **buffer**, send ack5

rcv pkt2; **deliver** pkt2,
 pkt3, pkt4, pkt5; send ack2

Q: what happens when ack2 arrives?


Selective repeat: dilemma (win size vs. seq. no.)

Example:


- ❖ Seq #'s: 0, 1, 2, 3
- ❖ Window size=3
- ❖ Receiver sees no difference in two scenarios!
- ❖ **Take retransmitted data** accepted as **new** data in (b)

Q: What relationship between **seq # size** and **window size** to avoid problem in (b)?

sequence number space $\geq 2 \times$ window size



*receiver can't see sender side.
receiver behavior identical in both cases!
something's (very) wrong!*



this is old *pkt0*, which is not expected by receiver (new *pkt0*)

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

❖ Point-to-point

- One sender, one receiver

❖ Reliable, in-order byte stream

- No “message boundaries” (doesn’t matter how many packets)
- Byte stream based, each byte is assigned a seq#

❖ Pipelined

- TCP congestion and flow control set window size (bytes)

❖ Send & receive buffers (selective repeat)

❖ Full duplex data

- Bi-directional data flow in same connection
- MSS: maximum segment size (TCP transmission data unit) – MSS should be determined at the time the connection is setup

❖ Connection-oriented

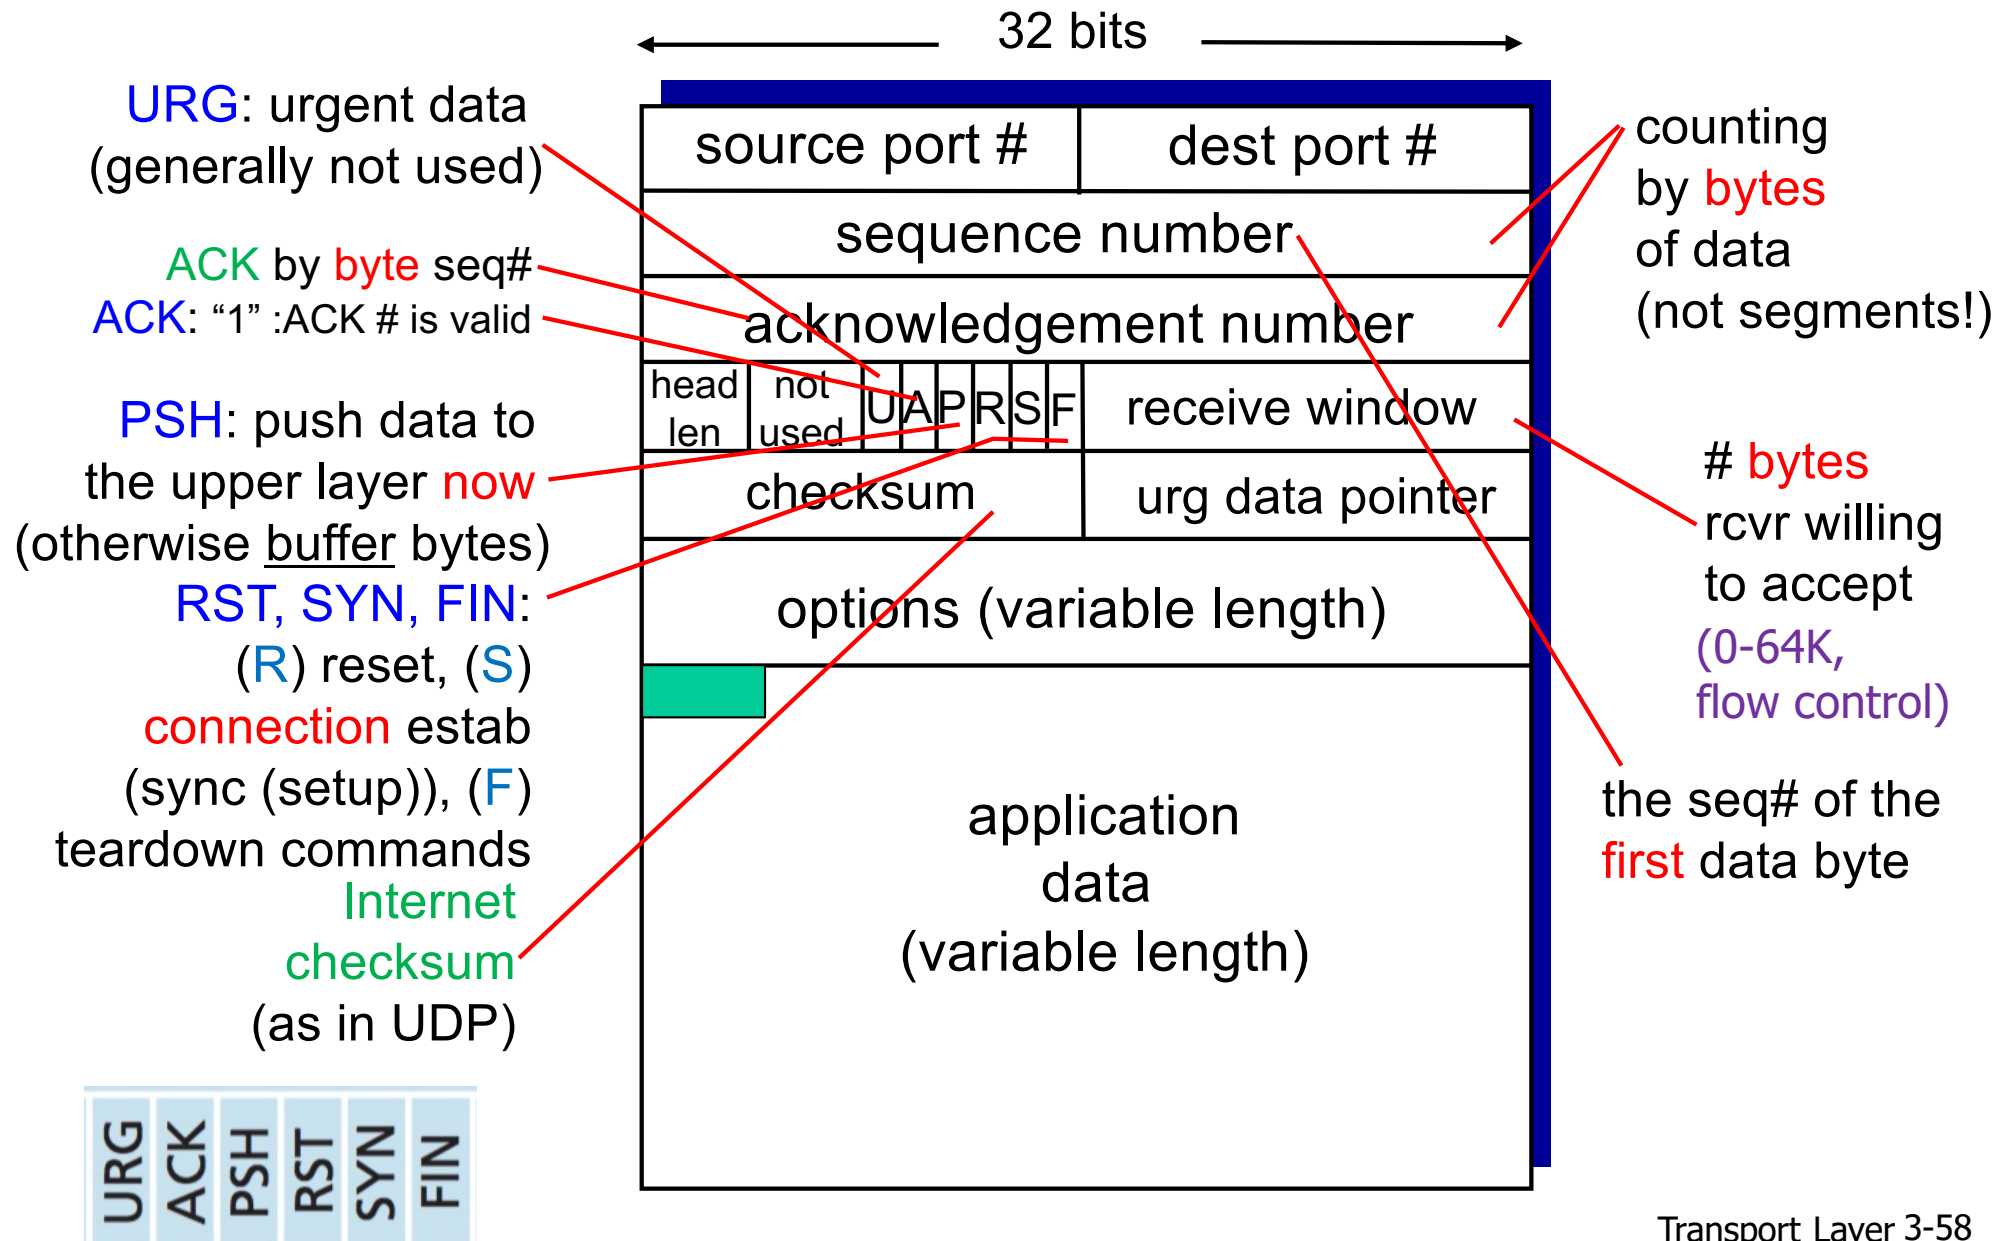
- Handshaking (exchange of control msgs) initializes sender, receiver state before data exchange

❖ Flow controlled


- Sender will not overwhelm receiver

TCP segment structure

No IP address fields
(Network Layer header)



TCP send and receive buffers



TCP seq. numbers, ACKs

Sequence numbers

- Byte stream “number” of **first byte** in segment's data
- Each byte is associated with a seq #

Acknowledgements

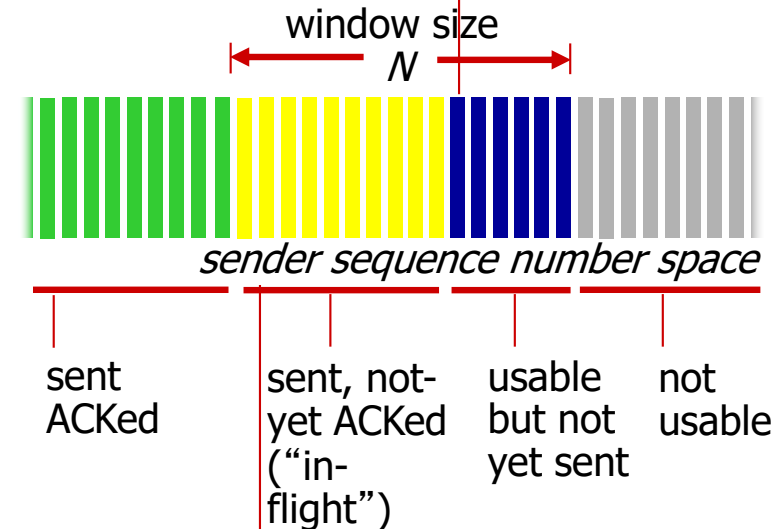
- Seq # of next byte expected from other side
- Cumulative ACK (Go-Back-N)

Q: How receiver handles out-of-order segments

- A: TCP spec doesn't say, - up to implementer

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

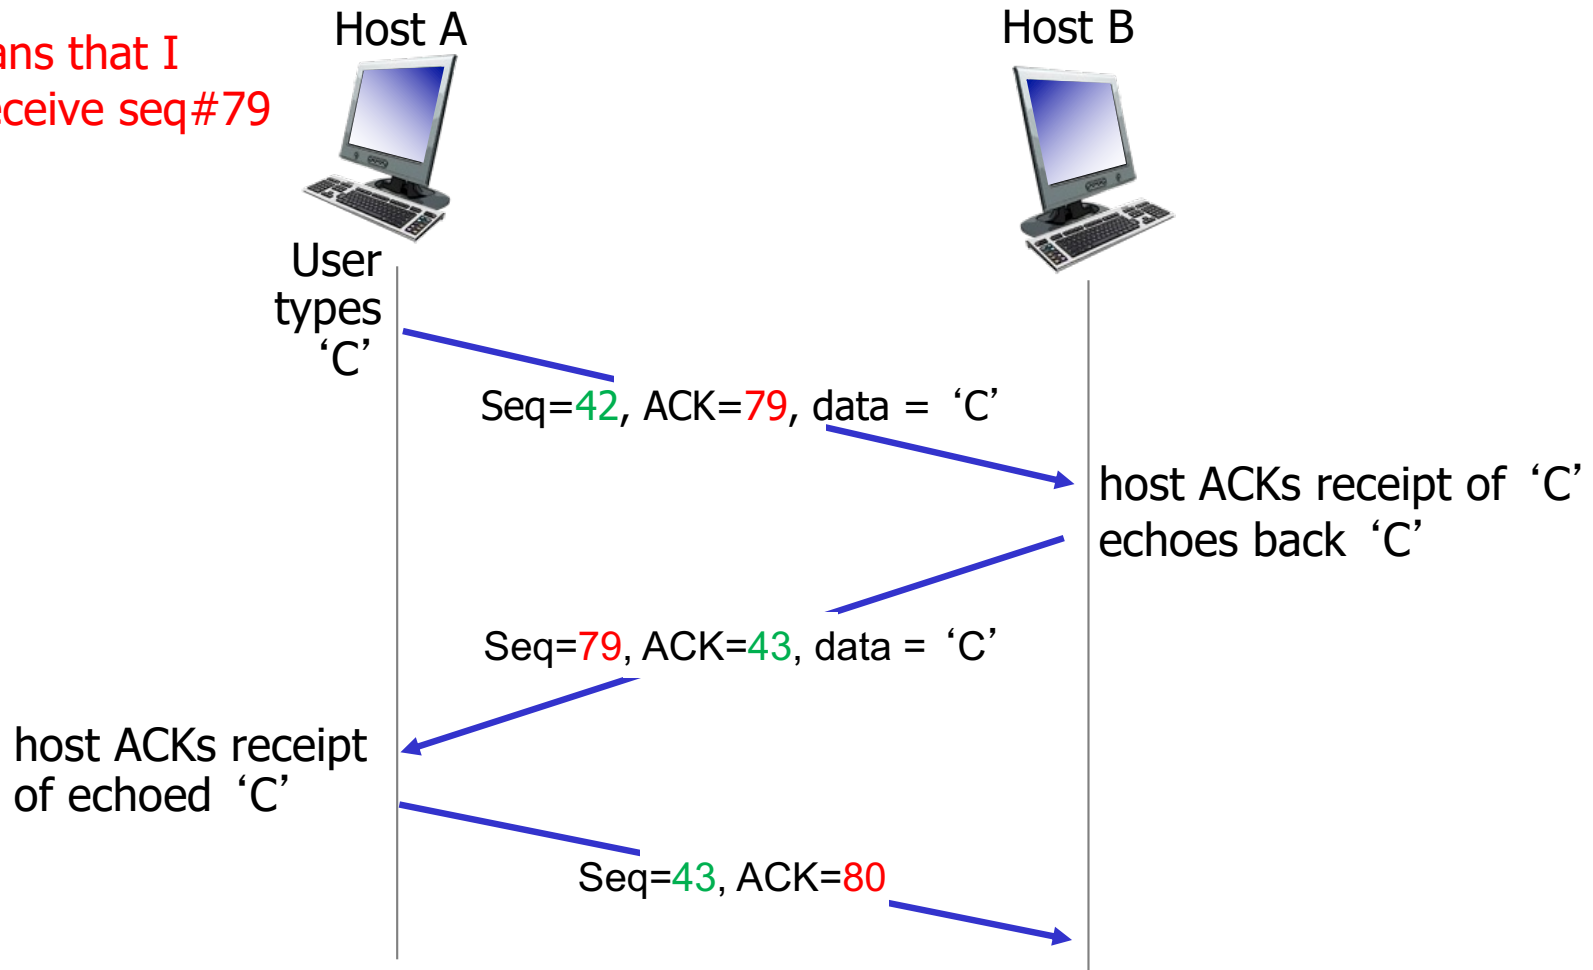


incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

TCP seq. numbers, ACKs

ACK-79 means that I expect to receive seq#79



simple telnet scenario

TCP round trip time, timeout

Q: How to set TCP timeout value?

- ❖ Longer than RTT
 - But RTT varies
- ❖ *Timer too short:* premature timeout, unnecessary retransmissions
- ❖ *Timer too long:* slow reaction to segment loss


Q: How to estimate RTT?

- ❖ **SampleRTT:** measured time from segment transmission until ACK receipt
 - Ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT “smoother”
 - Average several *recent* measurements, not just current **SampleRTT**

TCP round trip time, timeout

$$\text{EstimatedRTT} = (1-\alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ Exponential weighted moving average
- ❖ Influence of past sample decreases exponentially fast
- ❖ Typical value: $\alpha = 0.125$



TCP round trip time, timeout

- ❖ Timeout interval: EstimatedRTT plus “safety margin”
 - Large variation in EstimatedRTT → larger safety margin
- ❖ Estimate SampleRTT deviation from EstimatedRTT:

$$\begin{aligned}\text{DevRTT} = & (1-\beta) * \text{DevRTT} + \\ & \beta * |\text{SampleRTT} - \text{EstimatedRTT}| \\ & (\text{typically, } \beta = 0.25)\end{aligned}$$

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑
estimated RTT

↑
“safety margin”

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

TCP reliable data transfer

- ❖ TCP creates rdt service on top of IP's unreliable service

- Pipelined segments [not stop and wait]
- Cumulative acks [ack(n) is “expecting” pkt n]
- Single retransmission timer [retransmit **only one pkt**]

- ❖ Retransmissions triggered by

- Timeout events
- Duplicate acks [pkts may be lost]

Let's initially consider simplified TCP sender

- Ignore duplicate acks
- Ignore flow control, congestion control

TCP sender events

Data rcvd from app

- ❖ Create segment with seq #
- ❖ Seq # is the byte-stream number of the first data byte in segment
- ❖ Start timer if not already running (only one timer)
 - Think of timer as for oldest unacked segment
 - Expiration interval: **TimeoutInterval**

Timeout

- ❖ Retransmit the oldest unacked segment
- ❖ Restart timer

Ack rcvd

- ❖ If ack acknowledges previously unacked segments
 - Update what is known to be ACKed
 - Start timer if there are still unacked segments

```
/* Assume sender is not constrained by TCP flow or congestion control, that data from above is less than MSS in size, and that data transfer is in one direction only. */
```

```
NextSeqNum=InitialSeqNumber
```

```
SendBase=InitialSeqNumber
```

Simplified TCP sender

```
loop (forever) {  
    switch(event)
```

```
    event: data received from application above
```

```
        create TCP segment with sequence number NextSeqNum
```

```
        if (timer currently not running)
```

```
            start timer
```

```
        pass segment to IP
```

```
        NextSeqNum=NextSeqNum+length(data)
```

```
        break;
```

```
    event: timer timeout
```

```
        retransmit not-yet-acknowledged segment with
```

```
            smallest sequence number
```

```
        start timer
```

```
        break;
```

```
    event: ACK received, with ACK field value of y
```

```
        if (y > SendBase) {
```

```
            SendBase=y
```

```
            if (there are currently any not-yet-acknowledged segments)
```

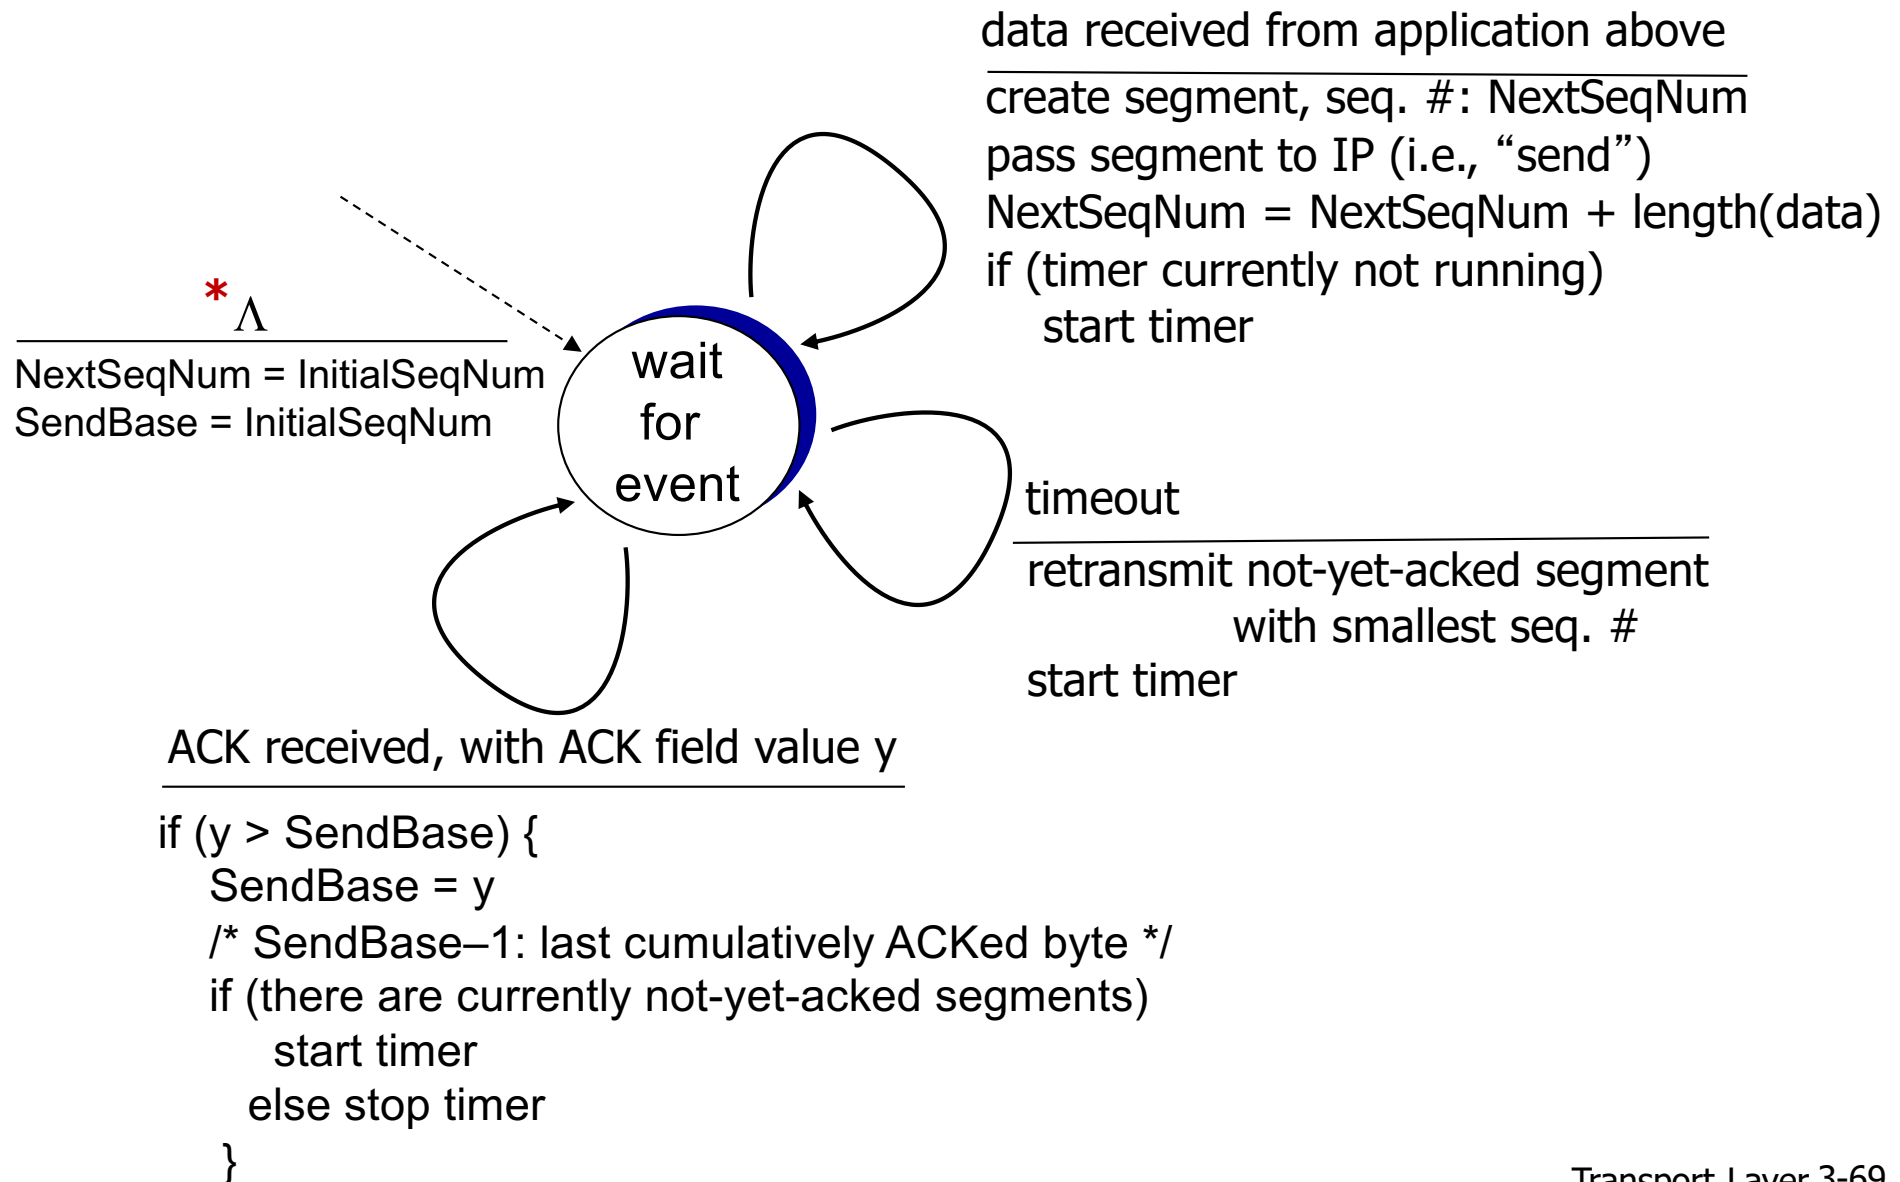
```
                start timer
```

```
        }
```

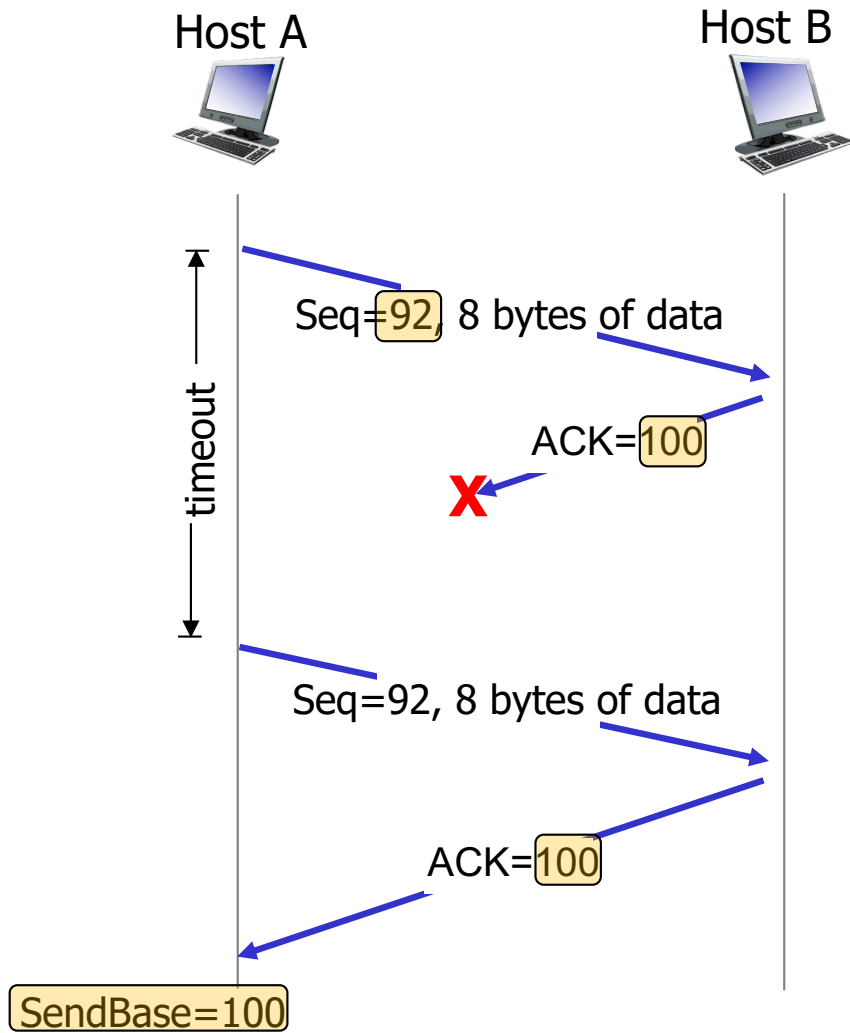
```
        break;
```

```
    } /* end of loop forever */
```

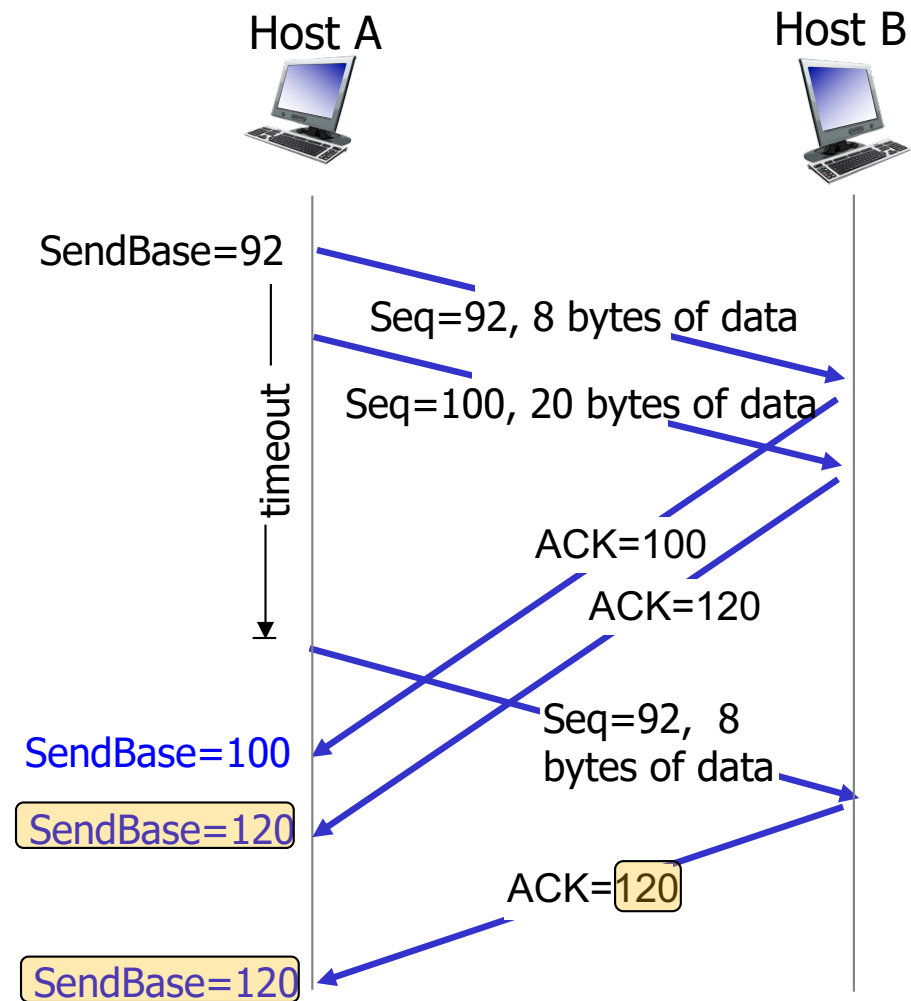
TCP sender (simplified)



TCP: retransmission scenarios

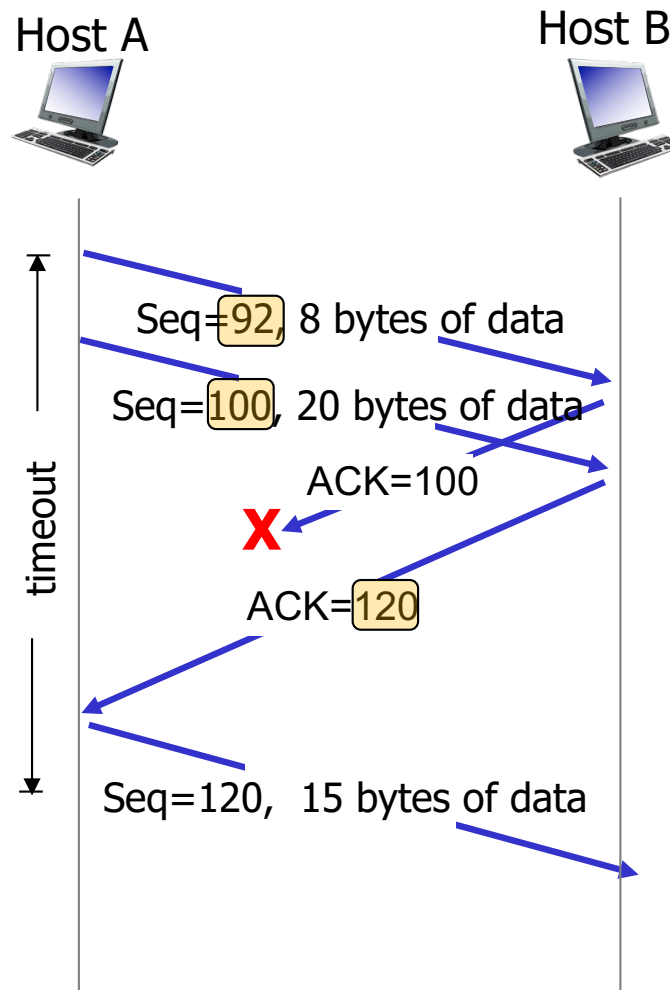


lost ACK scenario



premature timeout

TCP: retransmission scenarios



cumulative ACK

TCP ACK generation [RFC 1122, RFC 2581]

event at receiver

TCP receiver action

Arrival of **in-order** segment with expected seq #. All data up to expected seq # already ACKed.

Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK

Arrival of **in-order** segment with expected seq #. One other segment has ACK **pending**. [delayed ACK]

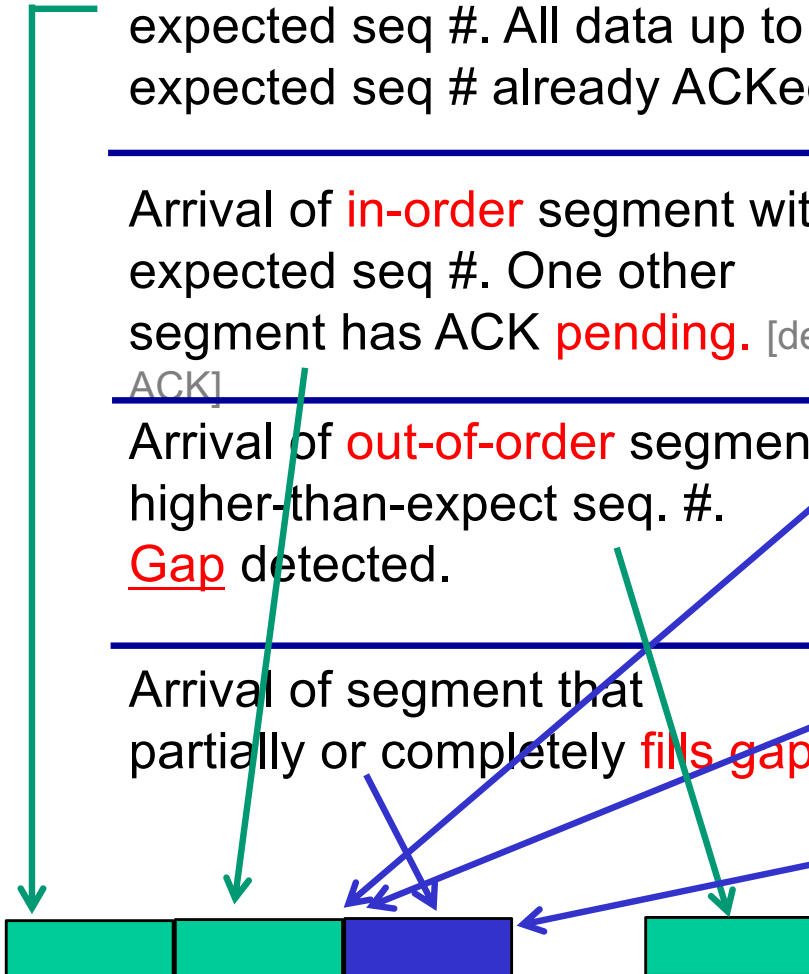
Immediately send **single cumulative ACK**, ACKing both in-order segments

Arrival of **out-of-order** segment higher-than-expect seq. #. **Gap** detected.

Immediately send **duplicate ACK**, indicating seq. # of next expected byte [imply pkts lost]

Arrival of segment that partially or completely **fills gap**.

Immediate send **ACK**, provided that segment starts at lower end of gap



TCP fast retransmit

- ❖ Time-out period often relatively long
 - Long delay before resending lost packet
- ❖ Detect lost segments via duplicate ACKs
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs

TCP fast retransmit

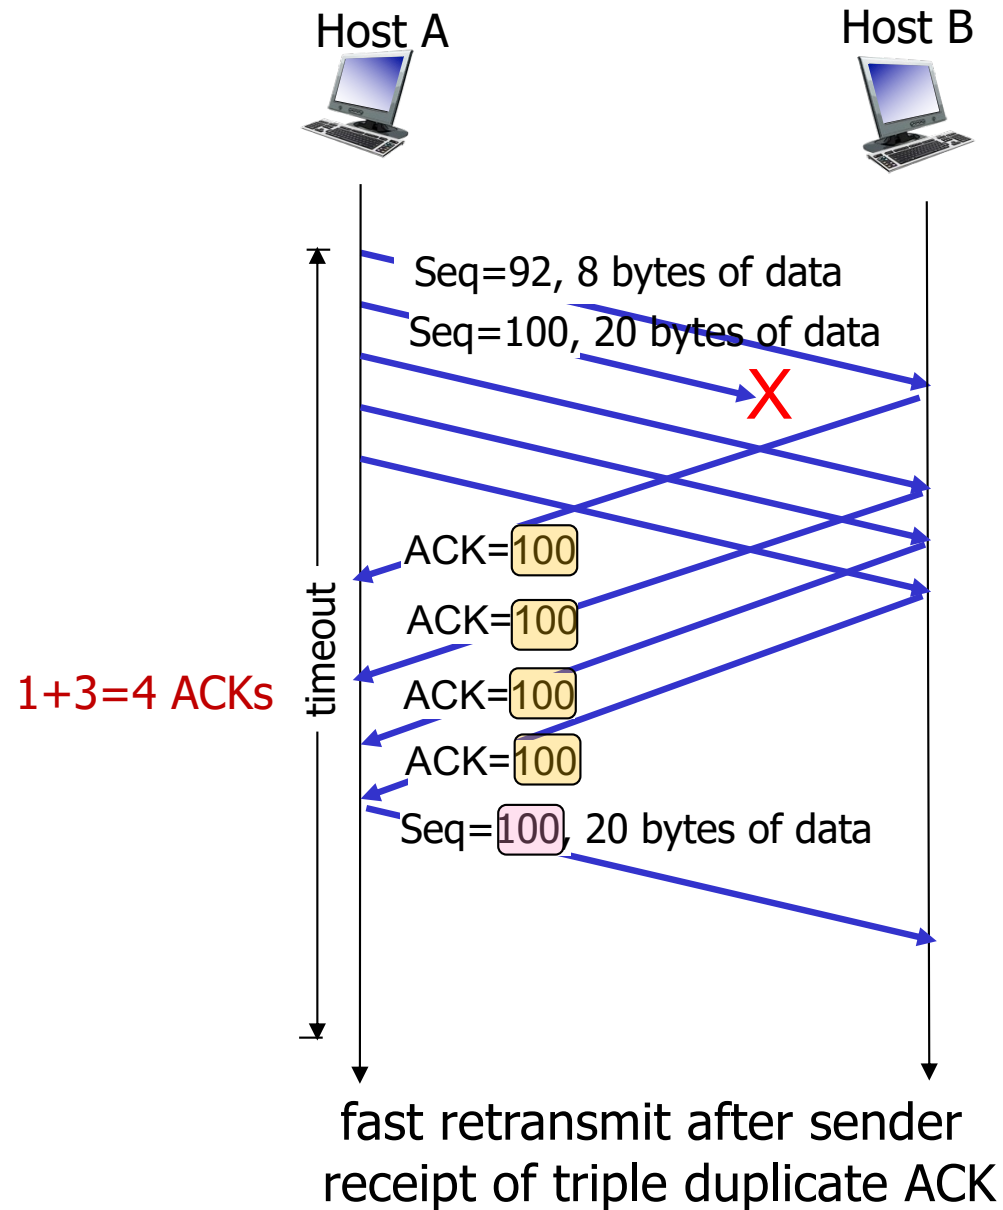
If sender receives **3 ACKs** for same data (“triple duplicate ACKs” - (1+3=4 ACKs)), resend unacked segment with smallest seq #

- Likely that unacked segment lost, so don't wait for timeout

Fast retransmit algorithm

```
event: ACK received, with ACK field value of y
    if (y > SendBase) {
        SendBase=y
        if (there are currently any not yet
            acknowledged segments)
            start timer
    }
else { /* a duplicate ACK for already ACKed
        segment */
    increment number of duplicate ACKs
        received for y
    if (number of duplicate ACKS received
        for y==3)
        /* TCP fast retransmit */
        resend segment with sequence number y
    }
break;
```


TCP fast retransmit



Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

Flow control

Receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast

Flow control

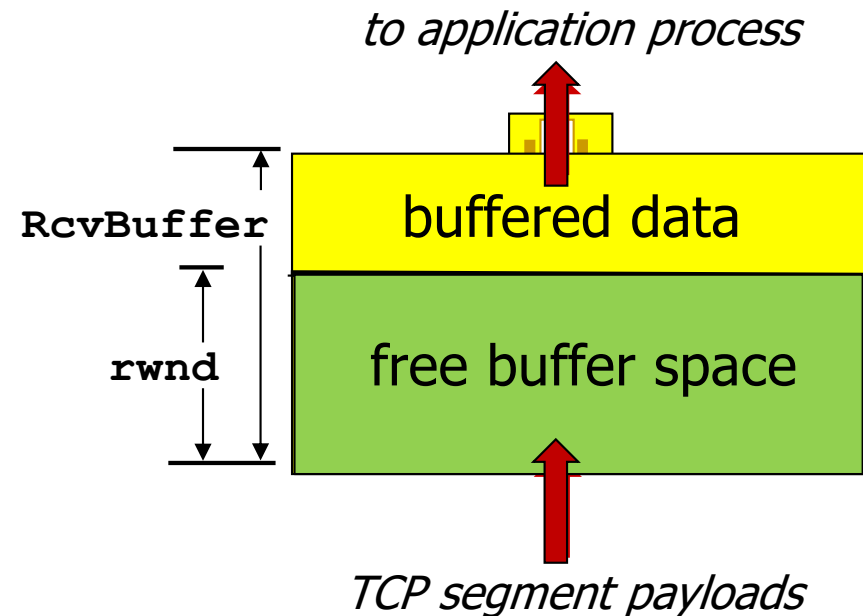
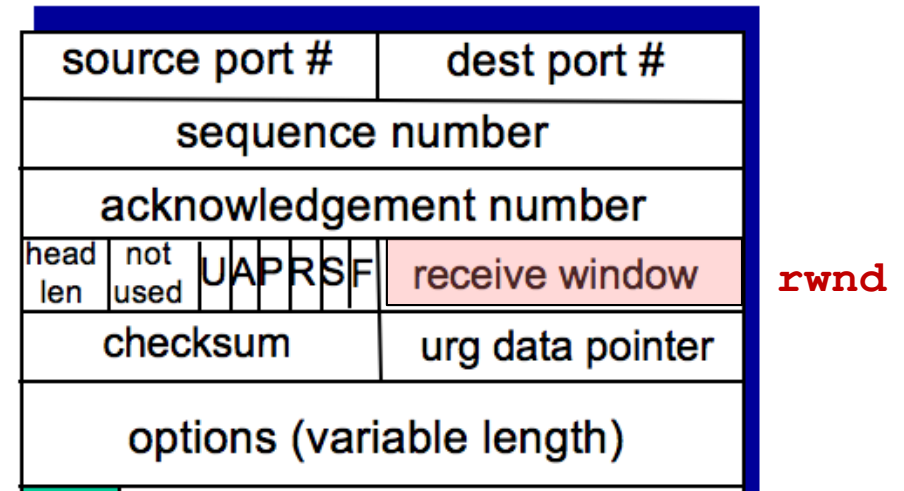
Receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



TCP flow control

- ❖ Receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
 - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
 - Many operating systems autoadjust **RcvBuffer**
- ❖ Sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- ❖ Guarantees receive buffer will not overflow

receiver segment



receiver-side buffering

Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer


3.5 Connection-oriented transport: TCP

- Segment structure
- Reliable data transfer
- Flow control
- Connection management

3.6 Principles of congestion control

3.7 TCP congestion control

TCP Connection Management




Recall: TCP sender, receiver establish “connection” before exchanging data segments

- ❖ Initialize TCP variables
 - Seq #s
 - Buffers, flow control info (e.g. RecWindow - rwnd)
- ❖ **Client:** connection initiator
`Socket clientSocket = new Socket(“hostname”, “port number”);`
- ❖ **Server:** connected by client
`Socket connectionSocket = welcomeSocket.accept();`

Three way handshake

Step 1: **client** host sends TCP **SYN** segment to server (set SYN = 1)

- Specifies initial seq#
- No data



Step 2: **server** host receives SYN, replies with **SYNACK** (set SYN = 1, ACK = 1)

- Server allocates buffers [DOS attack – SYN flooding]
- Specifies server initial seq#

Step 3: **client** receives SYNACK, replies with **ACK** segment, which may contain **data**




Connection Management

Before exchanging data, sender/receiver “handshake”

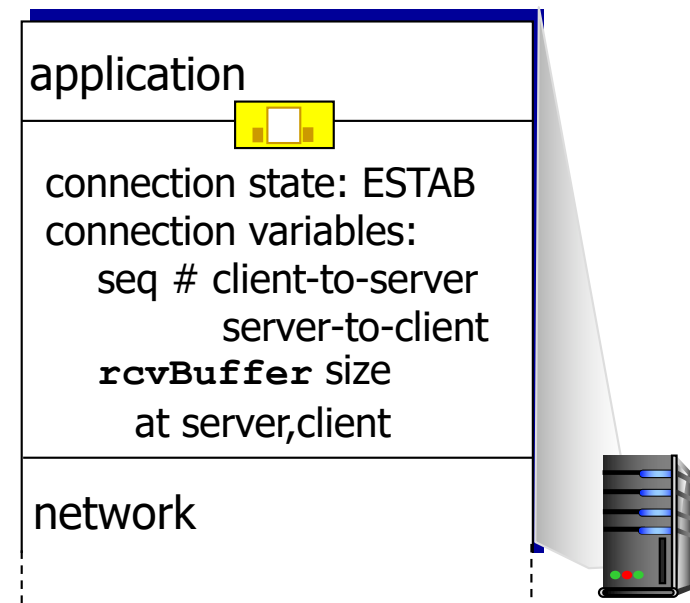
- ❖ Agree to establish connection (each knowing the other willing to establish connection)
- ❖ Agree on connection parameters

client



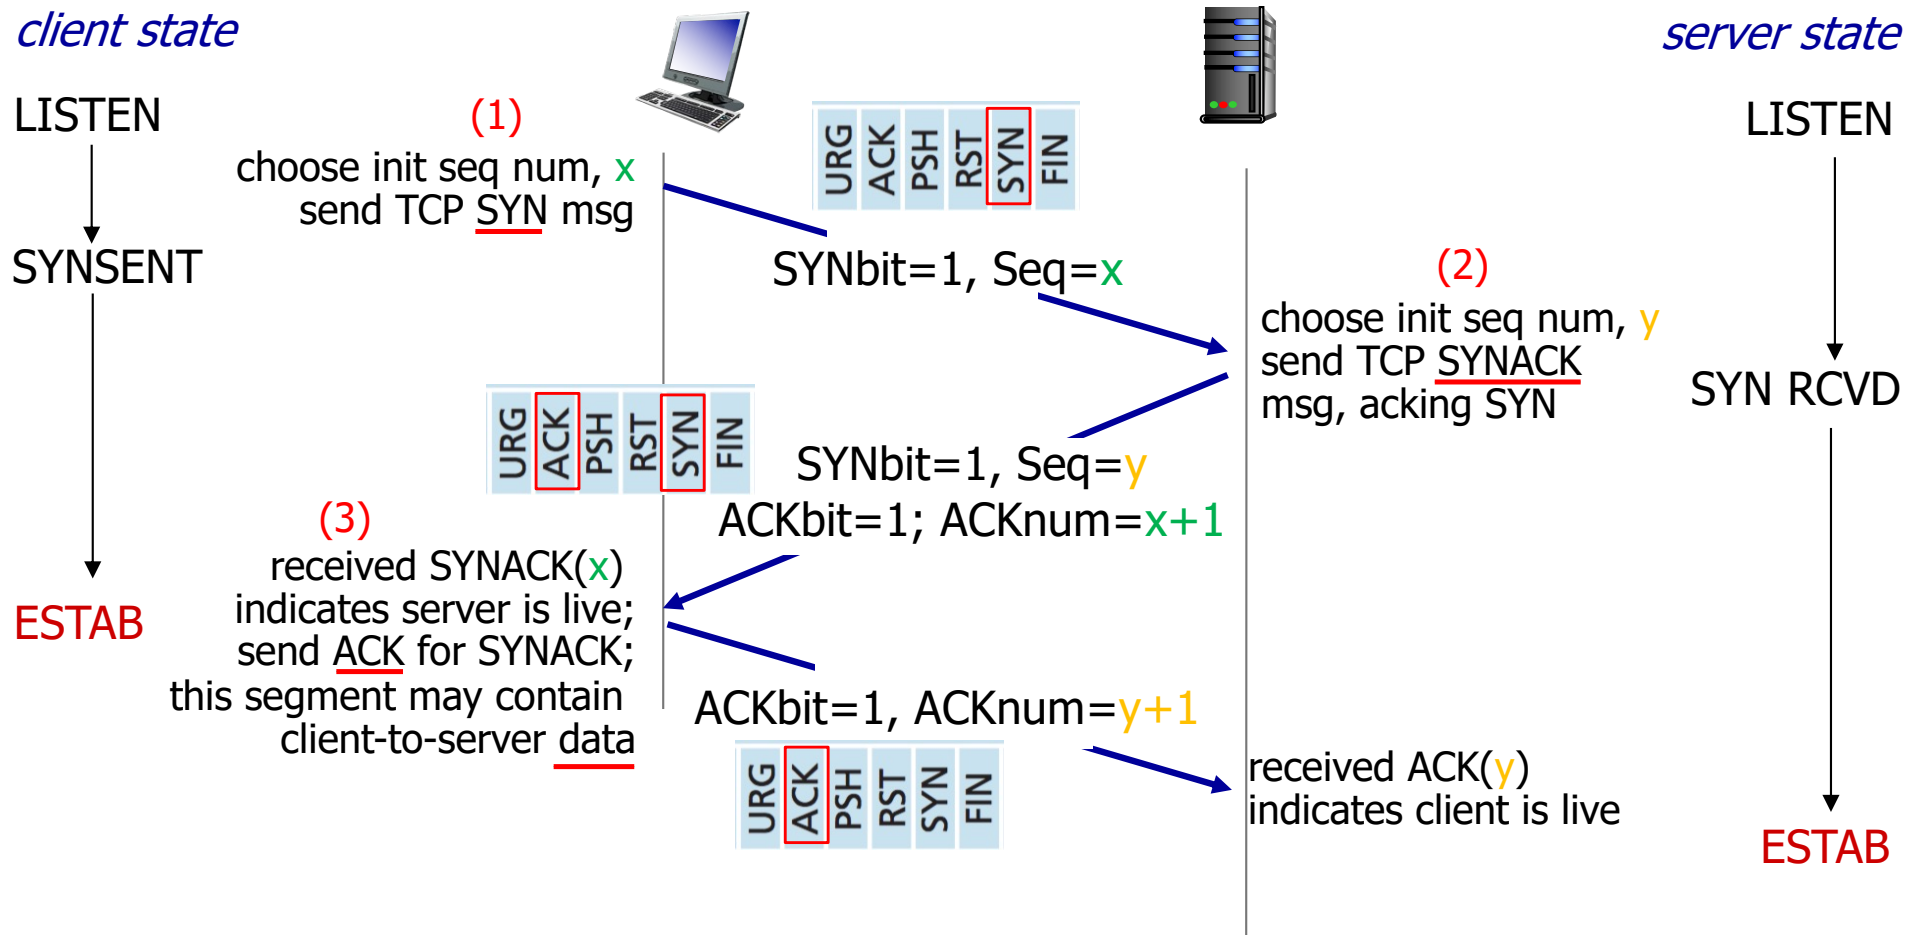
```
Socket clientSocket =  
    newSocket("hostname", "port  
    number");
```

server

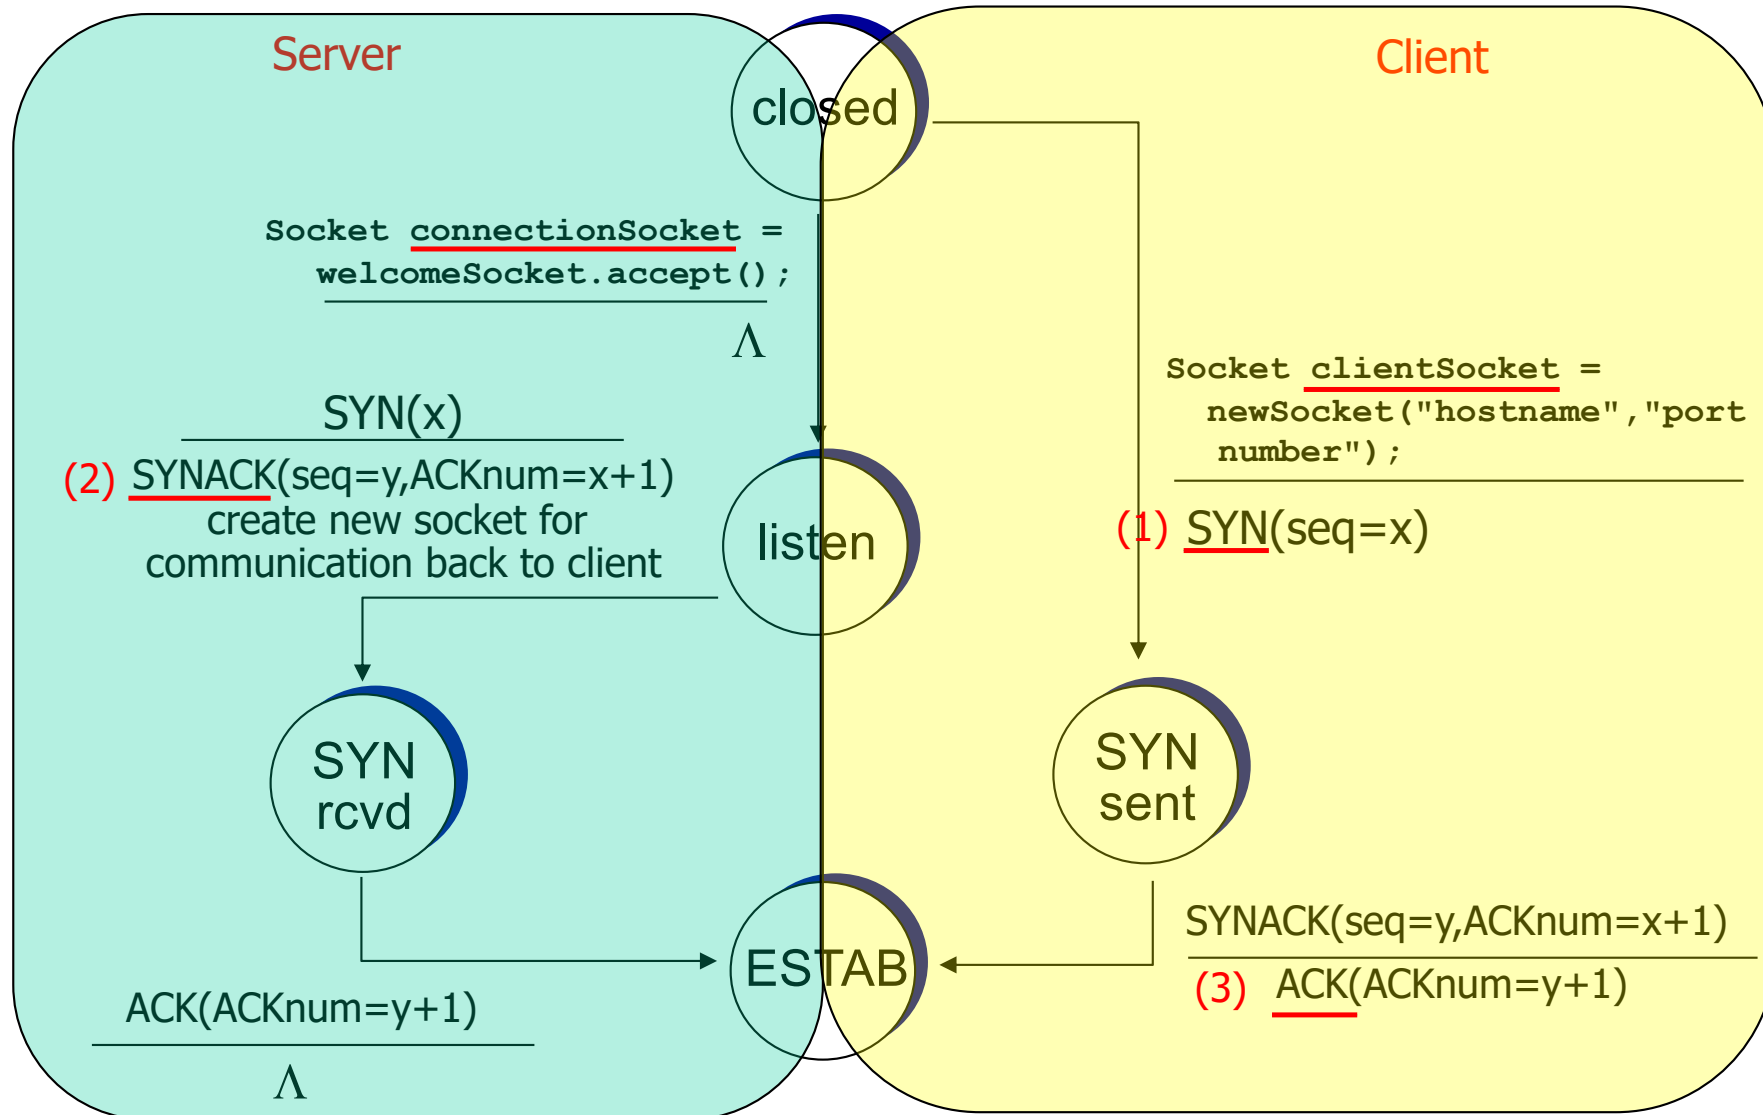


```
Socket connectionSocket =  
    welcomeSocket.accept();
```

TCP 3-way handshake



TCP 3-way handshake: FSM



TCP: closing a connection


- ❖ Client, server each close their side of **connection**
 - Send TCP segment with FIN bit = 1
- ❖ Respond to received FIN with ACK
 - On receiving FIN, ACK can be combined with own FIN
- ❖ **Closing a connection**

Client closes socket: `clientSocket.close();`

Step 1: **client** ends system
sends TCP FIN control segment to server



Step 2: **server** receives FIN
replies with ACK
closes connection



sends FIN



Step 3: client receives FIN
replies with ACK



enters “timed wait” – will respond with ACK to received FIN

Step 4 :server receives ACK
connection closed

TCP: closing a connection

client state

ESTAB

`clientSocket.close()`

FIN_WAIT_1

can no longer
send but can
receive data

FIN_WAIT_2

wait for server
close

TIMED_WAIT

timed wait
for $2 * \text{max}$
segment lifetime
(or 30ms)

CLOSED




FINbit=1, seq=x

ACKbit=1; ACKnum=x+1



FINbit=1, seq=y

ACKbit=1; ACKnum=y+1



server state

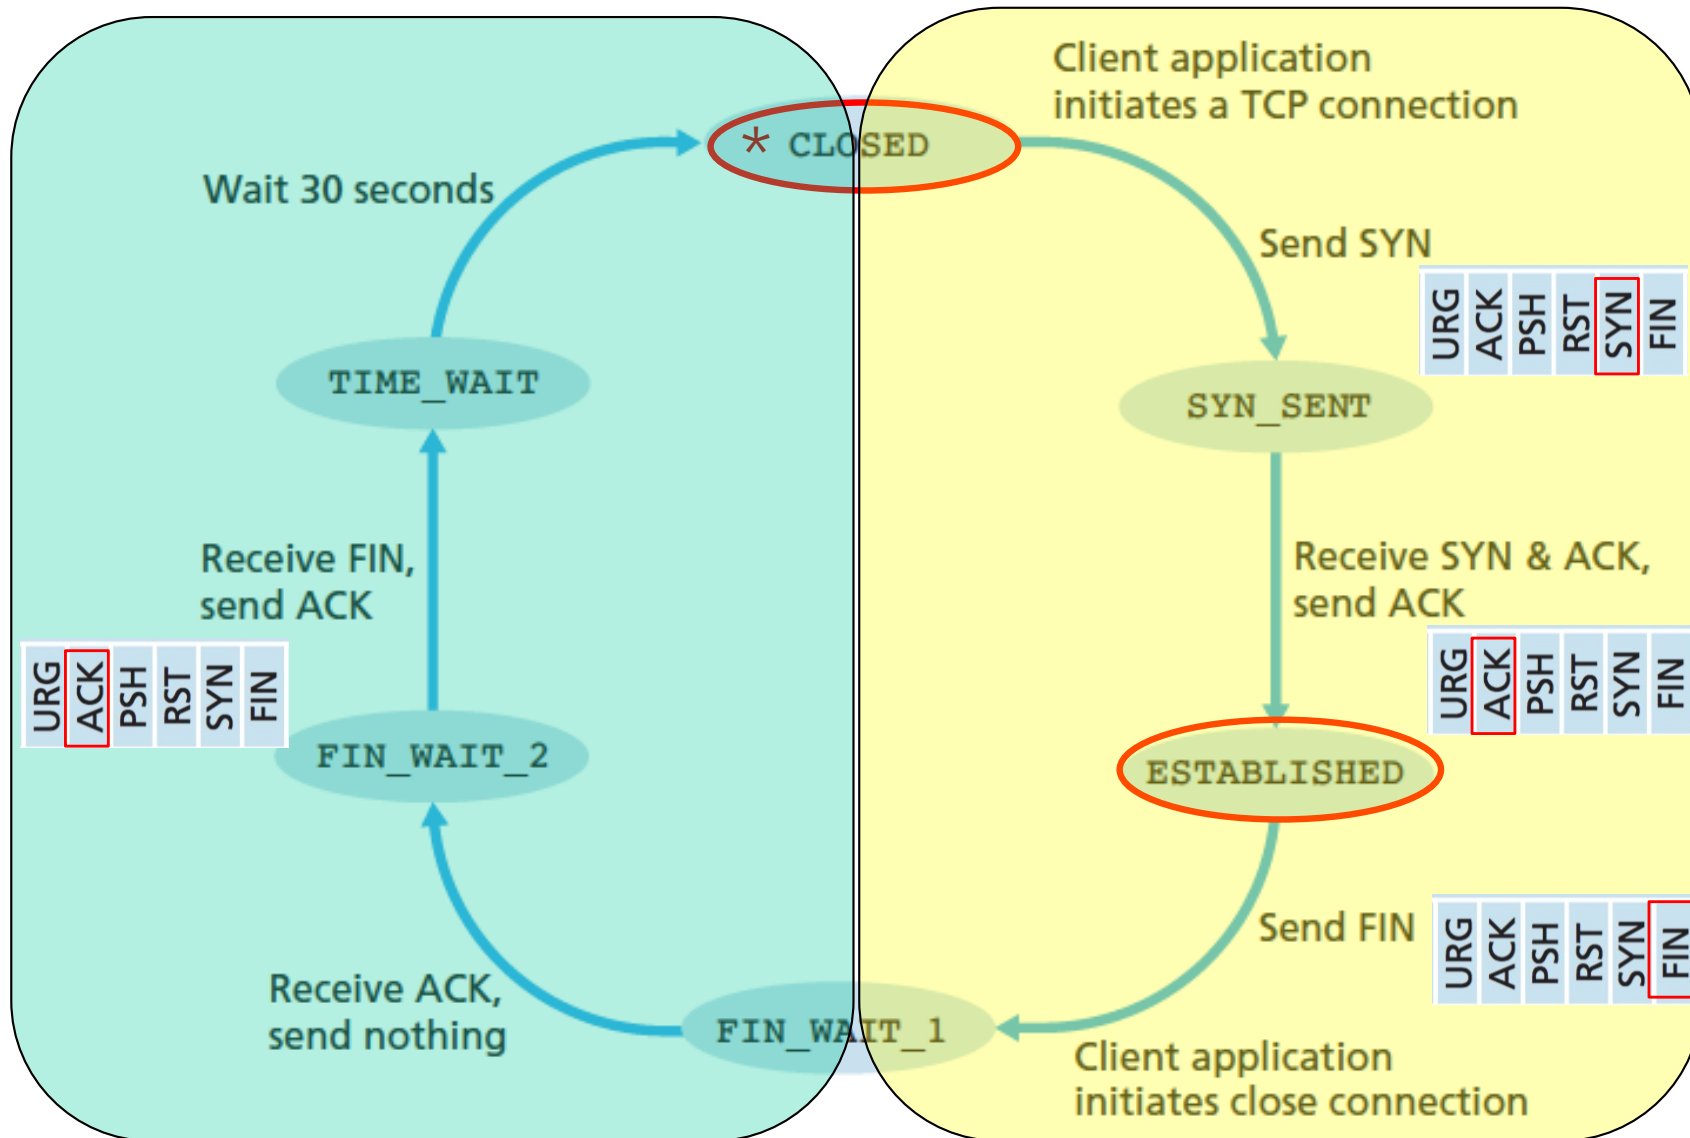
ESTAB

CLOSE_WAIT

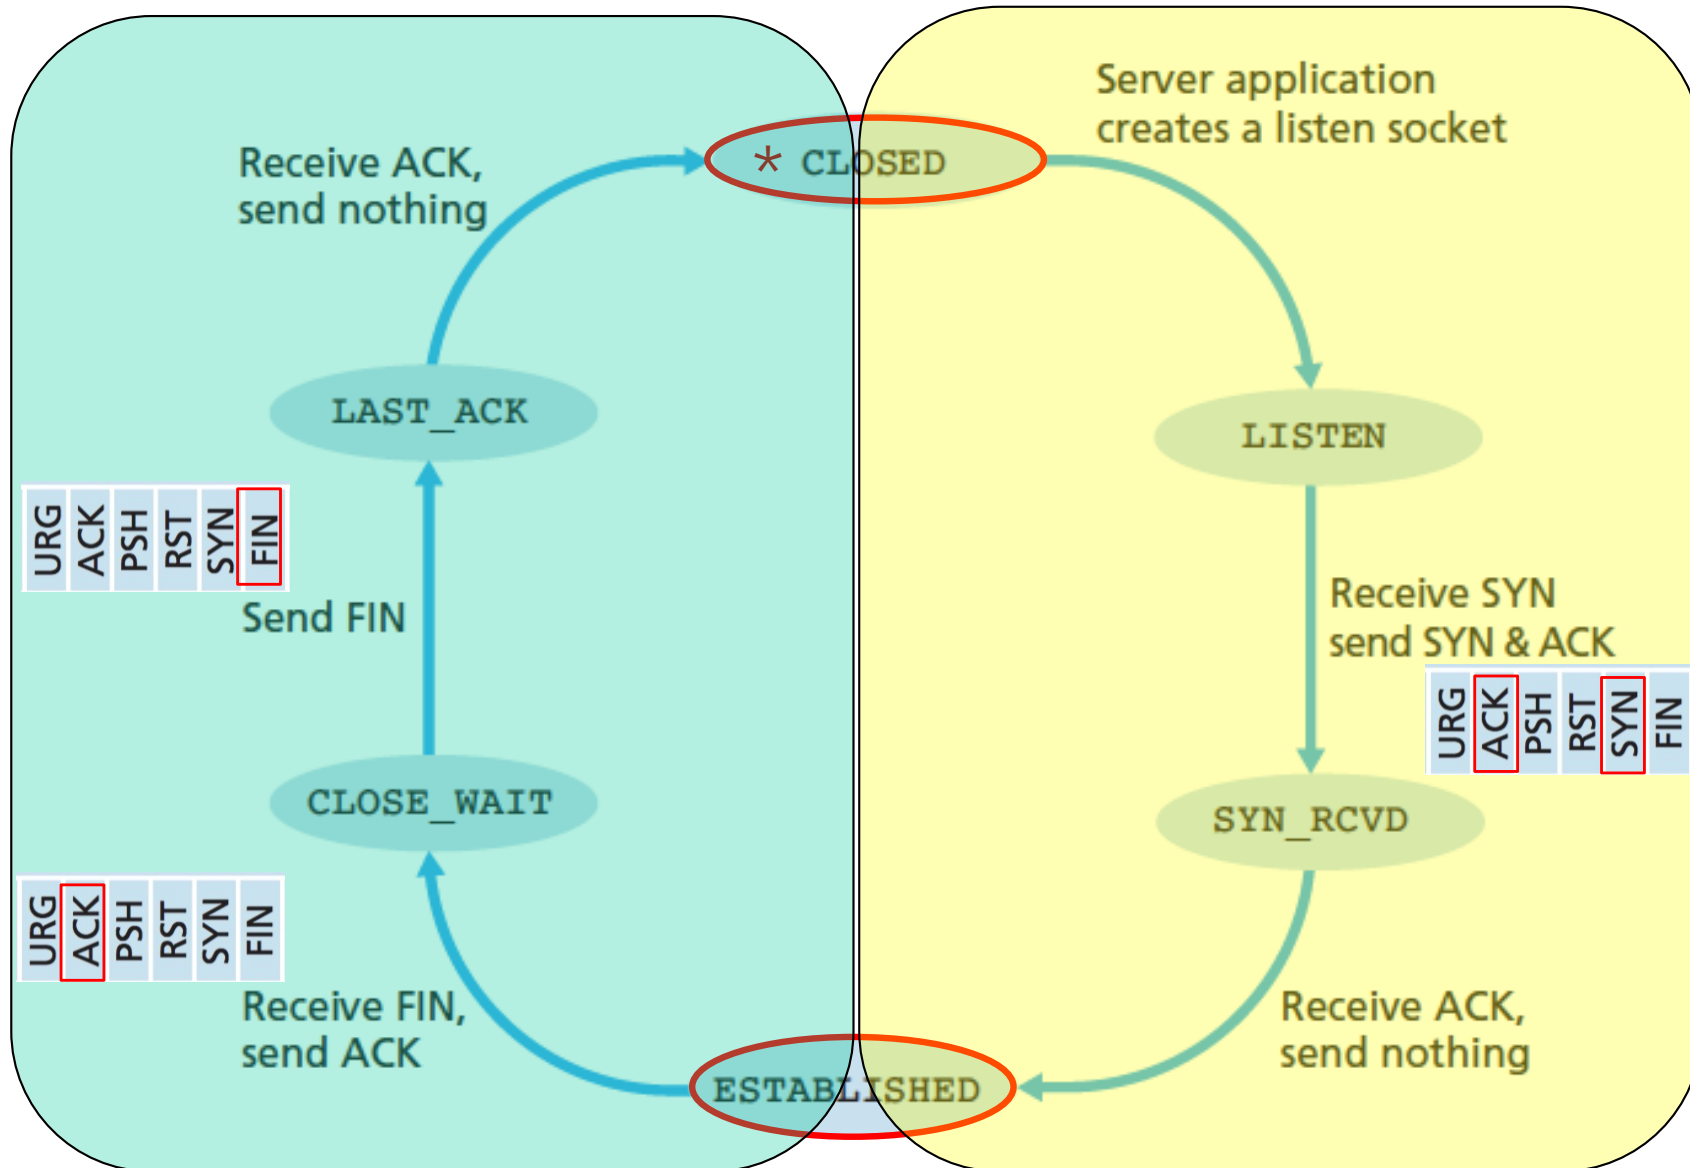
LAST_ACK

CLOSED

A typical sequence of TCP states visited by a **client** TCP



A typical sequence of TCP states visited by a server-side TCP



Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 Principles of congestion control

3.7 TCP congestion control



Principles of congestion control

Congestion


- ❖ Informally: “too many sources sending too much data too fast for *network* to handle”
- ❖ Different from flow control!
- ❖ Manifestations
 - Lost packets (buffer overflow at routers)
 - Long delays (queuing in router buffers)
- ❖ A top-10 problem!

Causes/costs of congestion: scenario I

- ❖ Two senders, two receivers
- ❖ One router, **infinite** buffers
- ❖ Output link capacity: R
- ❖ No retransmission




- ❖ maximum per-connection **throughput**: $R/2$





- ❖ large **delays** as arrival rate, λ_{in} , approaches capacity

Causes/costs of congestion: scenario 2



- ❖ One router, *finite* buffers
- ❖ Sender retransmission of timed-out packet
 - **Application-layer** input = application-layer output:
 $\lambda_{in} = \lambda_{out}$
 - **Transport-layer** input includes *retransmissions* : $\lambda'_{in} \geq \lambda_{in}$



- ❖ Sender sends only when router buffers available
- ❖ Packets can be lost, dropped at router due to full buffers
- ❖ Sender only resends if packet *known* to be lost



- ❖ Sender times out prematurely, sending two copies, both of which are delivered



“Costs” of congestion

- ❖ More work (retrans) for given “goodput”
- ❖ Unneeded retransmissions: link carries multiple copies of pkt
 - Decreasing goodput


Goodput : the **application level** throughput, i.e. the number of useful information bits delivered by the network to a certain destination per unit of time



Causes/costs of congestion: scenario 3

- ❖ Four senders
- ❖ Multihop paths
- ❖ Timeout/retransmit

Q: What happens as λ_{in} and λ_{in}' increase ?

A: As red λ_{in}' increases, all arriving blue pkts at upper queue are dropped, blue throughput $\rightarrow 0$





Another “cost” of congestion:

- ❖ When packet dropped, any “upstream transmission” capacity used for that packet was wasted!

Approaches towards congestion control

Two broad approaches towards congestion control:

End-end congestion control

- ❖ No explicit feedback from network
- ❖ Congestion inferred from end-system observed loss, delay
- ❖ Approach taken by **TCP**

Network-assisted congestion control

- ❖ Routers provide feedback to end systems
 - Single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - Explicit rate for sender to send at

ECN : Explicit Congestion Notification


Case study: ATM ABR congestion control

ABR: Available Bit Rate

- ❖ “Elastic service”
- ❖ If sender’s path “underloaded”
 - Sender should use available bandwidth
- ❖ If sender’s path congested
 - Sender throttled to minimum guaranteed rate

RM (Resource Management) cells

- ❖ Sent by sender, interspersed with data cells
- ❖ Bits in RM cell set by switches (“*network-assisted*”)
 - *NI bit*: No Increase in rate (mild congestion)
 - *CI bit*: Congestion Indication
- ❖ RM cells returned to sender by receiver, with bits intact

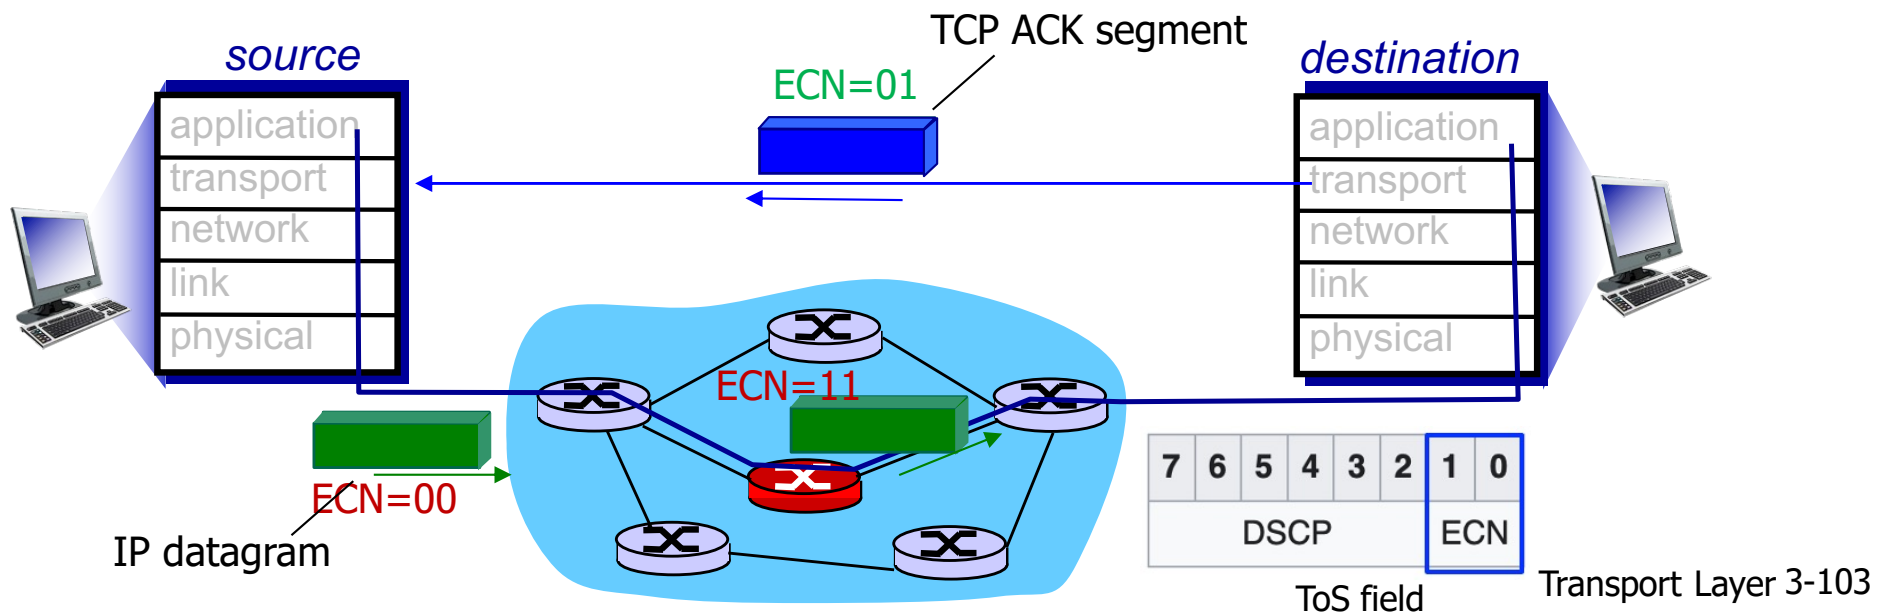


- ❖ Two-byte **ER (explicit rate)** field in **RM cell**
 - Congested switch may lower ER value in cell
 - Senders' send rate thus max supportable rate on path
- ❖ **EFCI bit** in **data cells**: set to 1 in congested switch
 - If data cell preceding RM cell has EFCI set, receiver sets **CI bit** in returned RM cell [**CI bit in RM cell, EFCI bit in data cell**]
 - EFCI bit set is useful since there might have congestion between the last switch and destination

Explicit Congestion Notification (ECN)

Network-assisted congestion control

- Two bits in IP header (**ToS field**) marked *by network router* to indicate congestion
- Congestion indication carried to receiving host
- Receiver (seeing congestion indication in IP datagram) sets **ECE** (ECN-Echo) bit on receiver-to-sender ACK segment to notify sender of congestion



Chapter 3 outline

3.1 Transport-layer services

3.2 Multiplexing and demultiplexing

3.3 Connectionless transport: UDP

3.4 Principles of reliable data transfer

3.5 Connection-oriented transport: TCP


- segment structure
- reliable data transfer
- flow control
- connection management

3.6 Principles of congestion control

3.7 TCP congestion control

TCP congestion control: Additive Increase Multiplicative Decrease (AIMD)


- ❖ *Approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *Additive increase*: increase **cwnd** by 1 MSS (Max Segment Size) every RTT until loss detected
 - *Multiplicative decrease*: cut **cwnd** in half after loss



AIMD

- saw tooth behavior
- probing for bandwidth

TCP Congestion Control: details



- ❖ Sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion


TCP sending rate

- ❖ Roughly: send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$


TCP Slow Start

- ❖ When connection begins, increase rate exponentially (double) until first loss event
 - Initially **cwnd** = 1 MSS
 - Double **cwnd** every RTT
 - Done by incrementing **cwnd** (1 MSS) for every ACK received
 - eg.: 1 MSS = 500 bytes, RTT = 200 msce then initial rate = 20 kbps [$\text{cwnd}/\text{RTT} = 500 \times 8 \text{ (bits)} / 0.2\text{s} = 20\text{kbps}$]
- ❖ Summary: initial rate is slow but ramps up exponentially fast (double)



TCP: detecting, reacting to loss

- ❖ Loss indicated by **timeout** : **TCP Tahoe**
 - **cwnd** set to 1 MSS (TCP Tahoe always sets cwnd to 1)
 - Window then grows exponentially (as in slow start) to threshold (**ssthresh**, cut in half), then grows **linearly (CA)**
- ❖ Loss indicated by **3 duplicate ACKs**: **TCP RENO**
 - Duplicate ACKs indicate network is still capable of delivering some segments
 - **cwnd** is cut in half window then grows **linearly (CA)**




TCP: switching from slow start to CA

Q: When should the exponential increase (SS) switch to linear (CA)? [TCP Tahoe]

A: When **cwnd** gets to 1/2 of its value before timeout

Implementation

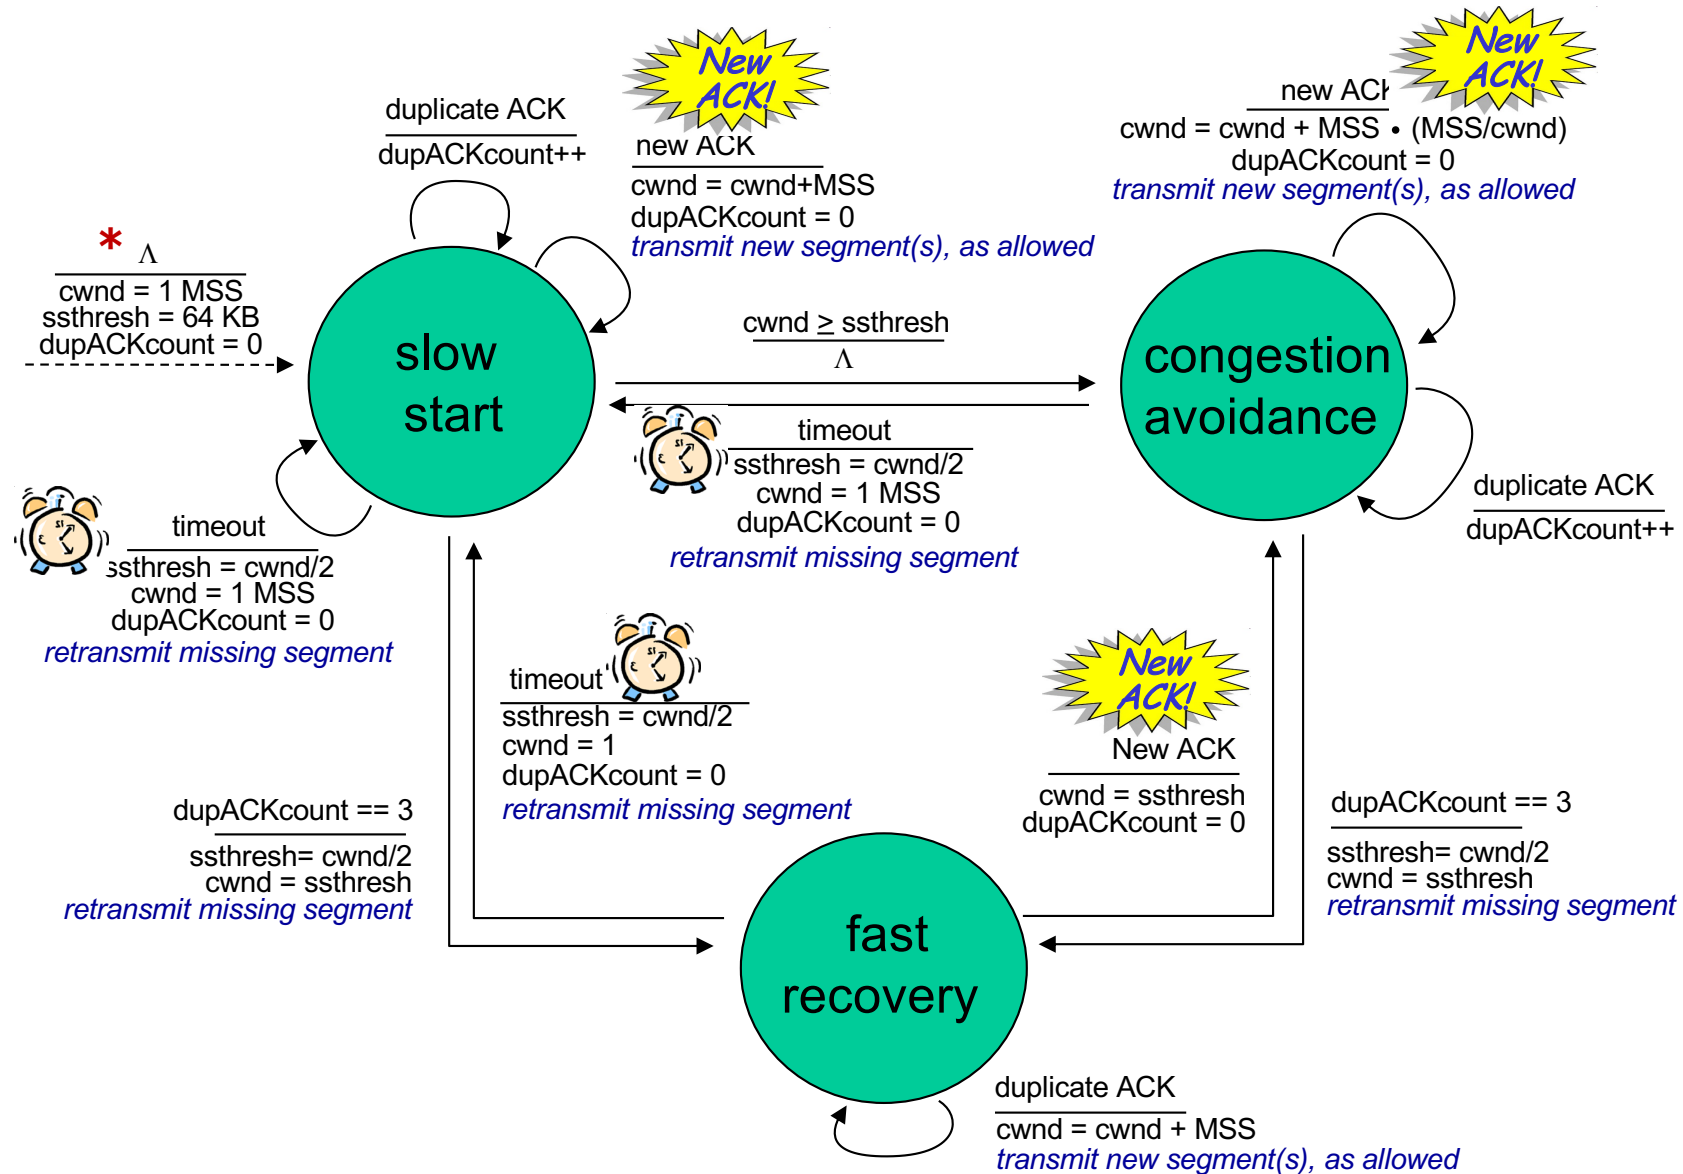
- ❖ Variable **ssthresh**
- ❖ On loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



TCP sender congestion control


State	Event	TCP sender action	Commentary
Slow Start (SS)	ACK receipt for previously unacked data	$\text{cwnd} = \text{cwnd} + \text{MSS}$, if $(\text{cwnd} > \text{ssthresh})$ set state to CA	resulting in a doubling of cwnd every RTT
Congestion Avoidance (CA)	ACK receipt for previously unacked data	$\text{cwnd} = \text{cwnd} + \text{MSS} * (\text{MSS} / \text{cwnd})$	additive increase, resulting in increase of cwnd by 1 MSS every RTT
SS or CA	loss event detected by 3 duplicate ACK	$\text{ssthresh} = \text{cwnd} / 2$, $\text{cwnd} = \text{ssthresh}$, set state to CA	fast recovery, implementing multiplicative decrease, cwnd will <u>not</u> drop below 1 MSS
SS or CA	timeout	$\text{ssthresh} = \text{cwnd} / 2$, $\text{cwnd} = 1 \text{ MSS}$, set state to SS	enter slow start
SS or CA	duplicate ACK	increment <u>duplicate ACK count</u> for segment being acked	cwnd and ssthresh not changed

Summary: TCP Congestion Control



TCP throughput

- ❖ Avg. TCP throughput as function of window size and RTT
 - Ignore slow start, assume always data to send
- ❖ W : window size (measured in bytes) where loss occurs
 - Avg. window size (# in-flight bytes) is $\frac{3}{4}W$ [$= (W + 0.5W) / 2$]
 - Avg. thruput [$cwnd/RTT$] is $\frac{3}{4}W$ per RTT
avg TCP thruput = $\frac{3}{4} \frac{W}{RTT}$ bytes/sec



TCP Futures: TCP over “long, fat pipes”

- ❖ Example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput [cwnd/RTT]

- ❖ Requires window size $W = 83,333$ in-flight segments [83,333 (seg.) * 1500 (byte) * 8 (bit)] / 0.1 (s) = 10 Gbps

- ❖ Throughput in terms of segment loss probability, L [Mathis 1997]:


$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \sqrt{L}} \quad [\text{Problem P45}]$$

→ To achieve 10 Gbps throughput, need a loss rate of $L = 2 \cdot 10^{-10}$ – *a very small loss rate!*

- ❖ New versions of TCP for high-speed

TCP Fairness


Fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K



Why is TCP fair?

Two competing sessions

- ❖ Additive increase gives slope of 1, as throughput increases
- ❖ Multiplicative decrease decreases throughput proportionally



R is the router capacity.

Fairness (more)

Fairness and UDP

- ❖ Multimedia apps often do not use TCP
 - Do not want rate throttled by congestion control
- ❖ Instead use UDP
 - Send audio/video at constant rate, tolerate packet loss

Fairness, parallel TCP connections

- ❖ Application can open multiple parallel connections between two hosts
- ❖ e.g., link of rate R with 9 existing connections
 - New app asks for 1 TCP, gets rate $R/10$ $R/(1+9) = R/10$
 - New app asks for 2 TCPs, gets $2R/11$ $2R/(2+9) = 2R/11 > R/10$

Chapter 3: summary

- ❖ Principles behind transport layer services
 - Multiplexing, demultiplexing
 - Reliable data transfer
 - Flow control
 - Congestion control
- ❖ Instantiation, implementation in the Internet
 - UDP
 - TCP

Next

- ❖ Leaving the network “edge” (application, transport layers)
- ❖ Into the network “core”